




# **Chapter 5 – Part 2**

## **Large and Fast Memory: Exploiting Memory Hierarchy**

# Virtual Memory

- There is a need to run programs that are too large to fit in memory
- Solution adopted in the 1960s, split programs into little pieces, called overlays (pages)
  - Kept on the disk, swapped in and out of main memory
- Virtual memory : each program has its own address space, broken up into chunks called pages
  - Each page is a continuous range of addresses that are mapped onto physical memory

# Virtual Memory




# Virtual Memory

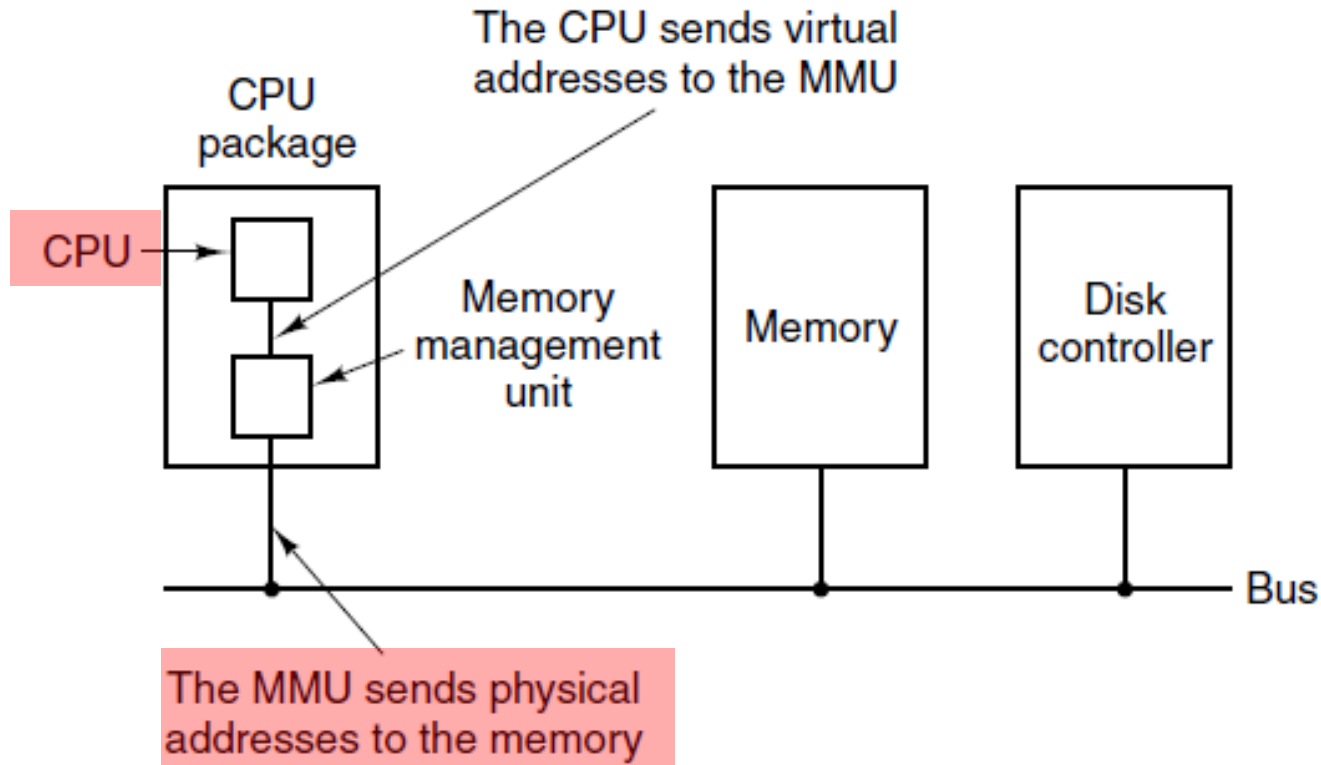
- Use main memory as a “cache” for secondary (disk) storage
- Programs share main memory
- CPU and OS translate virtual addresses to physical addresses
  - VM “block” is called a page

# Virtual Memory

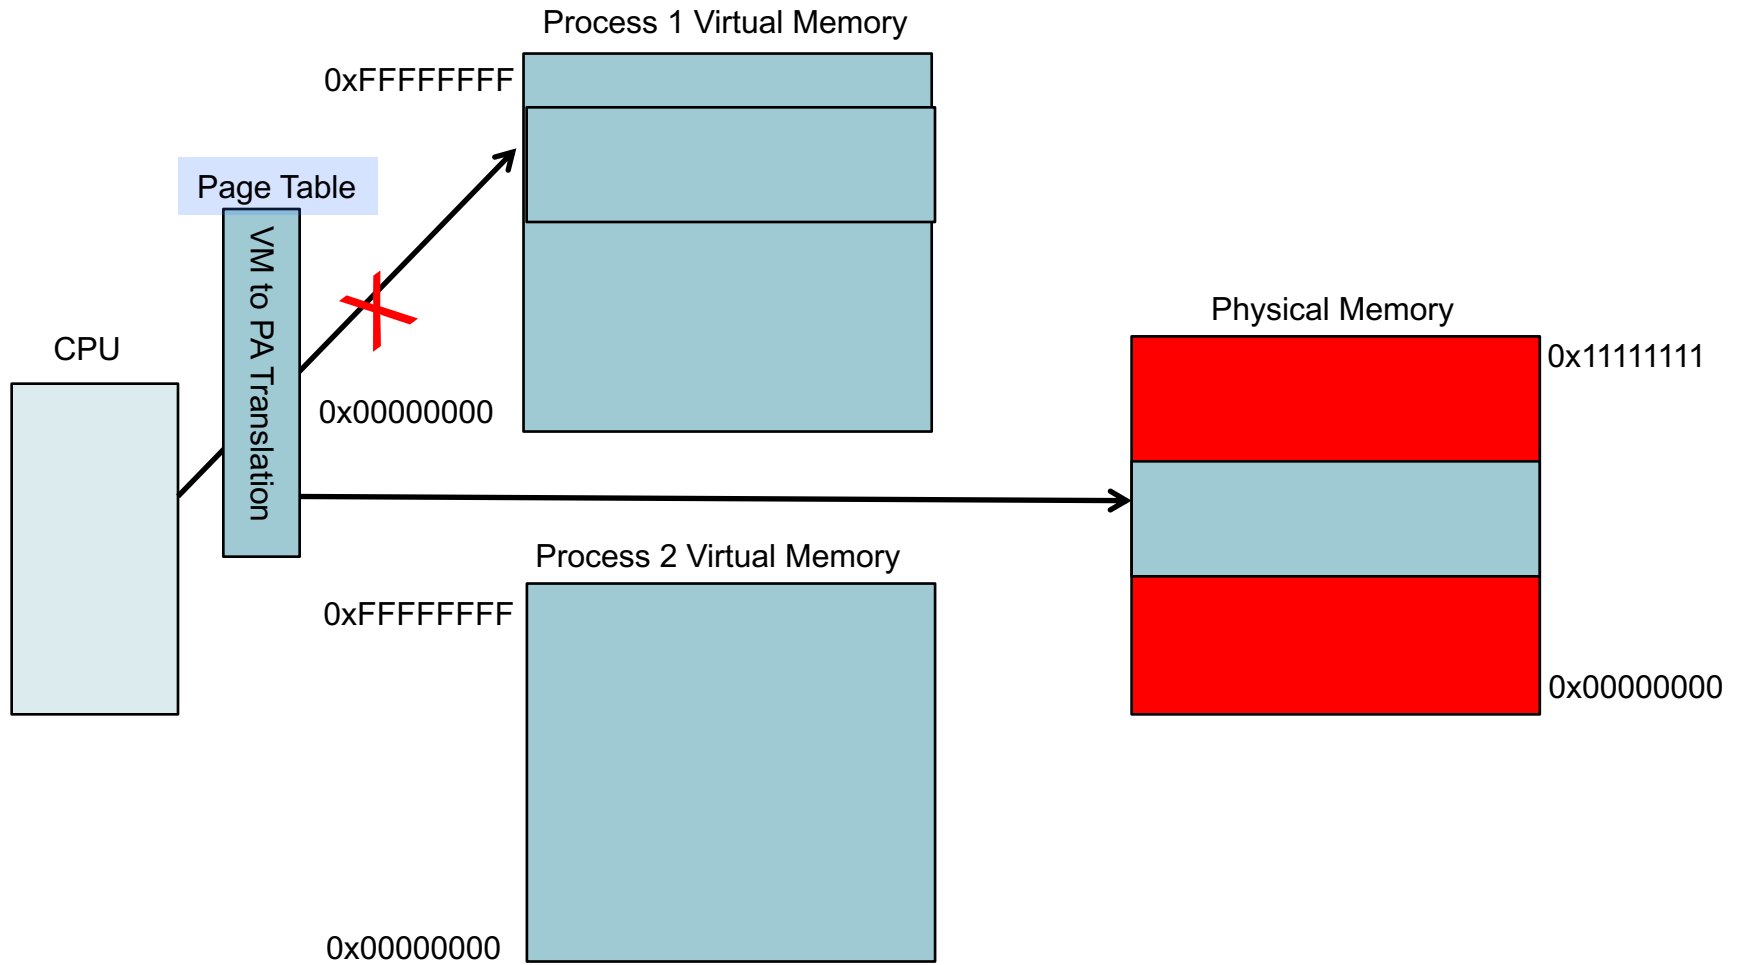
- VM adds a level of indirection between the virtual program addresses (VA) and the physical RAM addresses (PA)
- But there is a bit of a price (performance) to pay for
  - VA to PA translation (some bookkeeping required which impacts the performance)
- Three ways to achieve this VA to PA translation:
  1. VA to PA translation on the fly using a Memory Management Unit (MMU)
  2. Page tables
  3. Fast translation via TLB



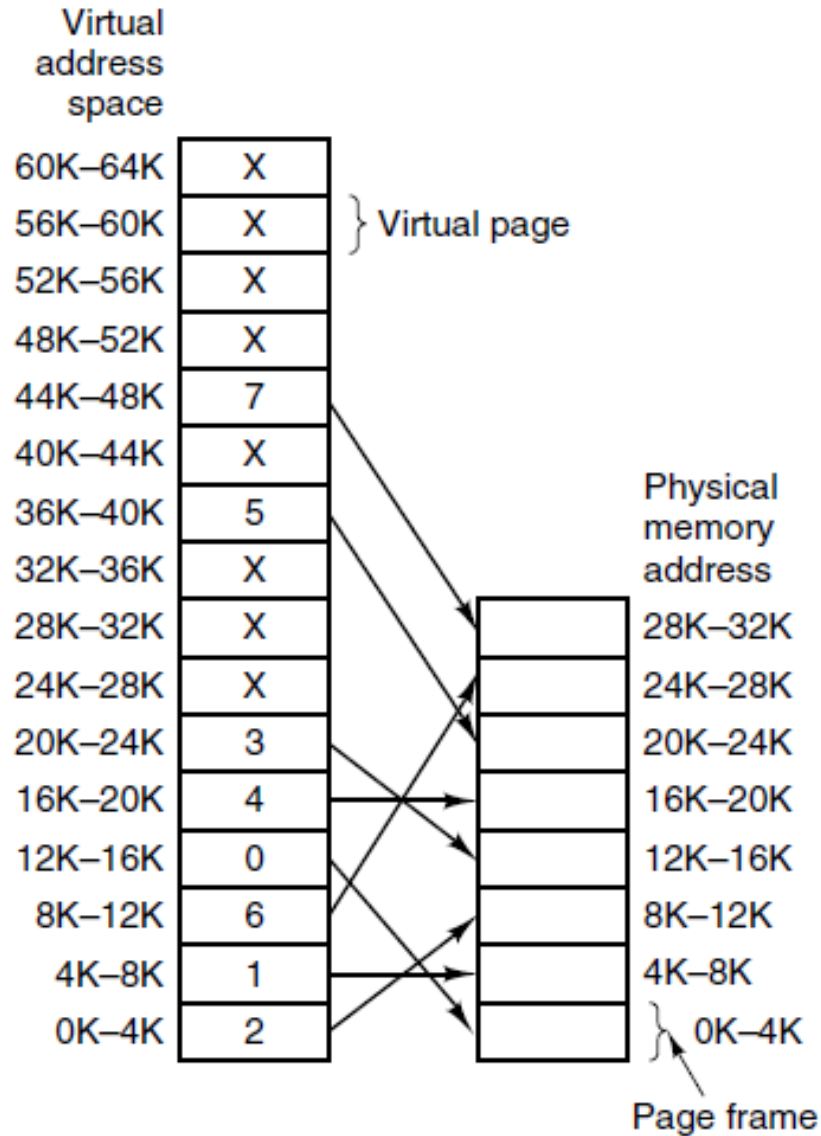
# Paging



# Paging

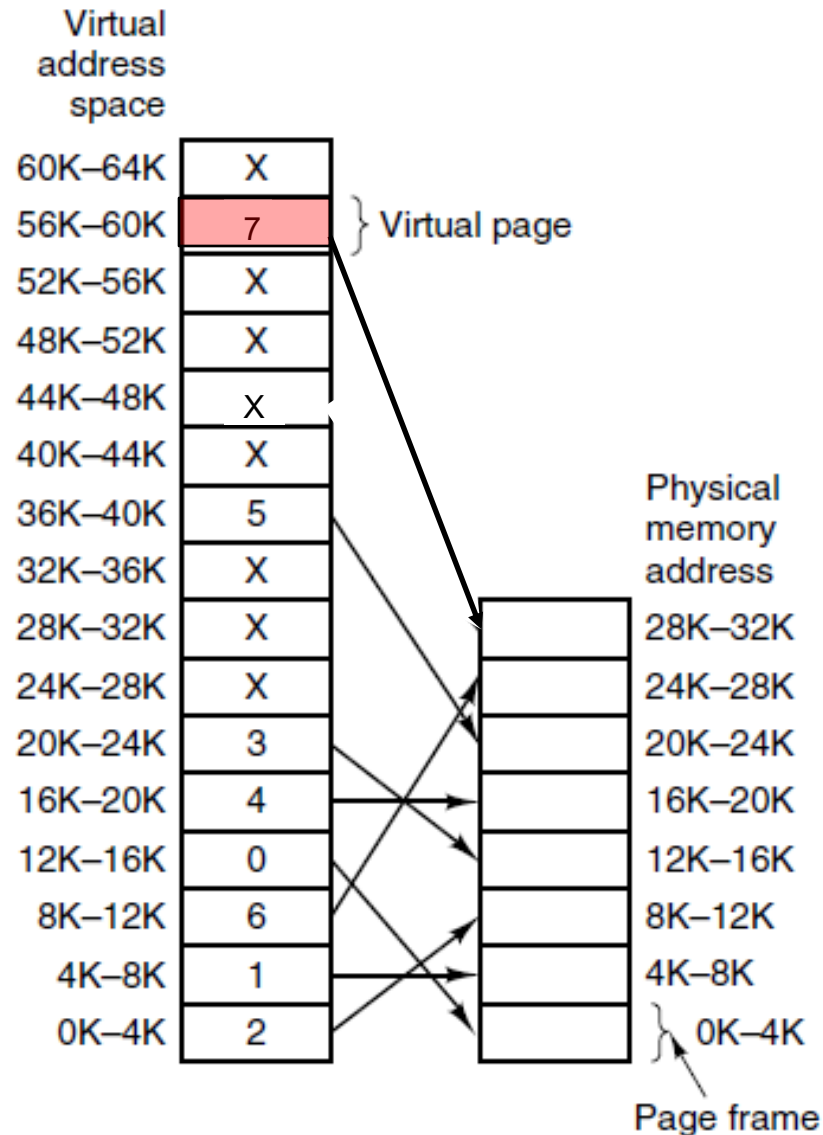


# Paging



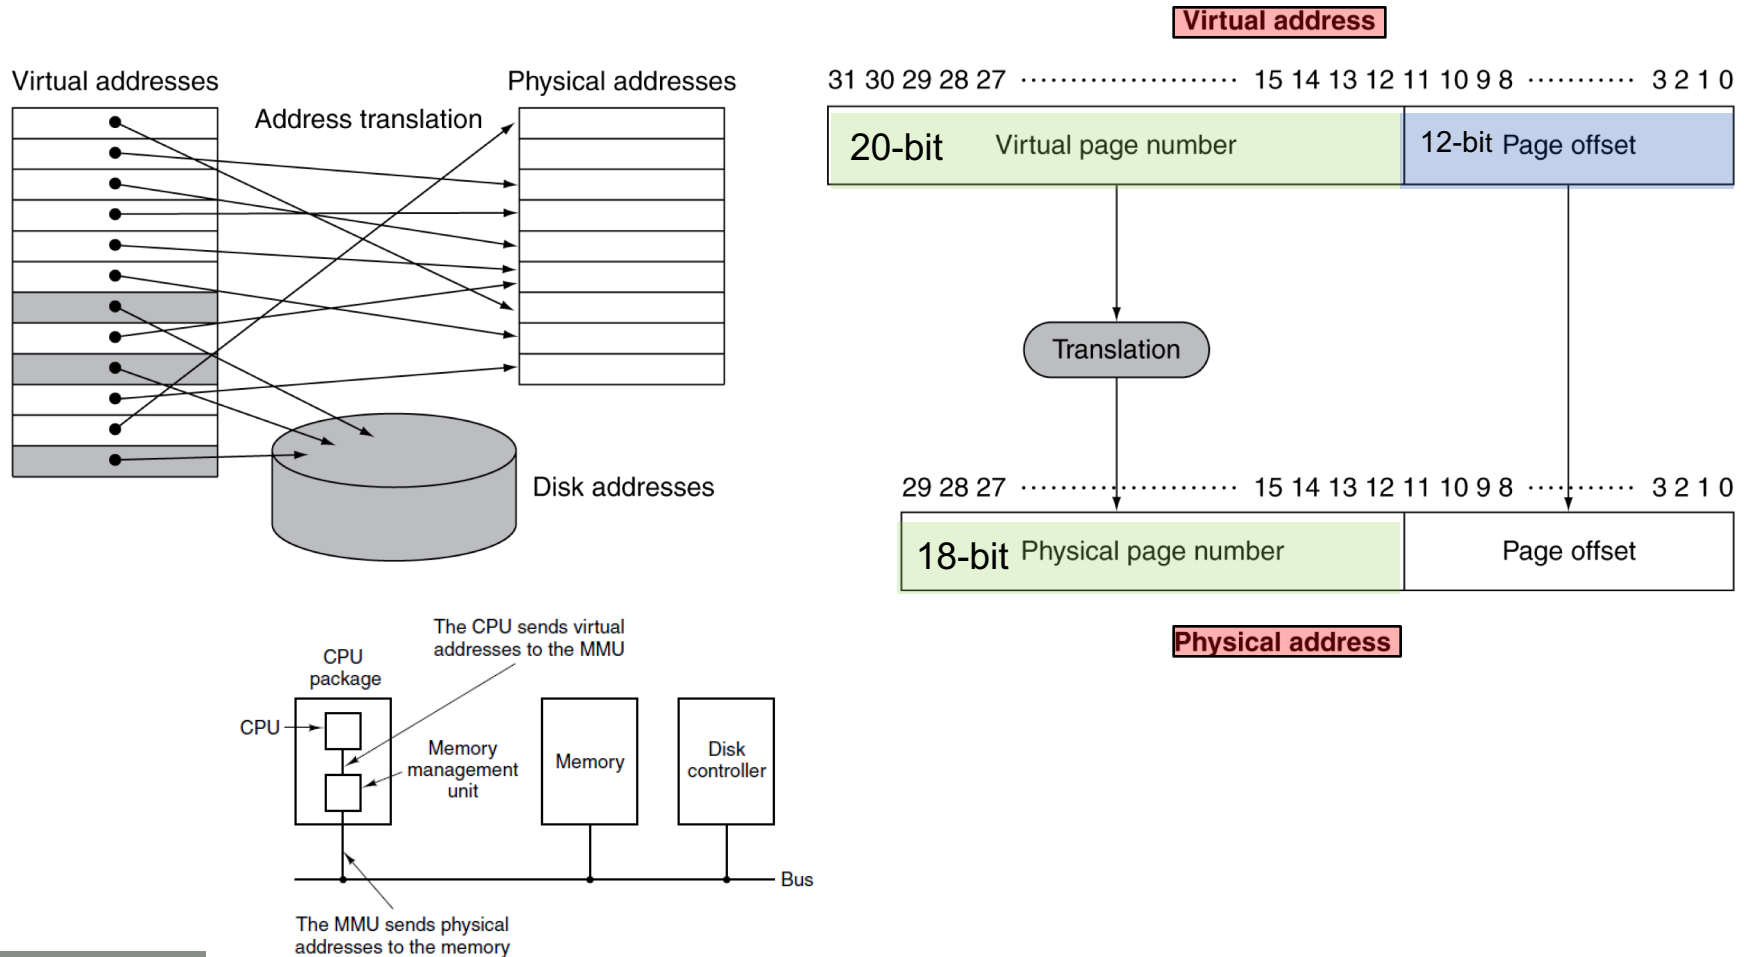


# Paging



# Address Translation

## ■ Fixed-size pages (e.g., 4K)



# Mapping Virtual to Physical Address


32-bit word addressing 4KB-page system

- ✦ Split address from CPU into two pieces
  - Page number ( $p$ )
  - Page offset ( $d$ )
- ✦ Page number
  - Index into page table
  - Page table contains base address of page in physical memory
- ✦ Page offset
  - Added to base address to get actual physical memory address
- ✦ Page size =  $2^d$  bytes


Example:

- 4 KB (=4096 byte) pages
- 32 bit logical addresses

$$2^d = 4096 \longrightarrow d = 12$$



$$32 - 12 = 20 \text{ bits} \quad 12 \text{ bits}$$

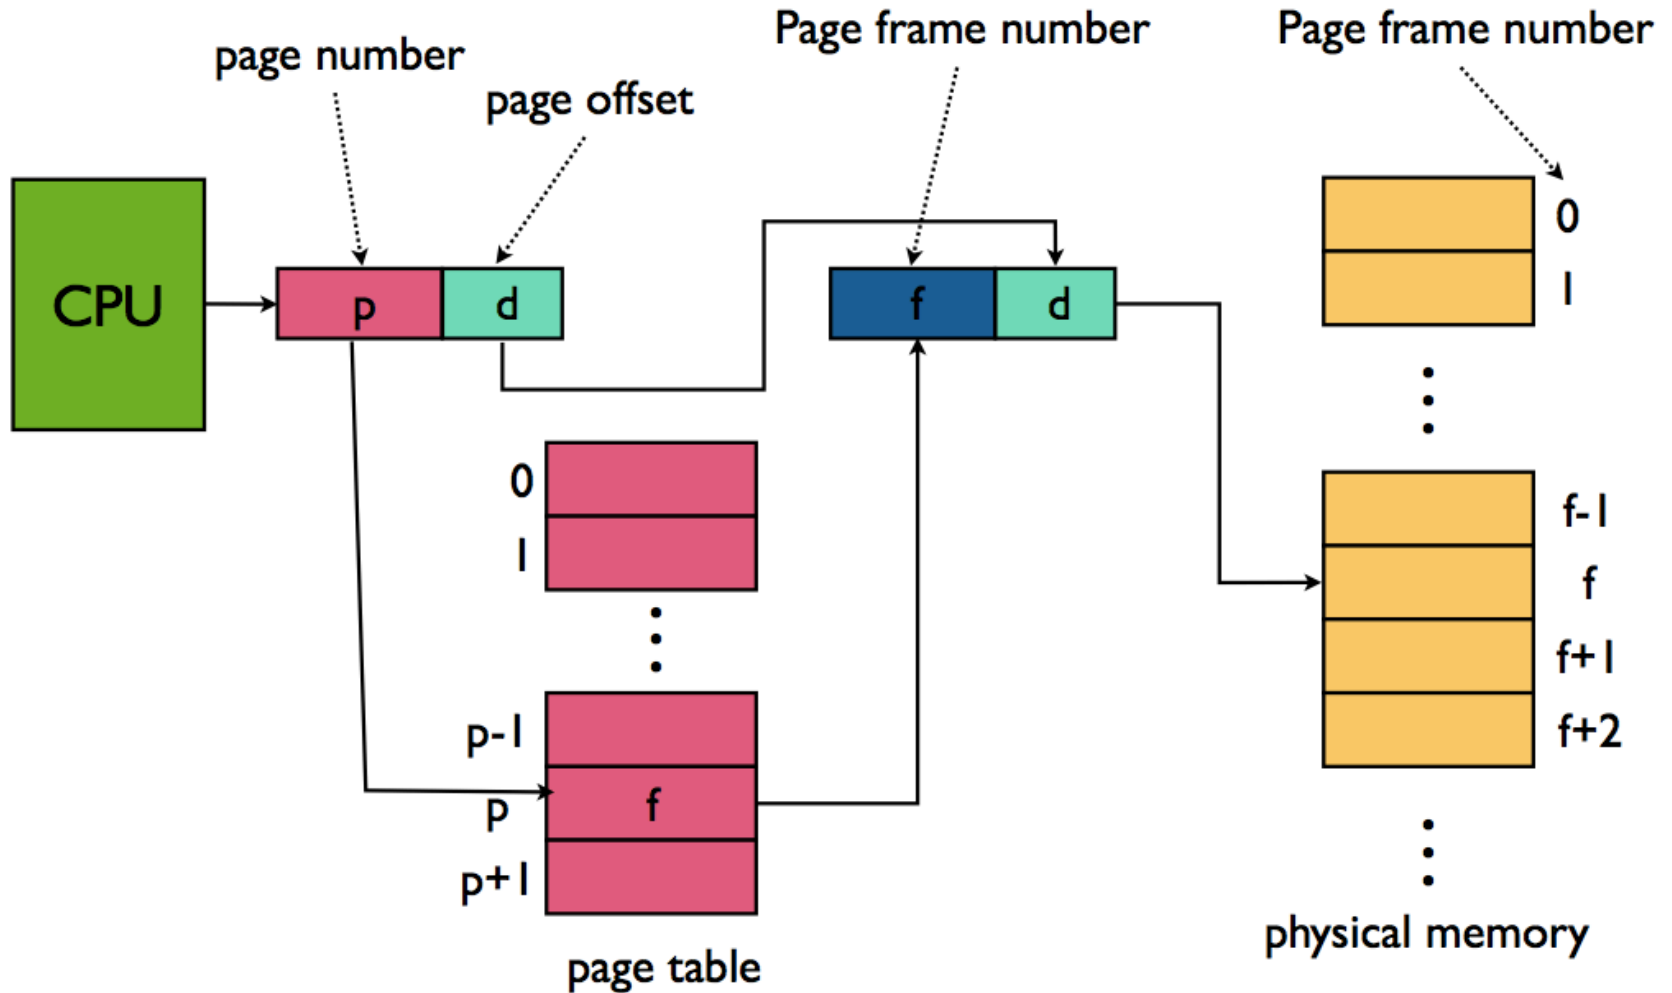


32 bit logical address

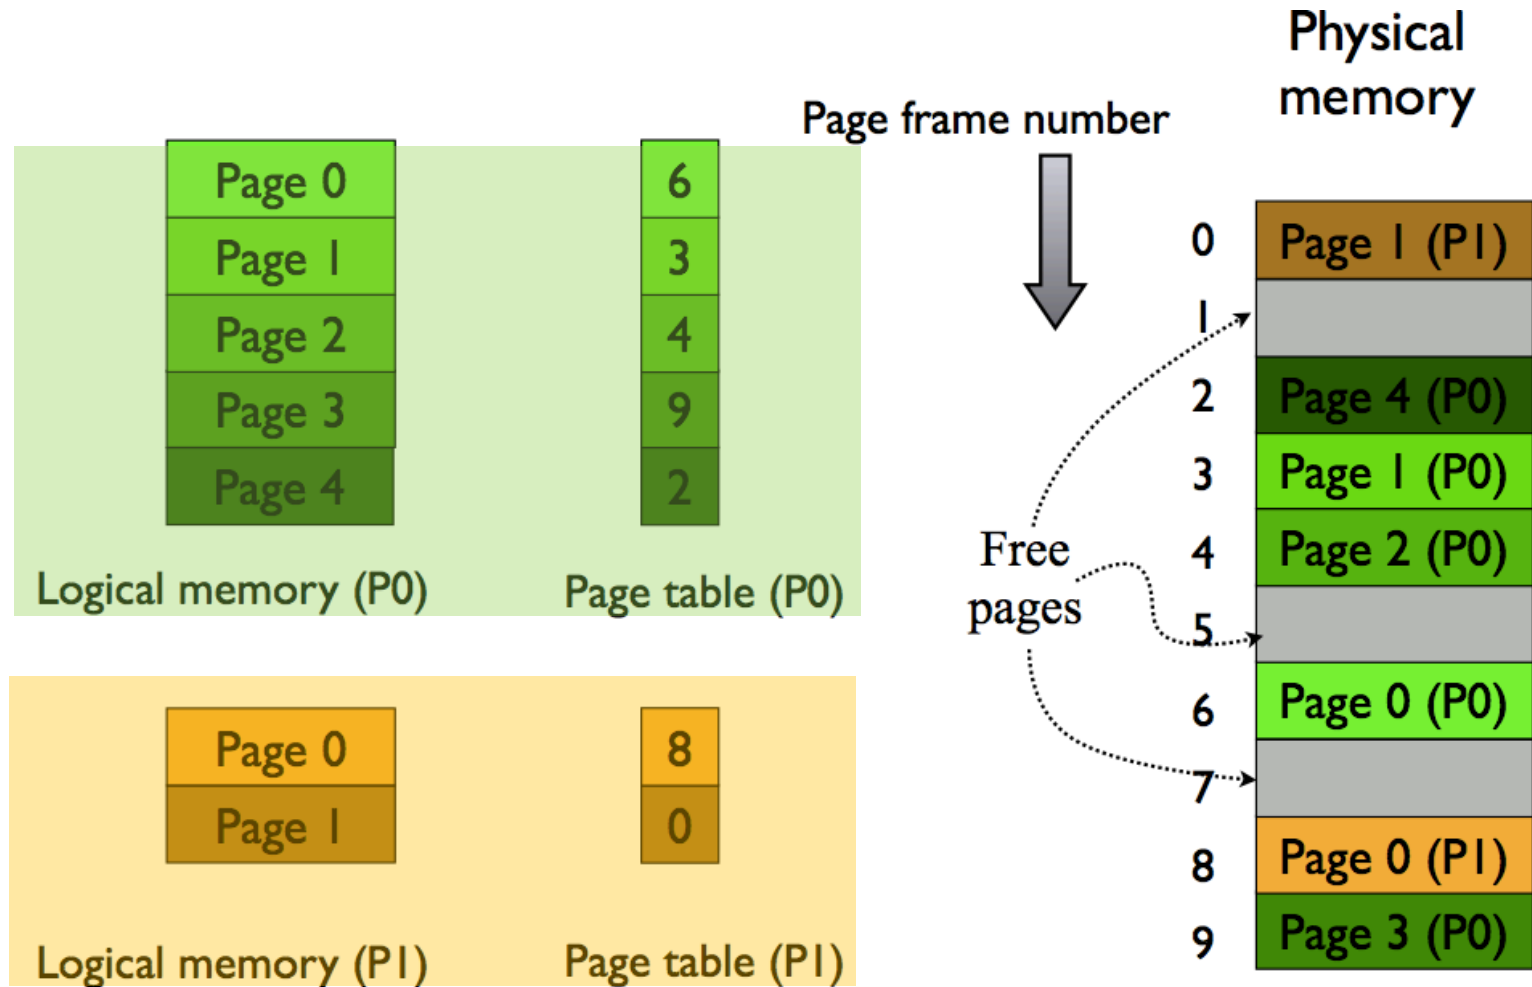
# Page Tables

- Stores placement information
- If page is present in memory
- If page is not present (page fault)

# Translation Using a Page Table



# Multiple Processes in Memory



# Page Replacement Algorithms

- Optimal algorithm
- Not recently used algorithm
- First-in, first-out (FIFO) algorithm
- Second-chance algorithm
- Clock algorithm
- Least recently used (LRU) algorithm
- Working set algorithm
- WSClock algorithm

# Optimal Algorithm

- Best possible but almost impossible to implement
- Almost impossible because the OS has no knowledge on when each page will be referenced next
- Possible solution
  - Run the program on simulator for a specific input data and keep track of those numbers
  - Implement optimal page replacement on the second run based on the info obtained the first time




# First-in, First-out (FIFO) Algorithm

- ✦ Maintain a linked list of all pages
  - Maintain the order in which they entered memory
- ✦ Page at front of list replaced
- ✦ Advantage: (really) easy to implement
- ✦ Disadvantage: page in memory the longest may be often used
  - This algorithm forces pages out regardless of usage
  - Usage may be helpful in determining which pages to keep

# Second-Chance Algorithm

- ♦ Modify FIFO to avoid throwing out heavily used pages
  - If reference bit is 0, throw the page out
  - If reference bit is 1
    - Reset the reference bit to 0
    - Move page to the tail of the list
    - Continue search for a free page
- ♦ Still easy to implement, and better than plain FIFO



# Not Recently Used Algorithm

- At page fault, the OS inspects pages and categorizes them based on the current values of their *R* (reference) and *M* (dirty) bits:

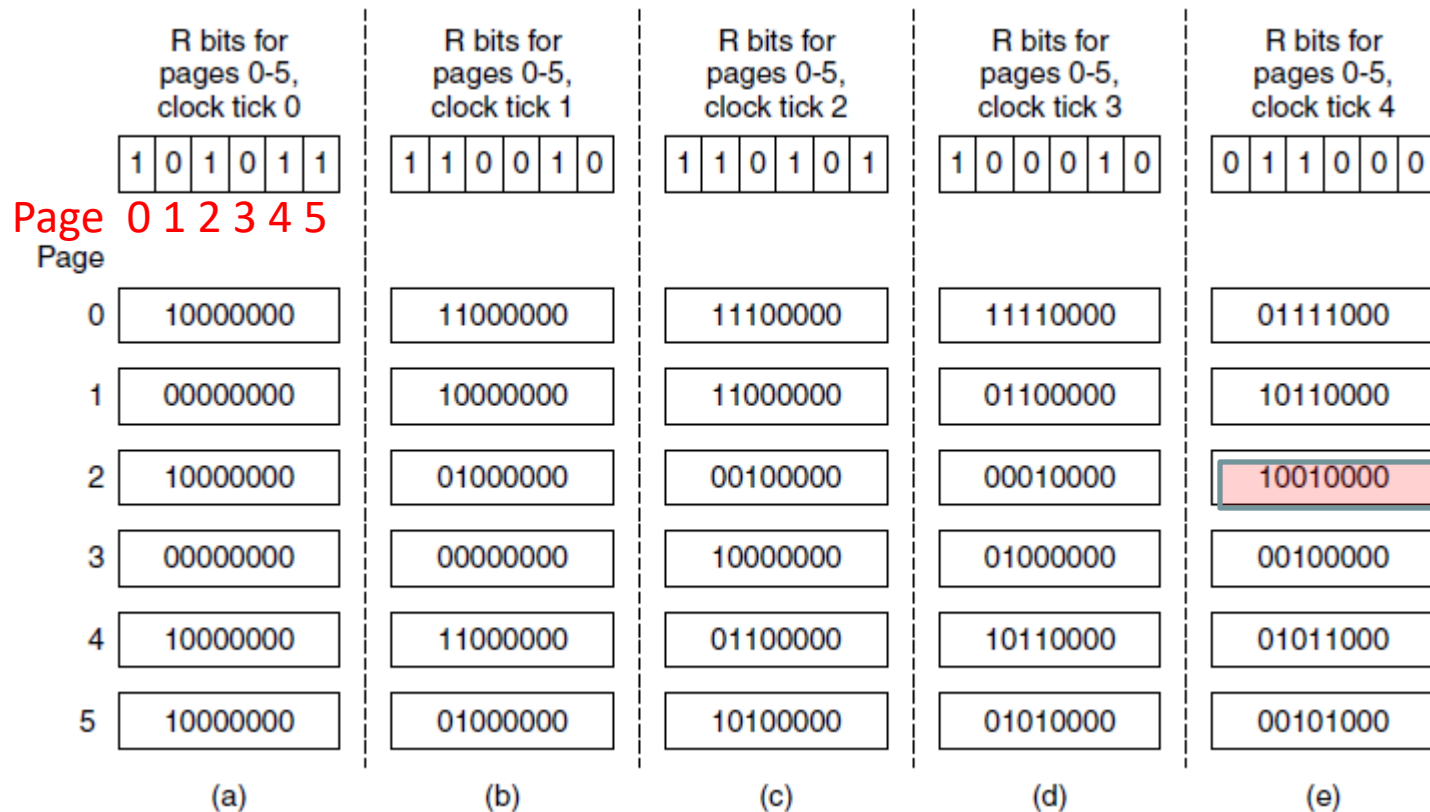
Class 0: not referenced, not modified.

Class 1: not referenced, modified.

Class 2: referenced, not modified.

Class 3: referenced, modified.

# Simulating LRU in Software



# Local versus Global Allocation Policies

	Age		
A0	10	A0	A0
A1	7	A1	A1
A2	5	A2	A2
A3	4	A3	A3
A4	6	A4	A4
A5	3	A6	A5
B0	9	B0	B0
B1	4	B1	B1
B2	6	B2	B2
B3	2	B3	A6
B4	5	B4	B4
B5	6	B5	B5
B6	12	B6	B6
C1	3	C1	C1
C2	5	C2	C2
C3	6	C3	C3
(a)		(b)	(c)

t = 3 (oldest in A)

t = 2 (oldest of all)

Local versus global page replacement.

(a) Original configuration. (b) Local page replacement.

(c) Global page replacement.

# Multi-level Page Tables

- Problem: page tables can be too large
  - In 4KB page size, with 32 bit logical addresses
    - $2^{12} = 4096$  (bytes), so 12 bits to get to the page byte offset
    - $2^{20} = 1$  million PTEs , so 20 bits to get to the page address
- Solution: use multi-level page tables
  - 1st level page table has pointers to 2nd level page tables
  - 2nd level page table has actual physical page numbers in it
- <http://www.youtube.com/watch?v=Z4kSOv49GNc>

# Fast Translation Using a TLB

- A widely implemented scheme for
  - Speeding up paging
  - Handling large virtual address spaces
- Small HW device for mapping virtual address to physical address without going through the page table in memory if possible
- <http://www.youtube.com/watch?v=uyrSn3qbZ8U>
- <http://www.youtube.com/watch?v=95QpHJX55bM>

# Concluding Remarks

- Fast memories are small, large memories are slow
  - We really want fast, large memories ☹️
  - Caching gives this illusion 😊
- Principle of locality
  - Programs use a small part of their memory space frequently
- Memory hierarchy
  - L1 cache ↔ L2 cache ↔ ... ↔ DRAM memory ↔ disk
- Memory system design is critical for multiprocessors