

Chapter 4

The Processor – Part 2

Control Hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
- In LEGv8 pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to find the outcome in earlier stages or use prediction

PC+4 :

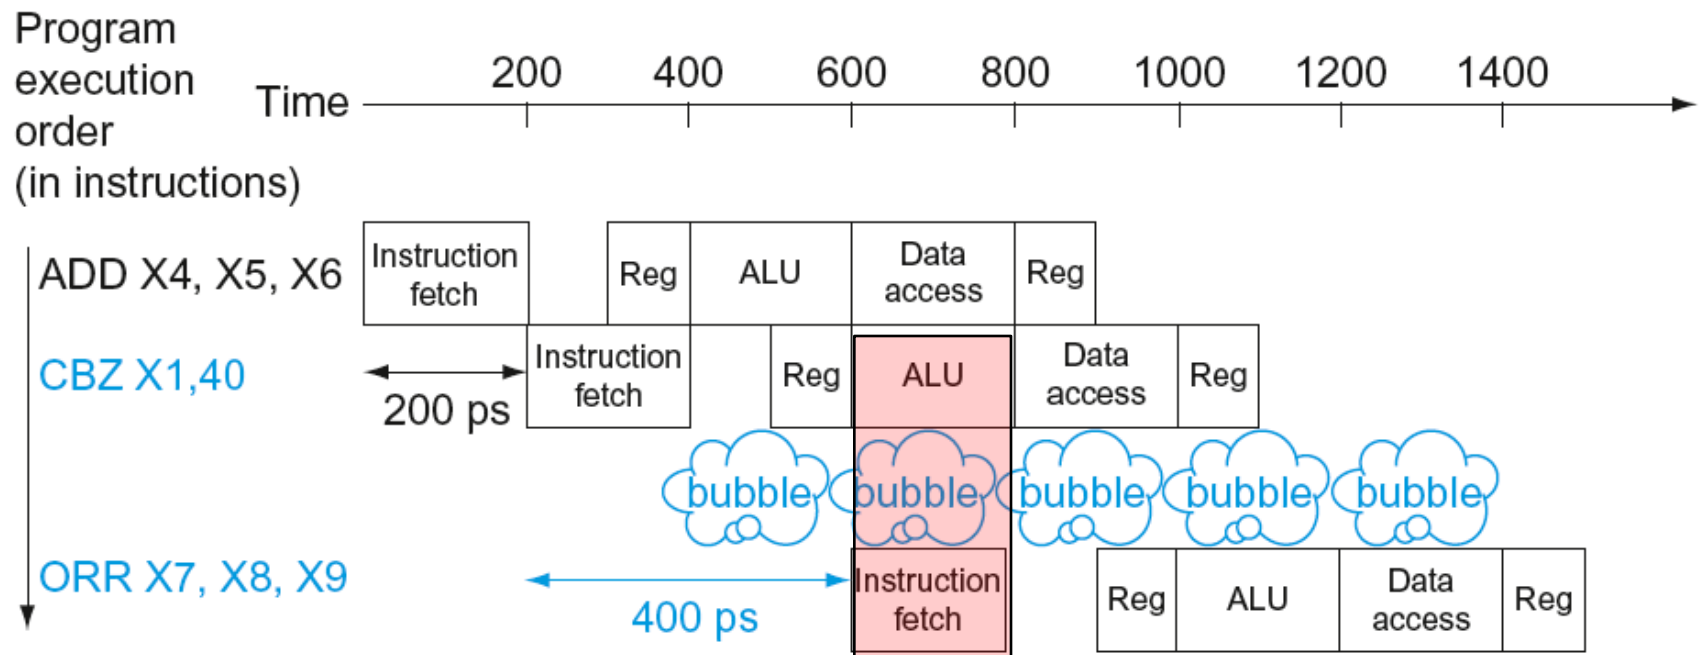
.
.
SUB X1, X2, X3
CBZ X1, 40
LDUR X3, [X0, #0]

PC+40 :

.
.
.
ORR X7, X8, X9

Stall on Branch

- Wait until branch outcome determined before fetching next instruction



Stall on Branch

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong

Branch Prediction

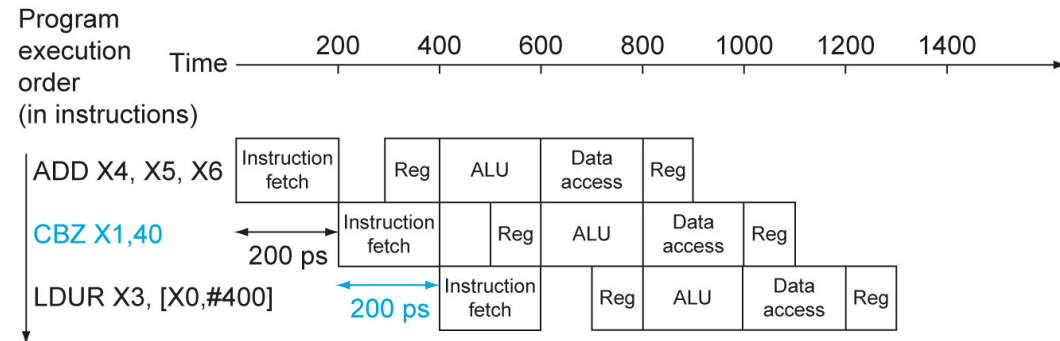
- Predict outcome of branch
 - Only stall if prediction is wrong
- In LEGv8 pipeline
 - Predict branches are not taken
 - Fetch instruction after branch instruction , with no delay

PC+4 : CBZ X1, 40
LDUR X3, [X0, #400]

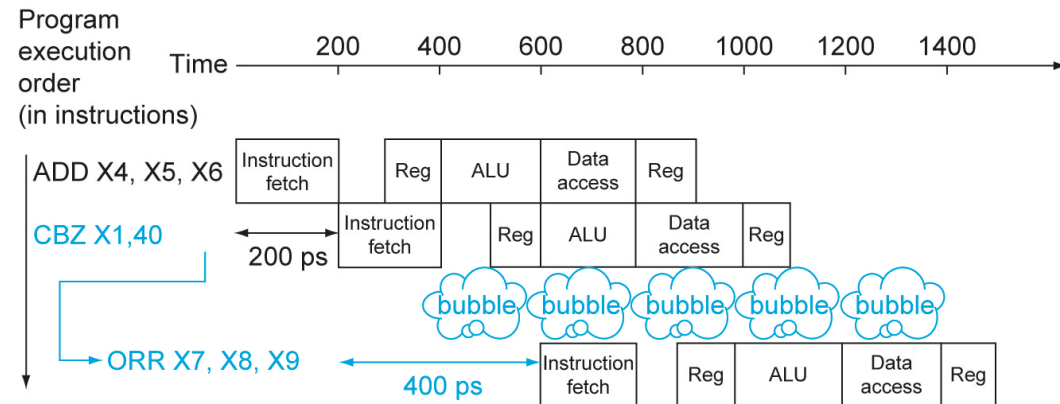
PC+40 : ORR X7, X8, X9

Branch Prediction

- Predict branches not taken



- What if the prediction is wrong?



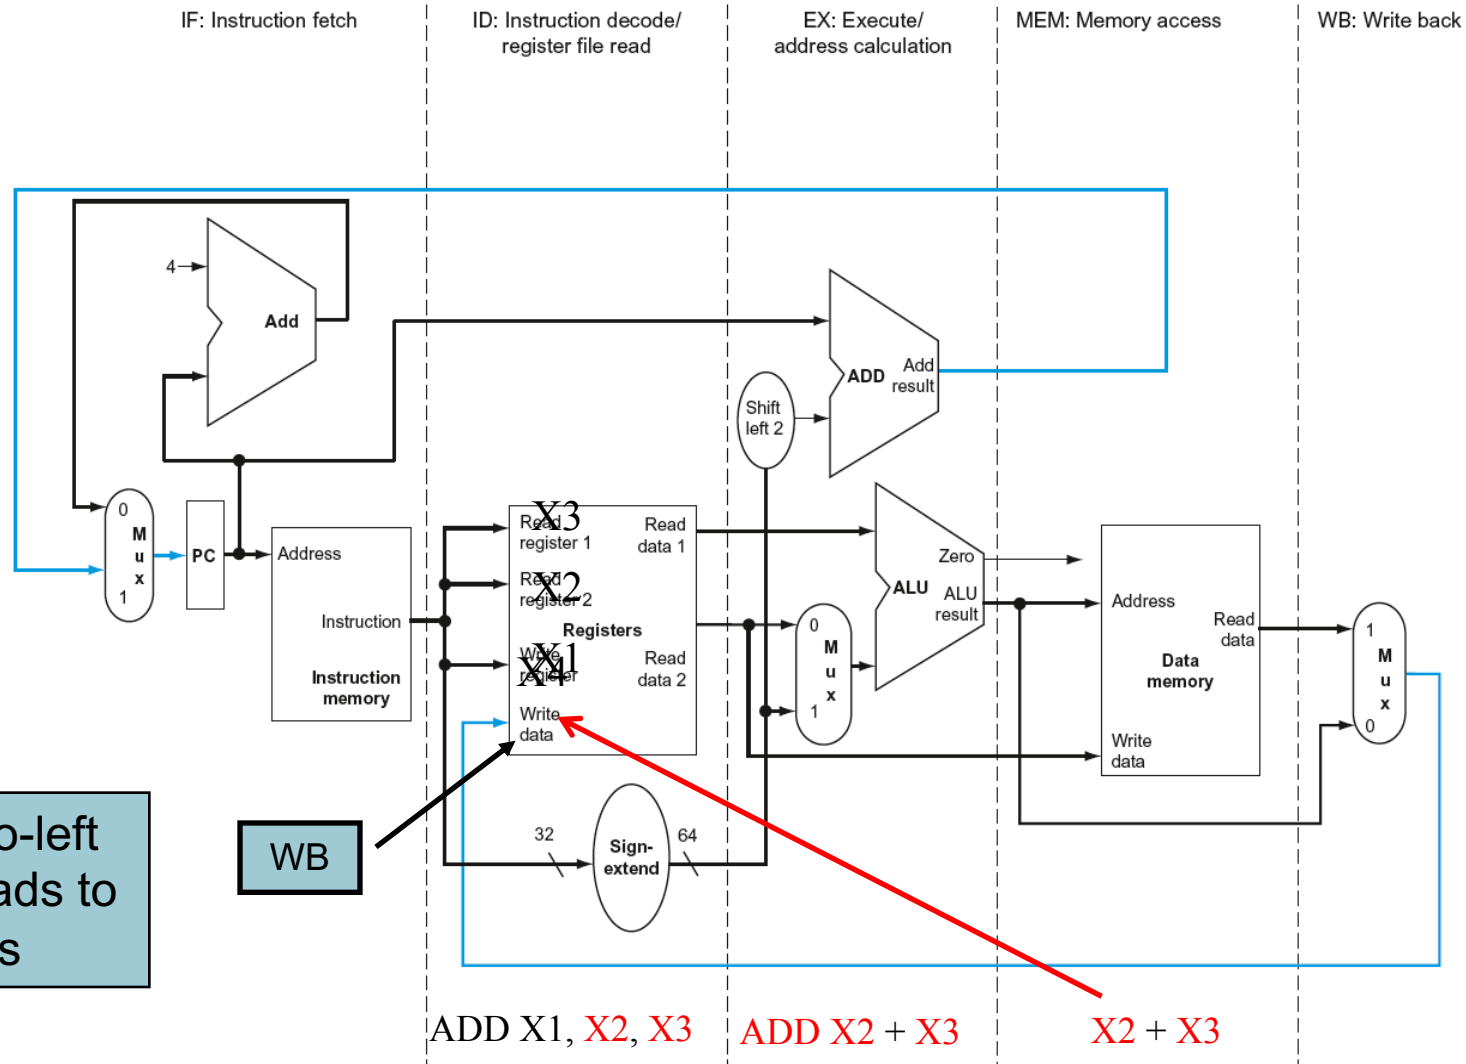
More-Realistic Branch Prediction

- Dynamic branch prediction
 - Hardware measures actual branch behavior
 - Assume future behavior will continue the trend

Pipeline Summary (So Far)

- Pipelining improves performance by increasing instruction throughput
- Subject to hazards
 - Structure, data, control
- Instruction set design affects complexity of pipeline implementation

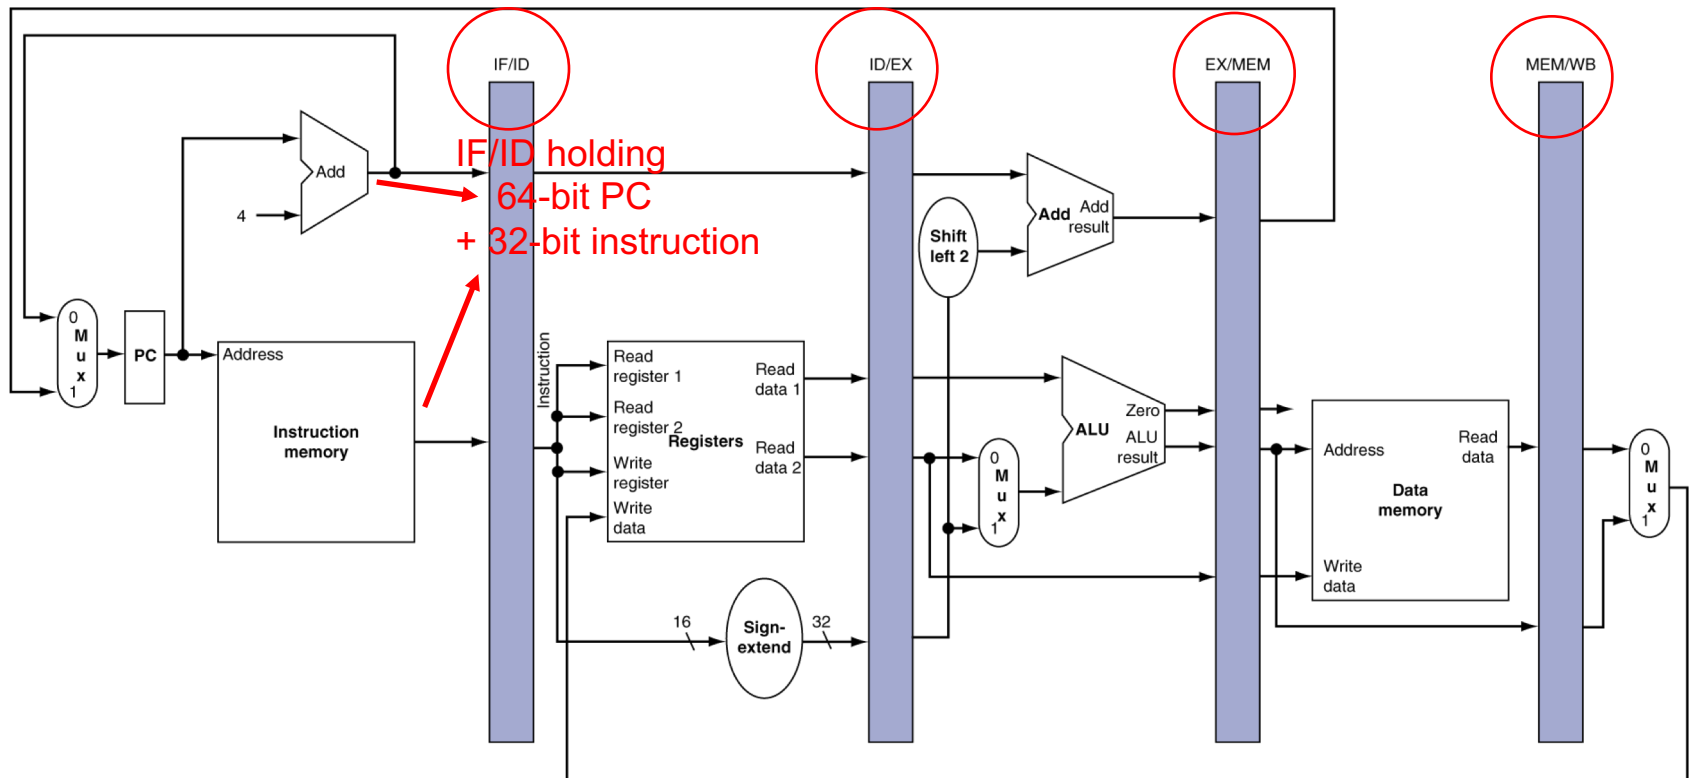
LEGv8 Pipelined Datapath



Right-to-left
flow leads to
hazards

Pipeline registers

- Need registers between stages
 - To hold information produced in previous cycle




Pipeline Operation

- Cycle-by-cycle flow of instructions through the pipelined datapath
 - “Single-clock-cycle” pipeline diagram
 - Shows pipeline usage (snapshot) in a single cycle
 - Highlight resources used
 - c.f. “multi-clock-cycle” diagram
 - Graph of operation/instructions over time

Single-Cycle Pipeline Diagram

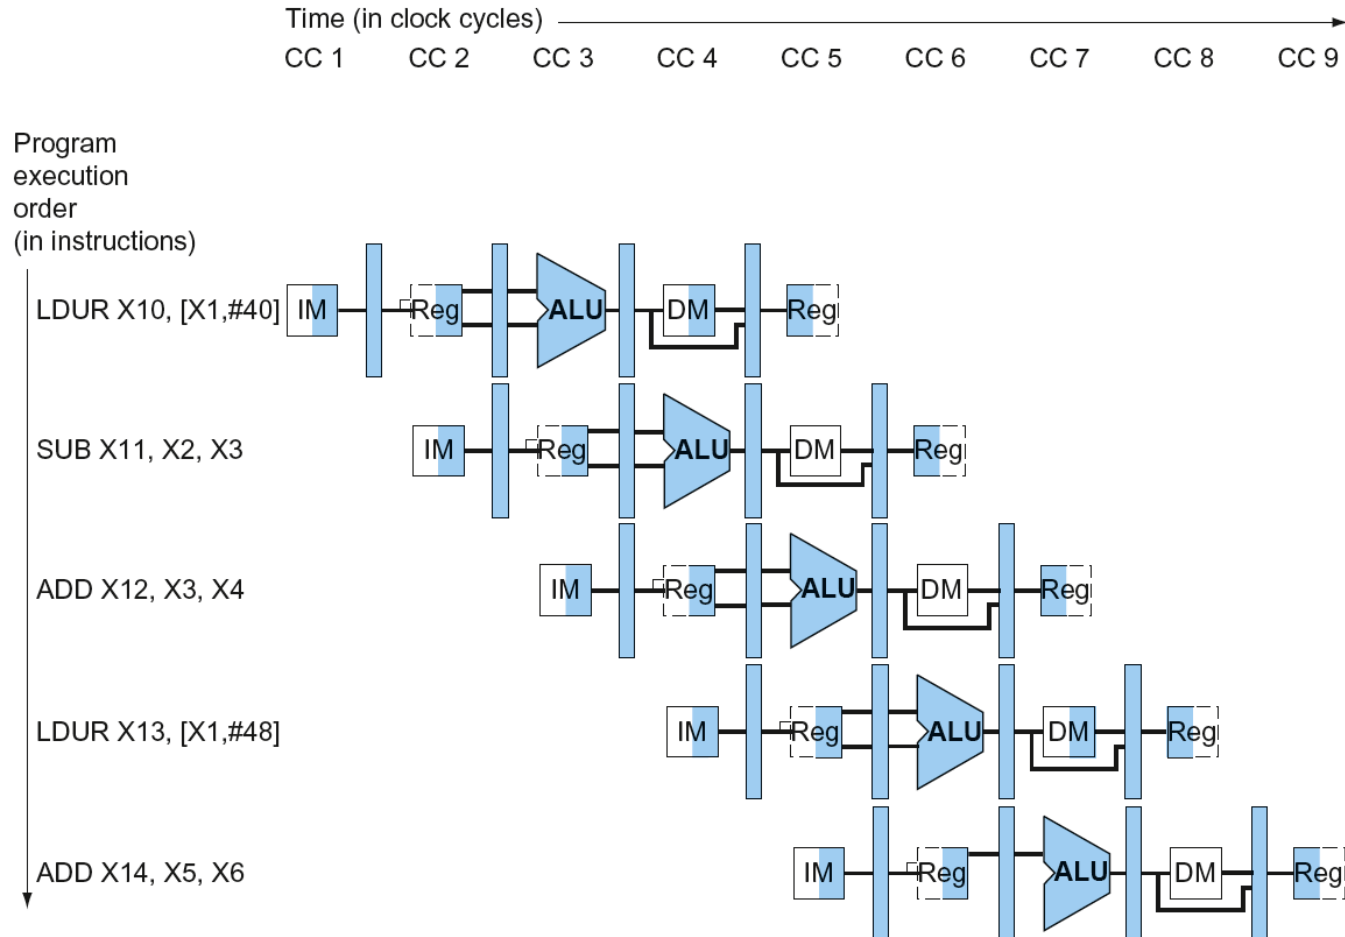
■ State of pipeline in a given cycle

ADD X14, X5, X6	LDUR X13, [X1,48]	ADD X12, X3, X4	SUB X11, X2, X3	LDUR X10, [X1,40]
Instruction fetch	Instruction decode	Execution	Memory	Write-back



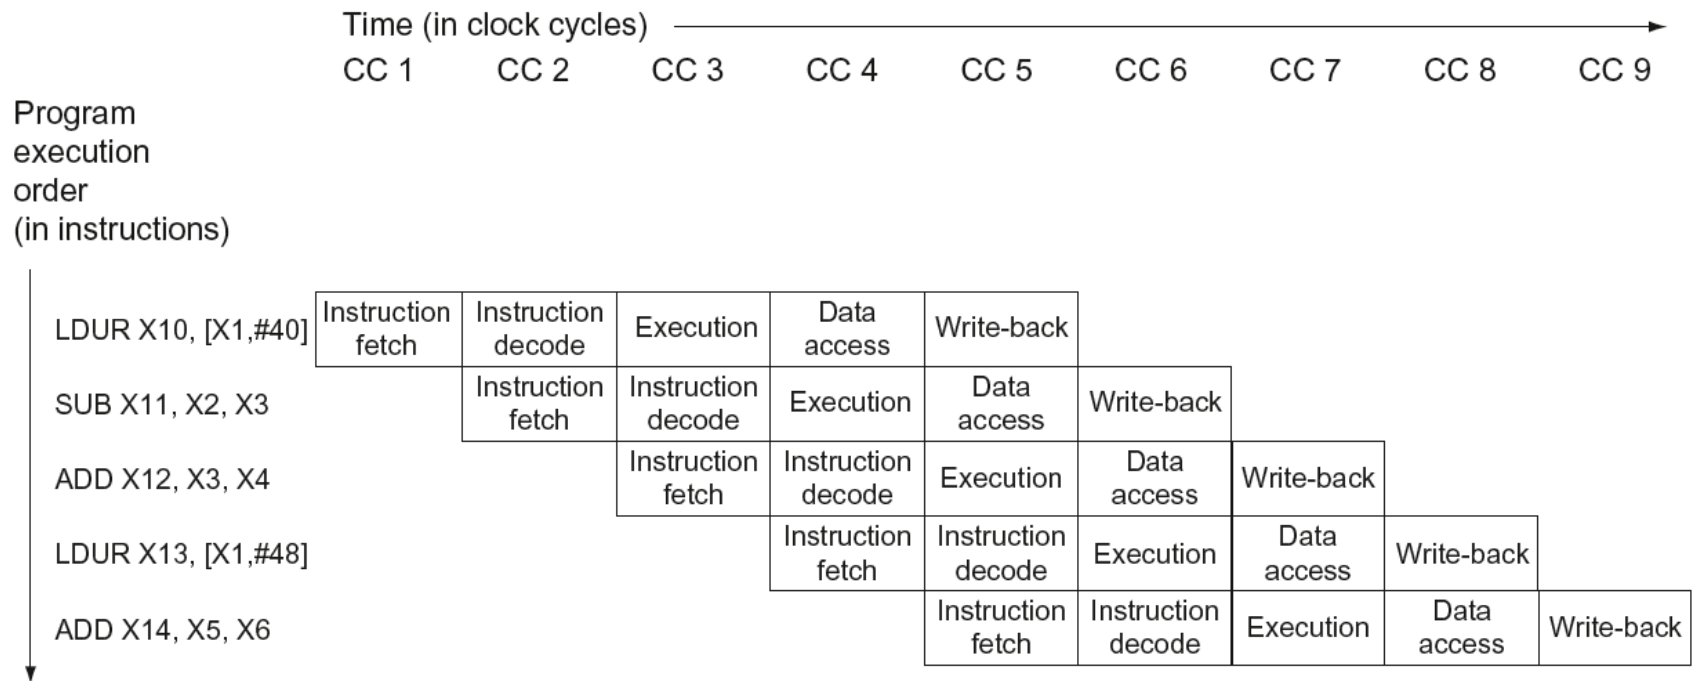
Multi-Cycle Pipeline Diagram

- Form showing resource usage

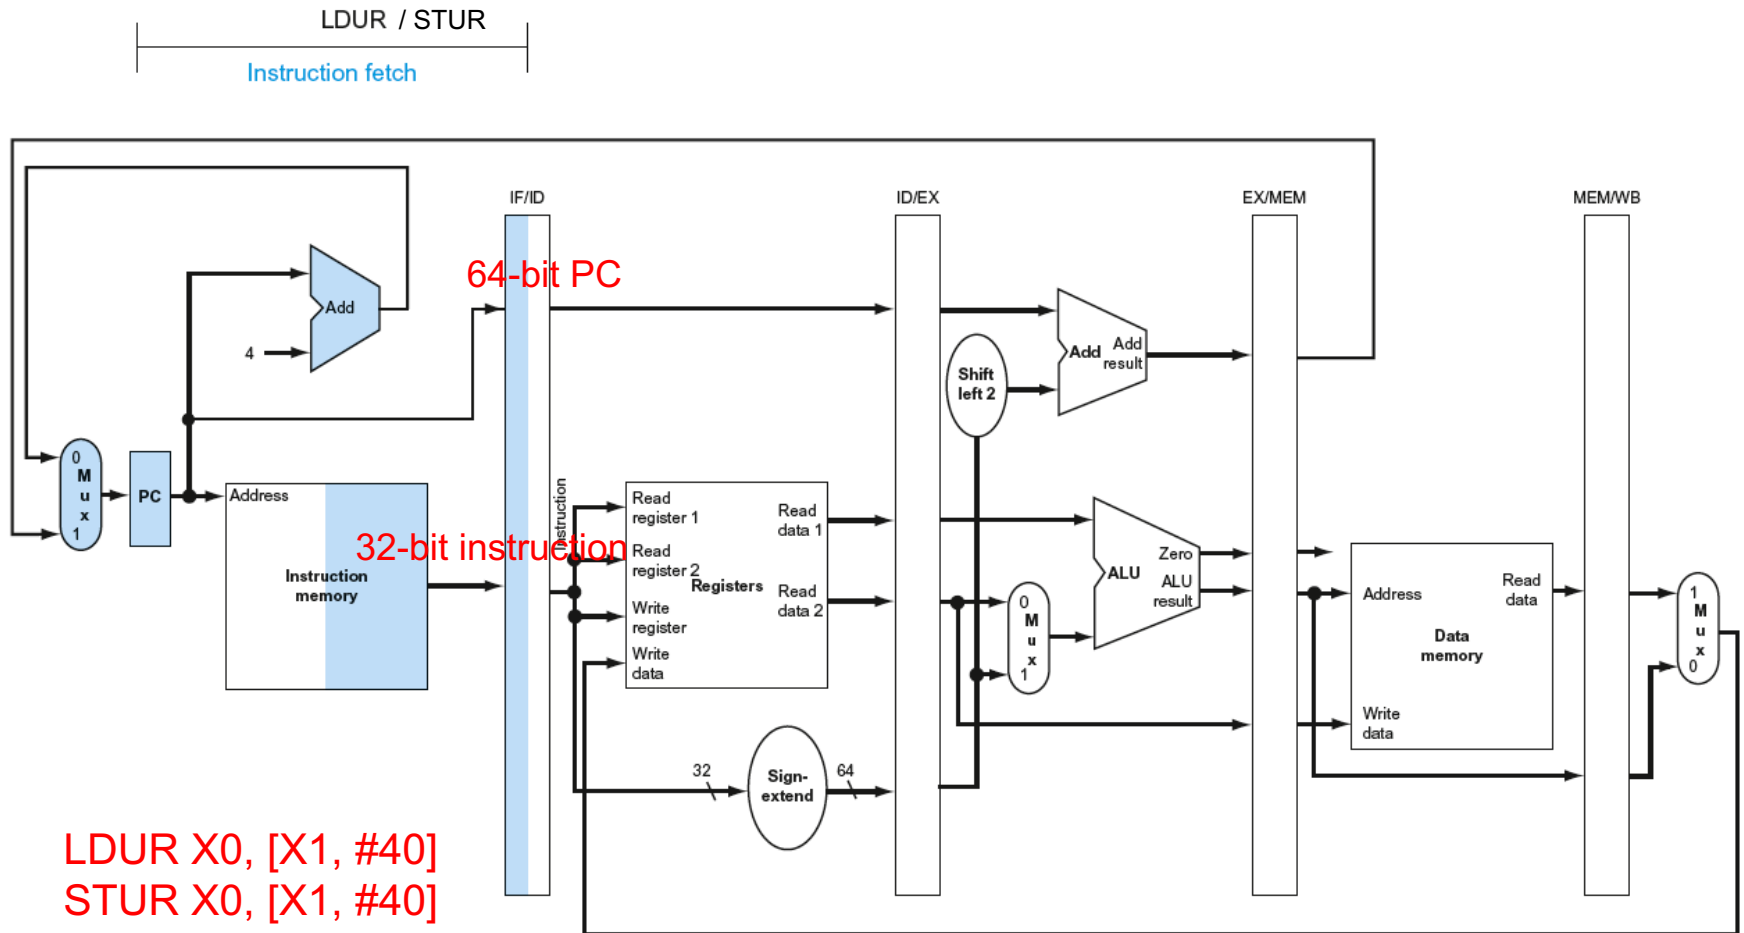


Multi-Cycle Pipeline Diagram

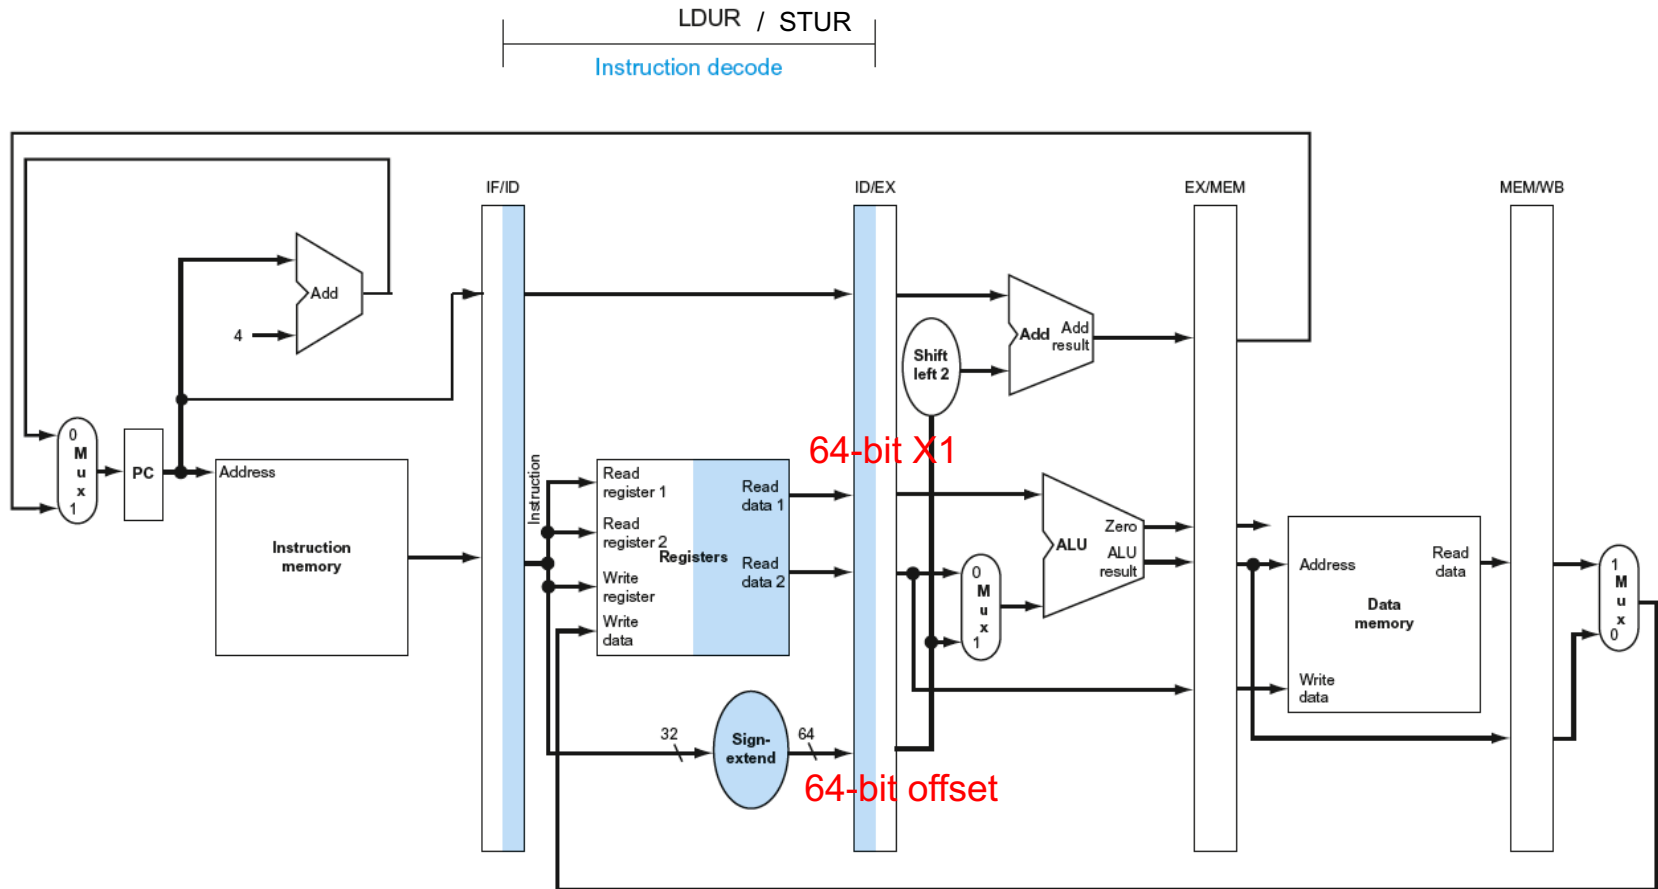
■ Traditional form



IF for Load, Store, ...

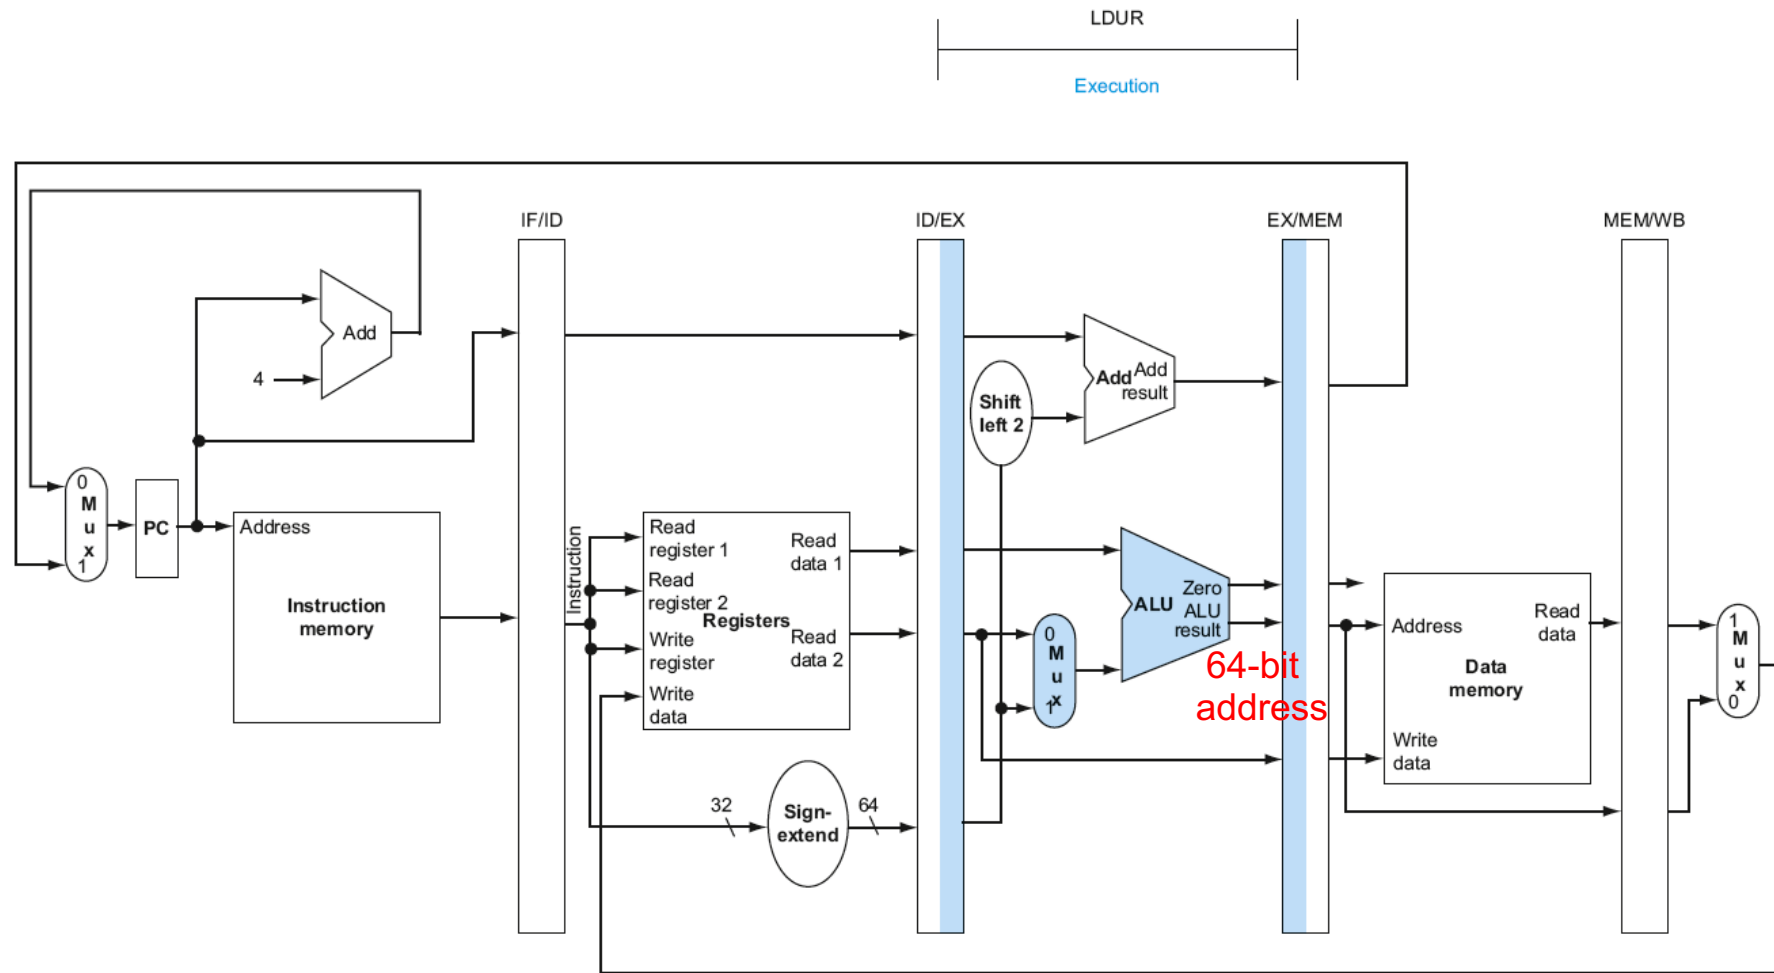


ID for Load, Store, ...



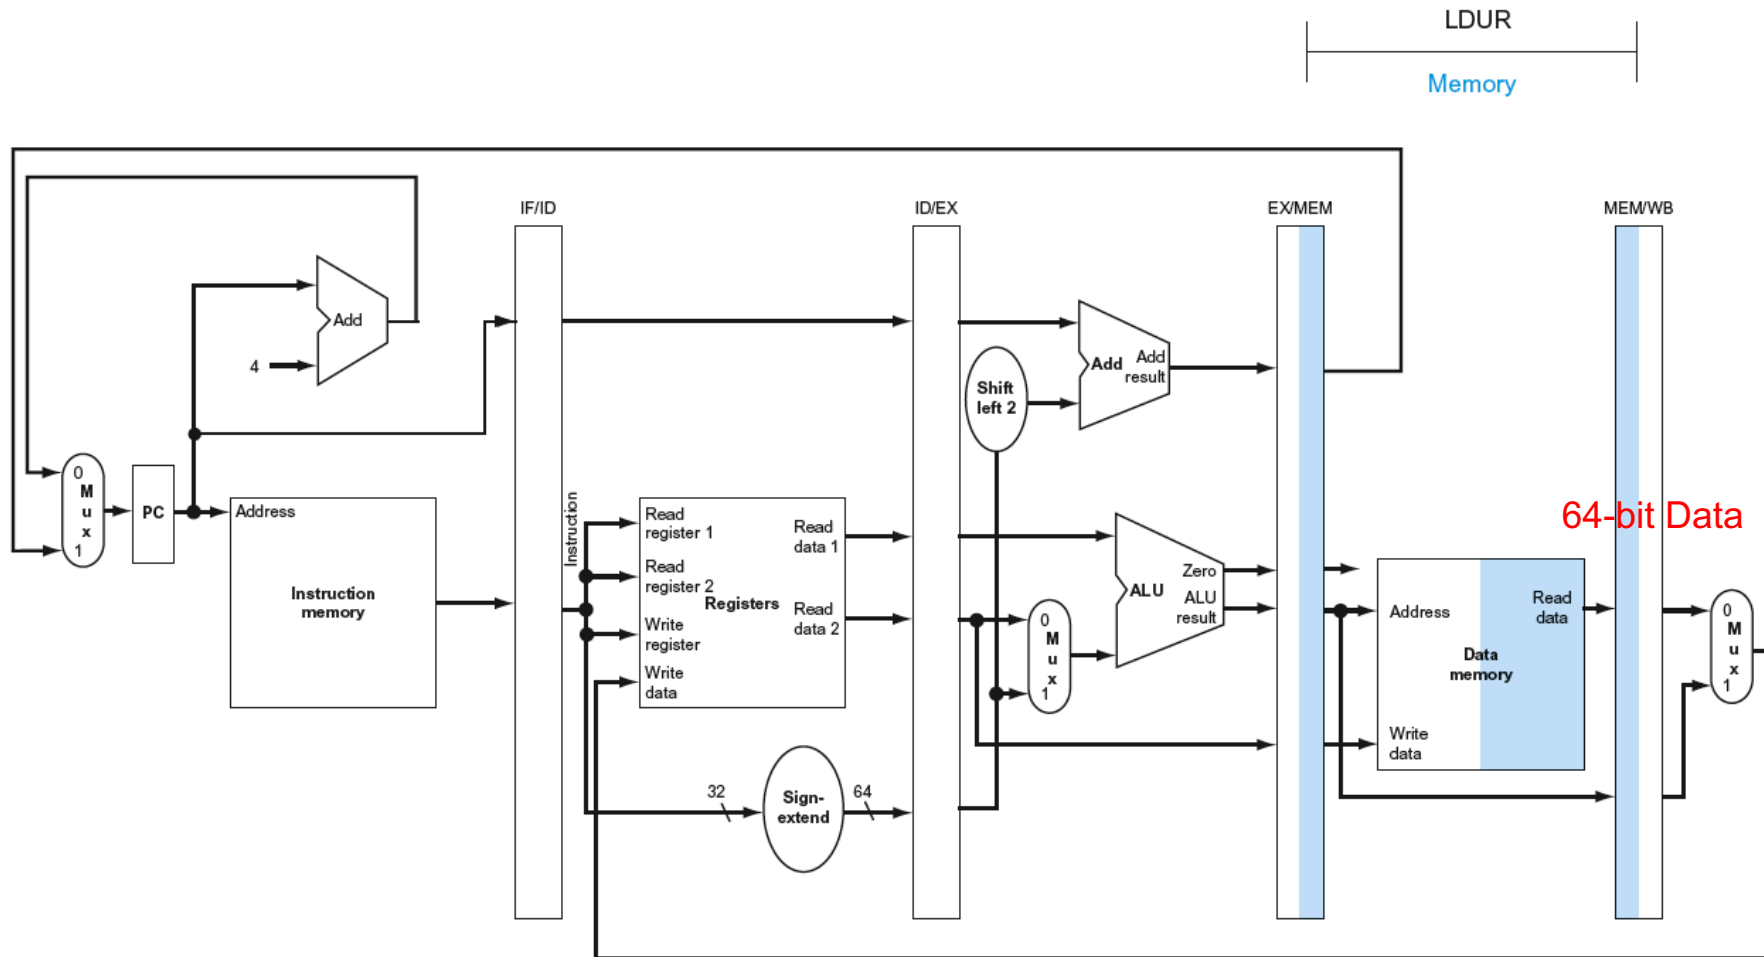
LDUR X0, [X1, #40]
STUR X0, [X1, #40]

EX for Load



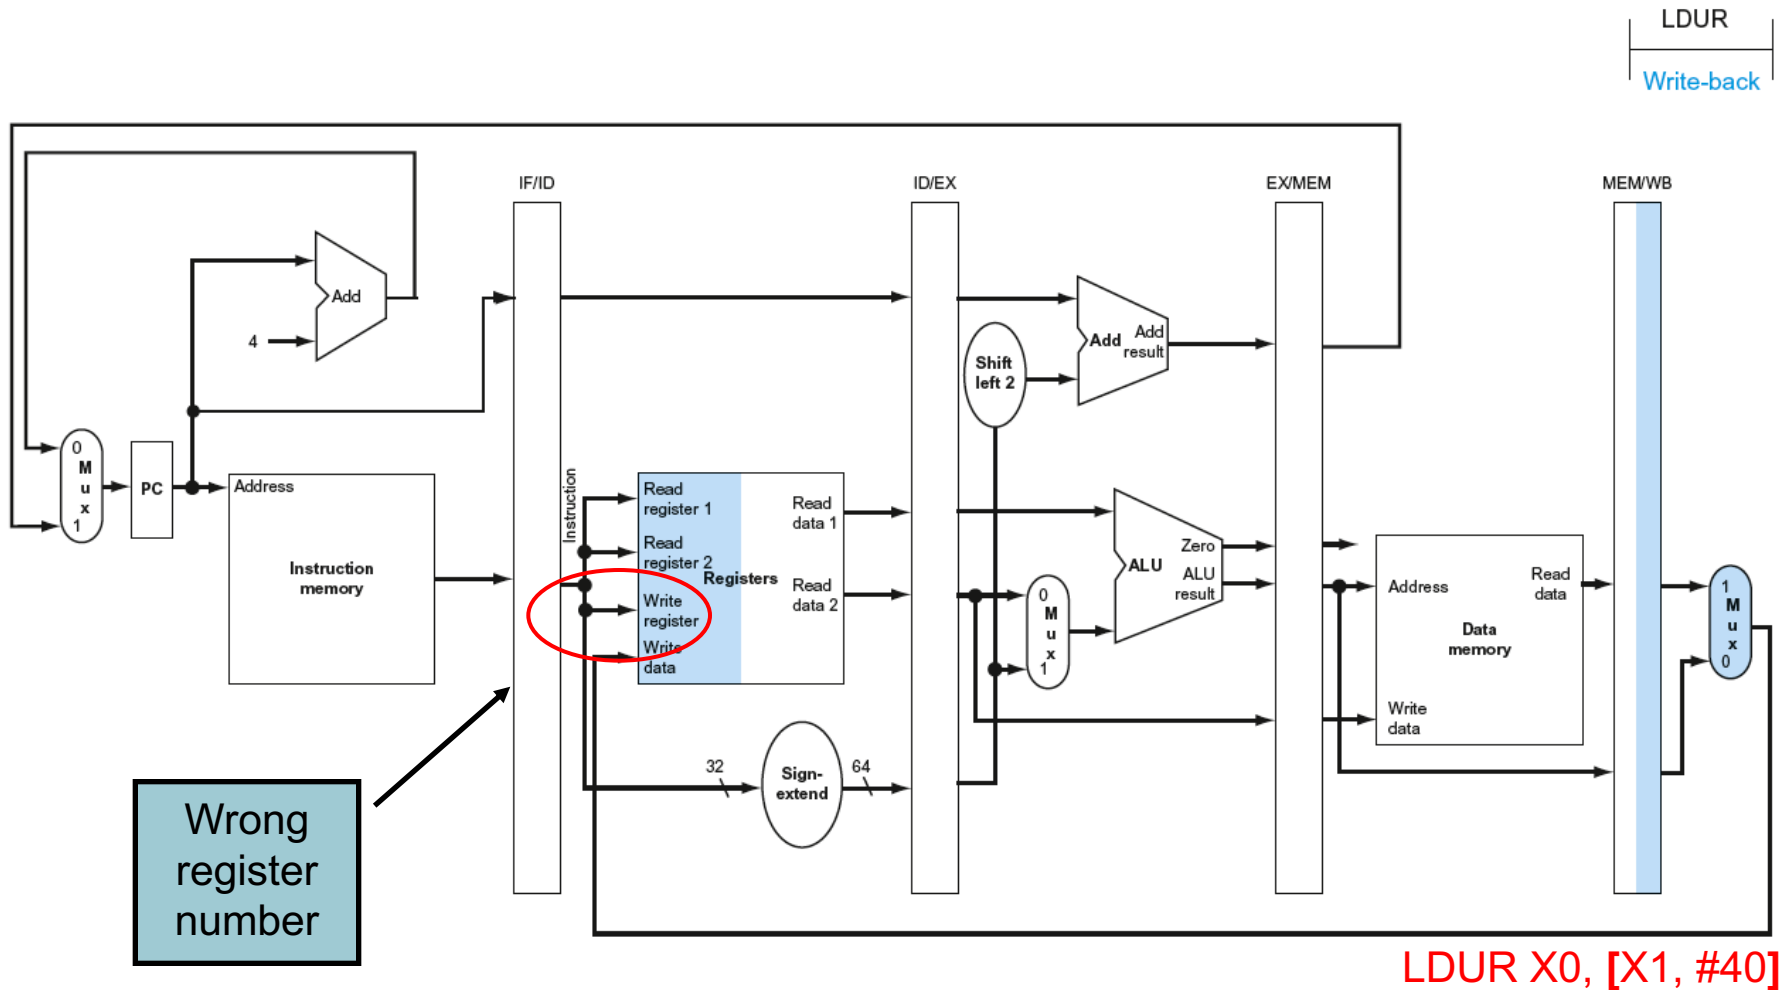
LDUR X0, [X1, #40]

MEM for Load

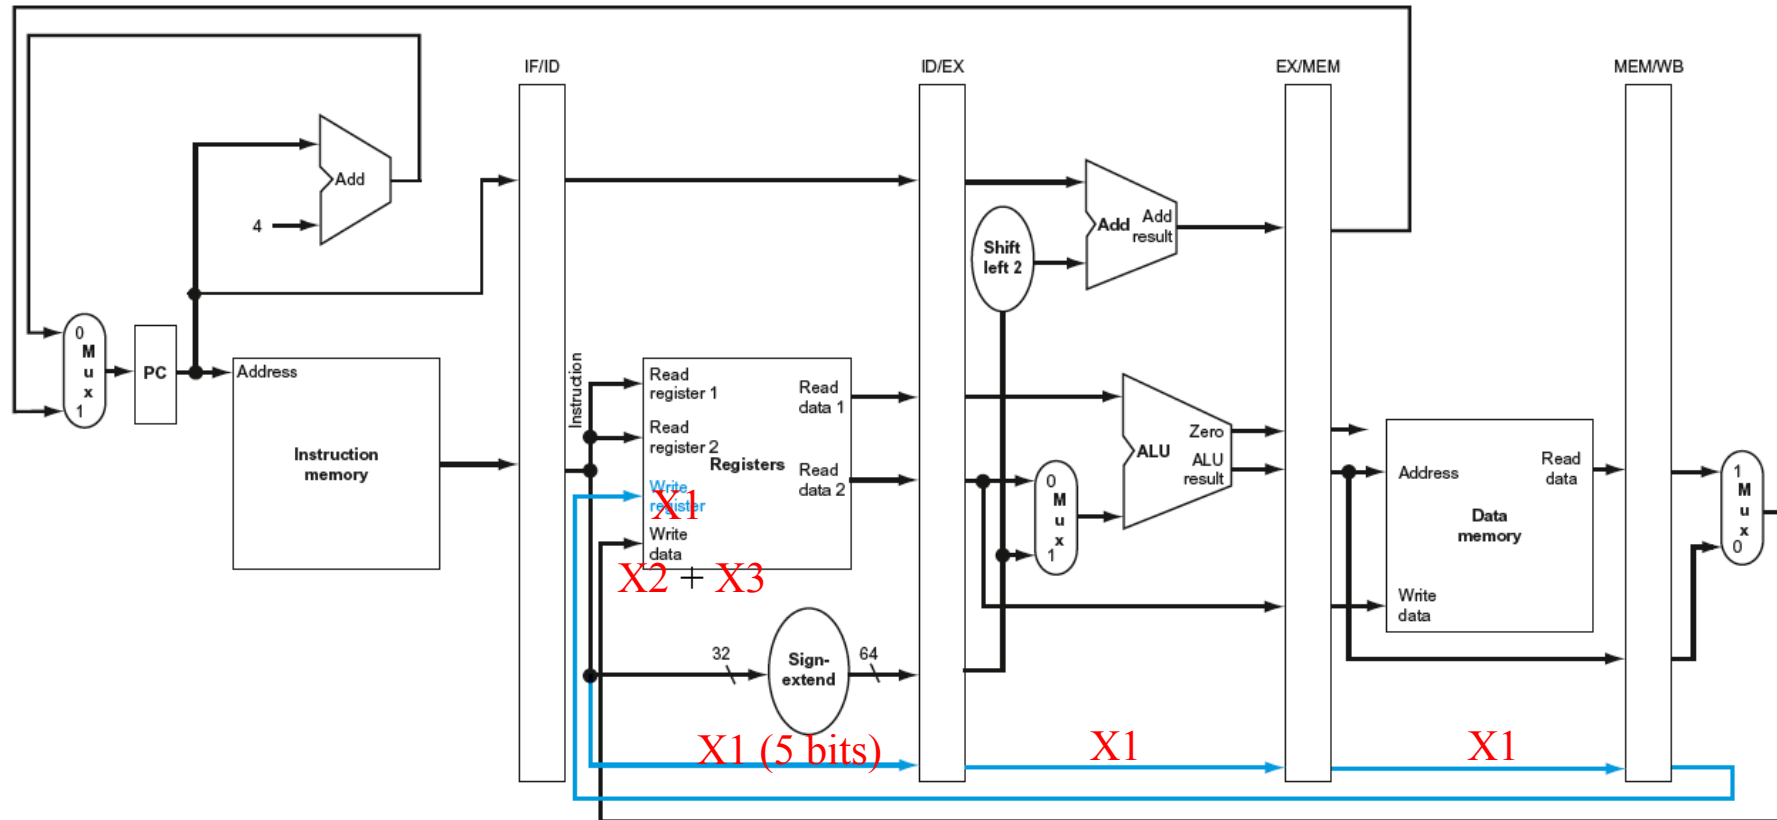


LDUR X0, [X1, #40]

WB for Load



Corrected Datapath for Load



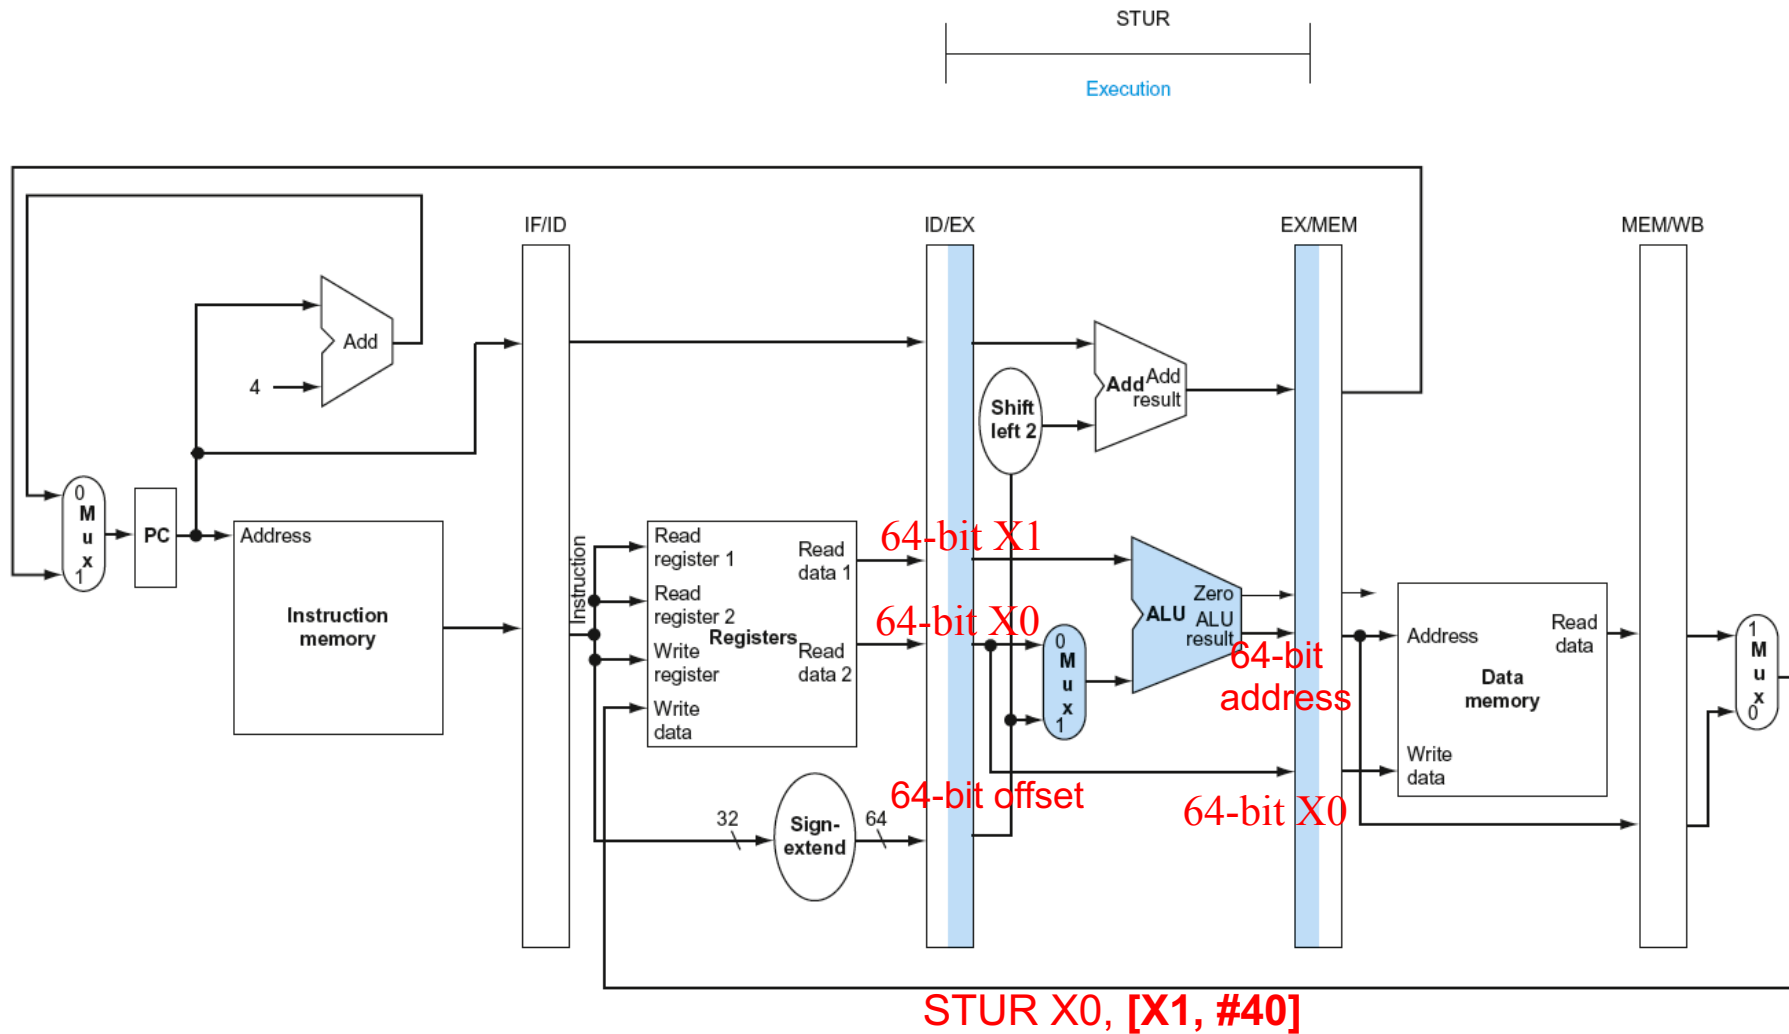
ADD X1, $X2$, $X3$

ADD $X2 + X3$

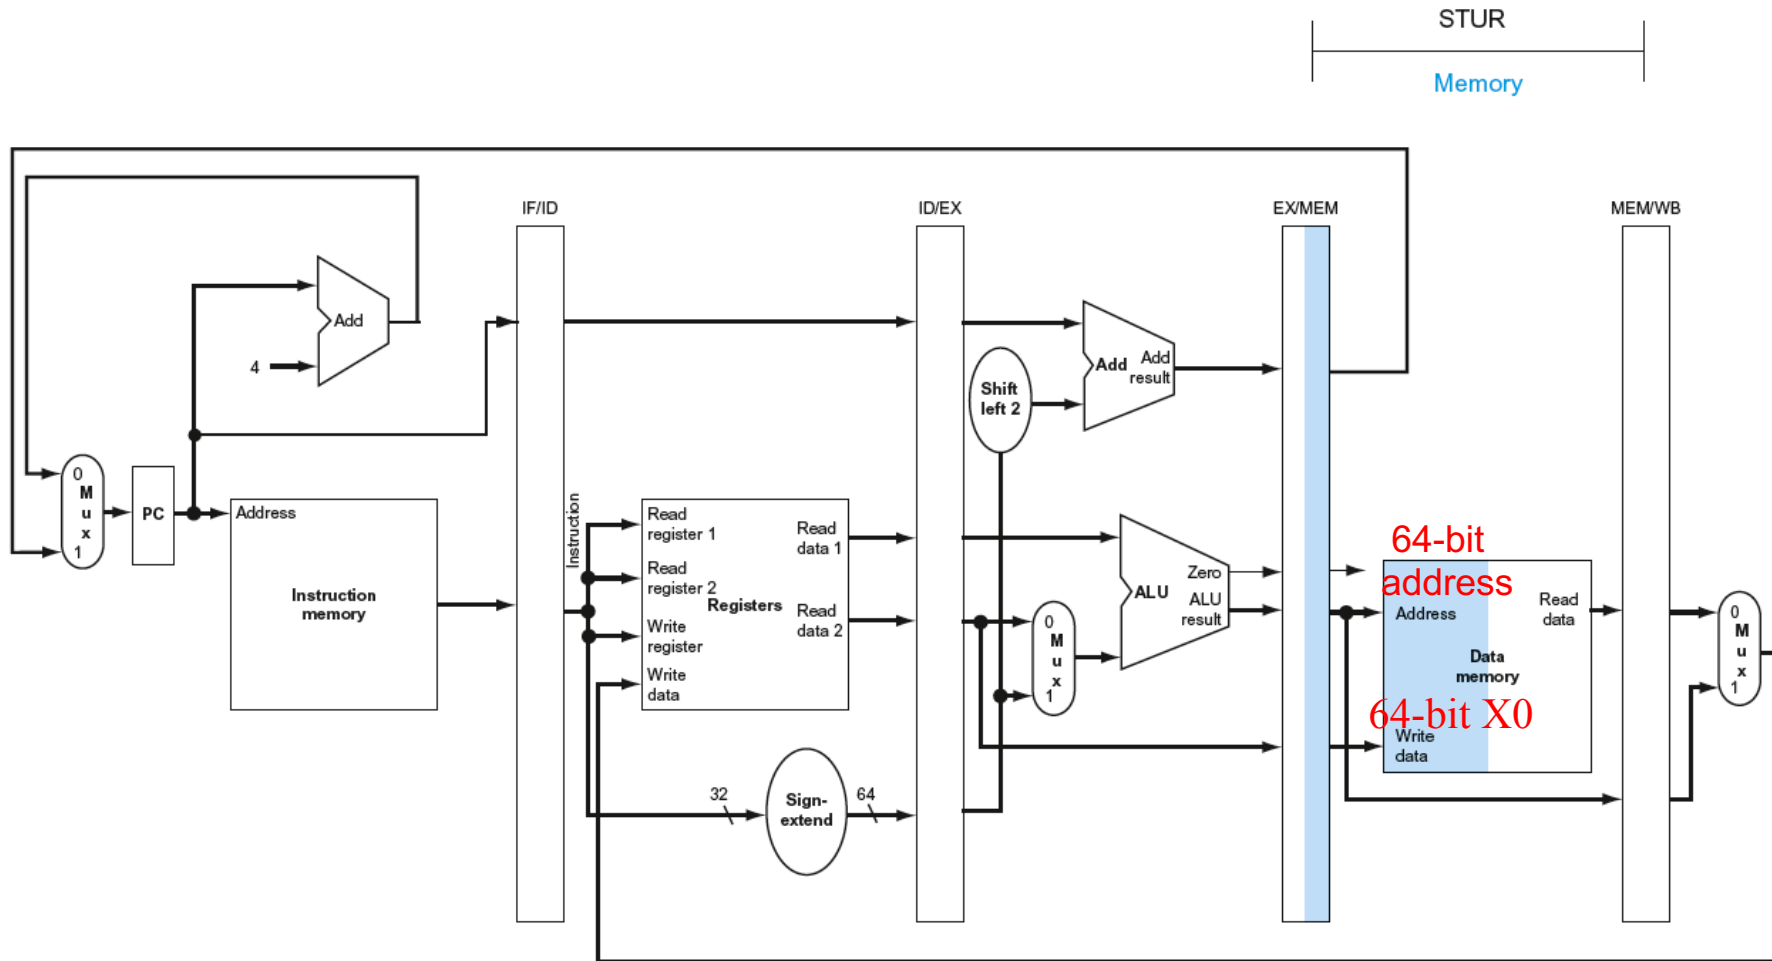
$X2 + X3$

ADD X4, $X5$, $X6$

EX for Store



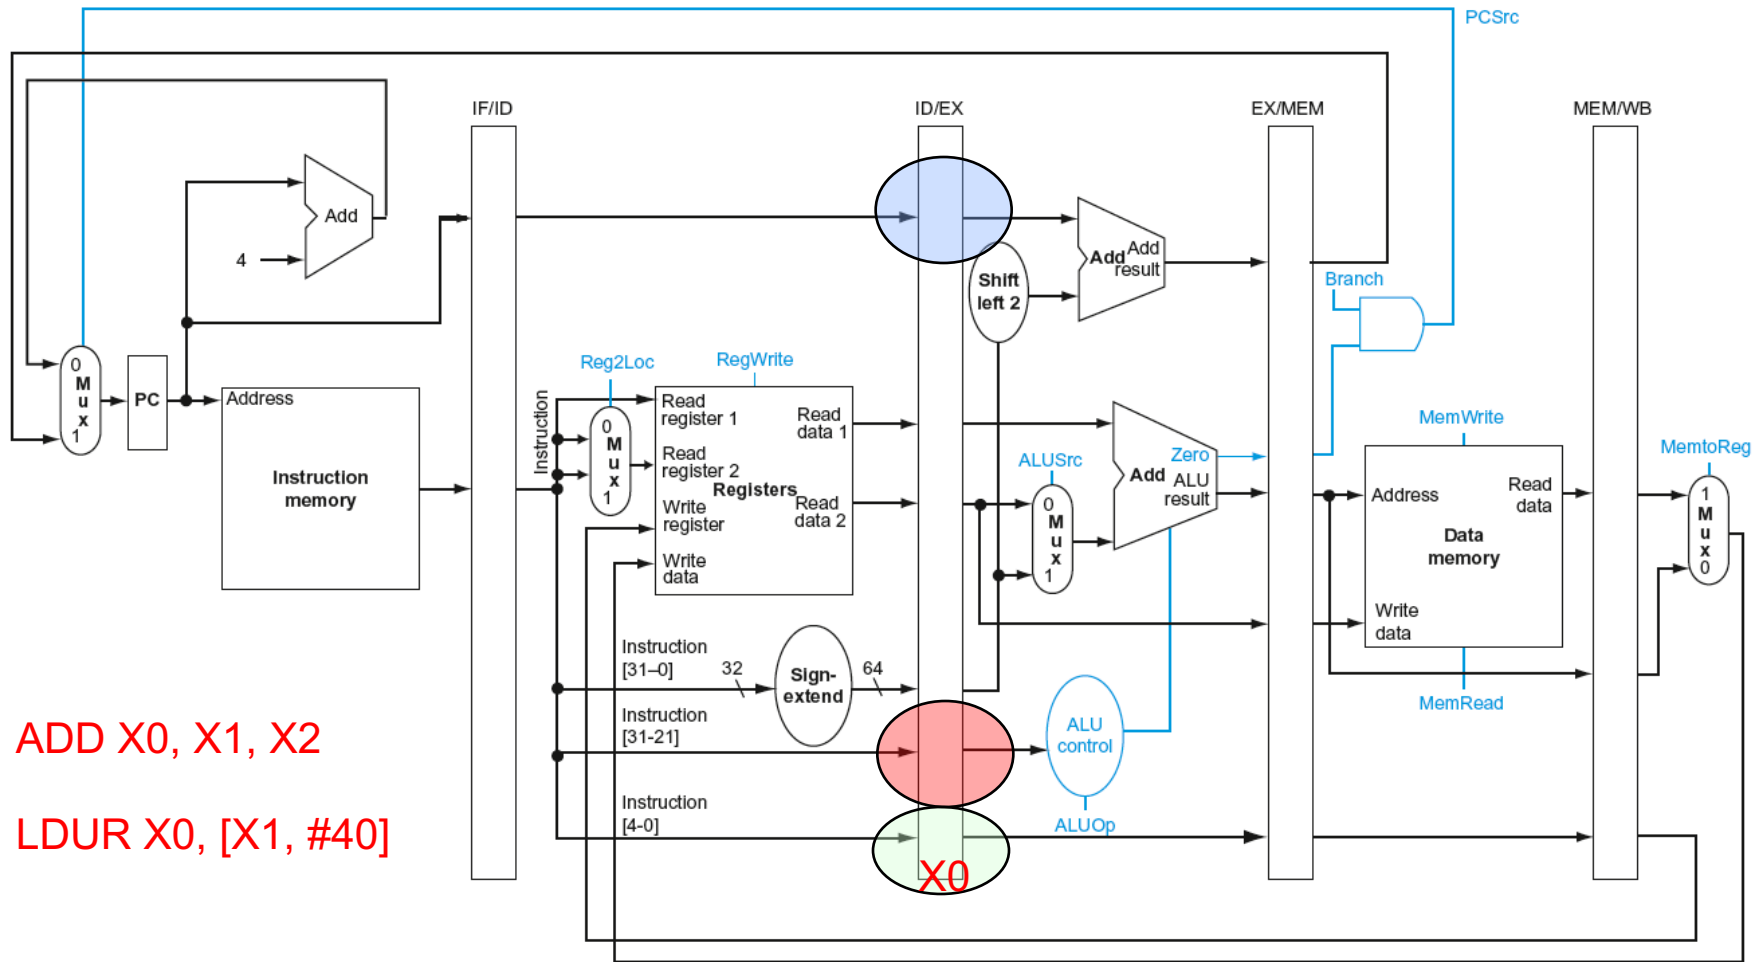
MEM for Store



STUR X0, [X1, #40]


MK

LDUR X0, [X1, #40]

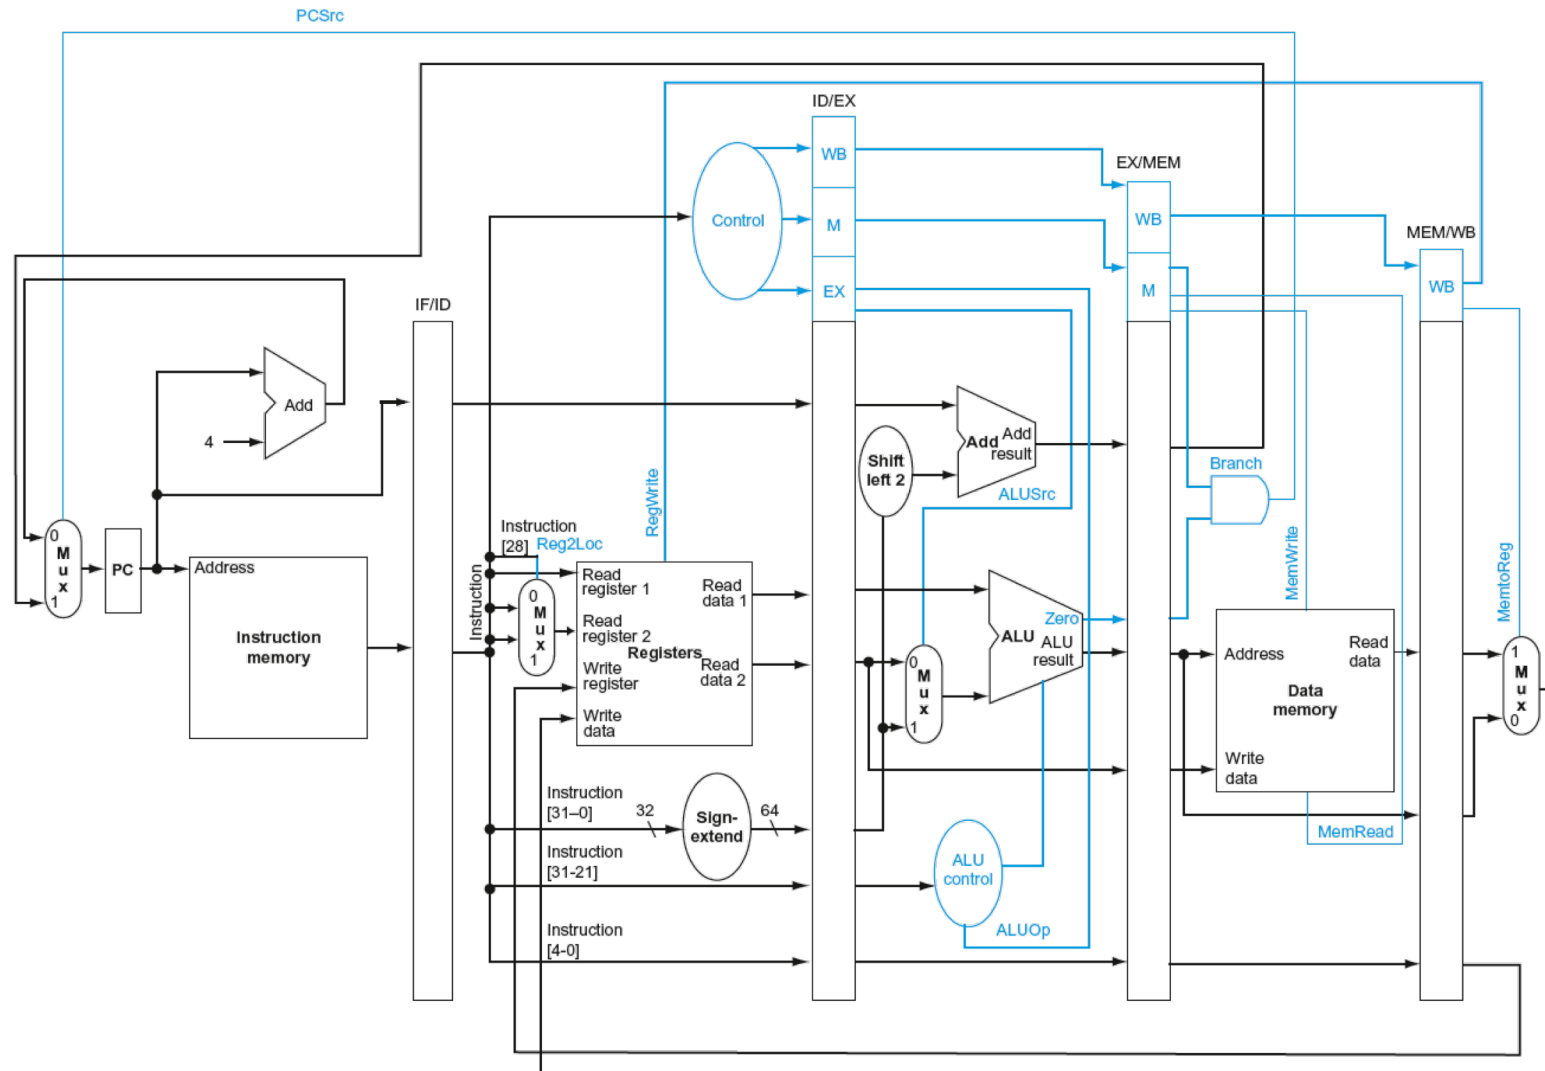


Pipelined Control

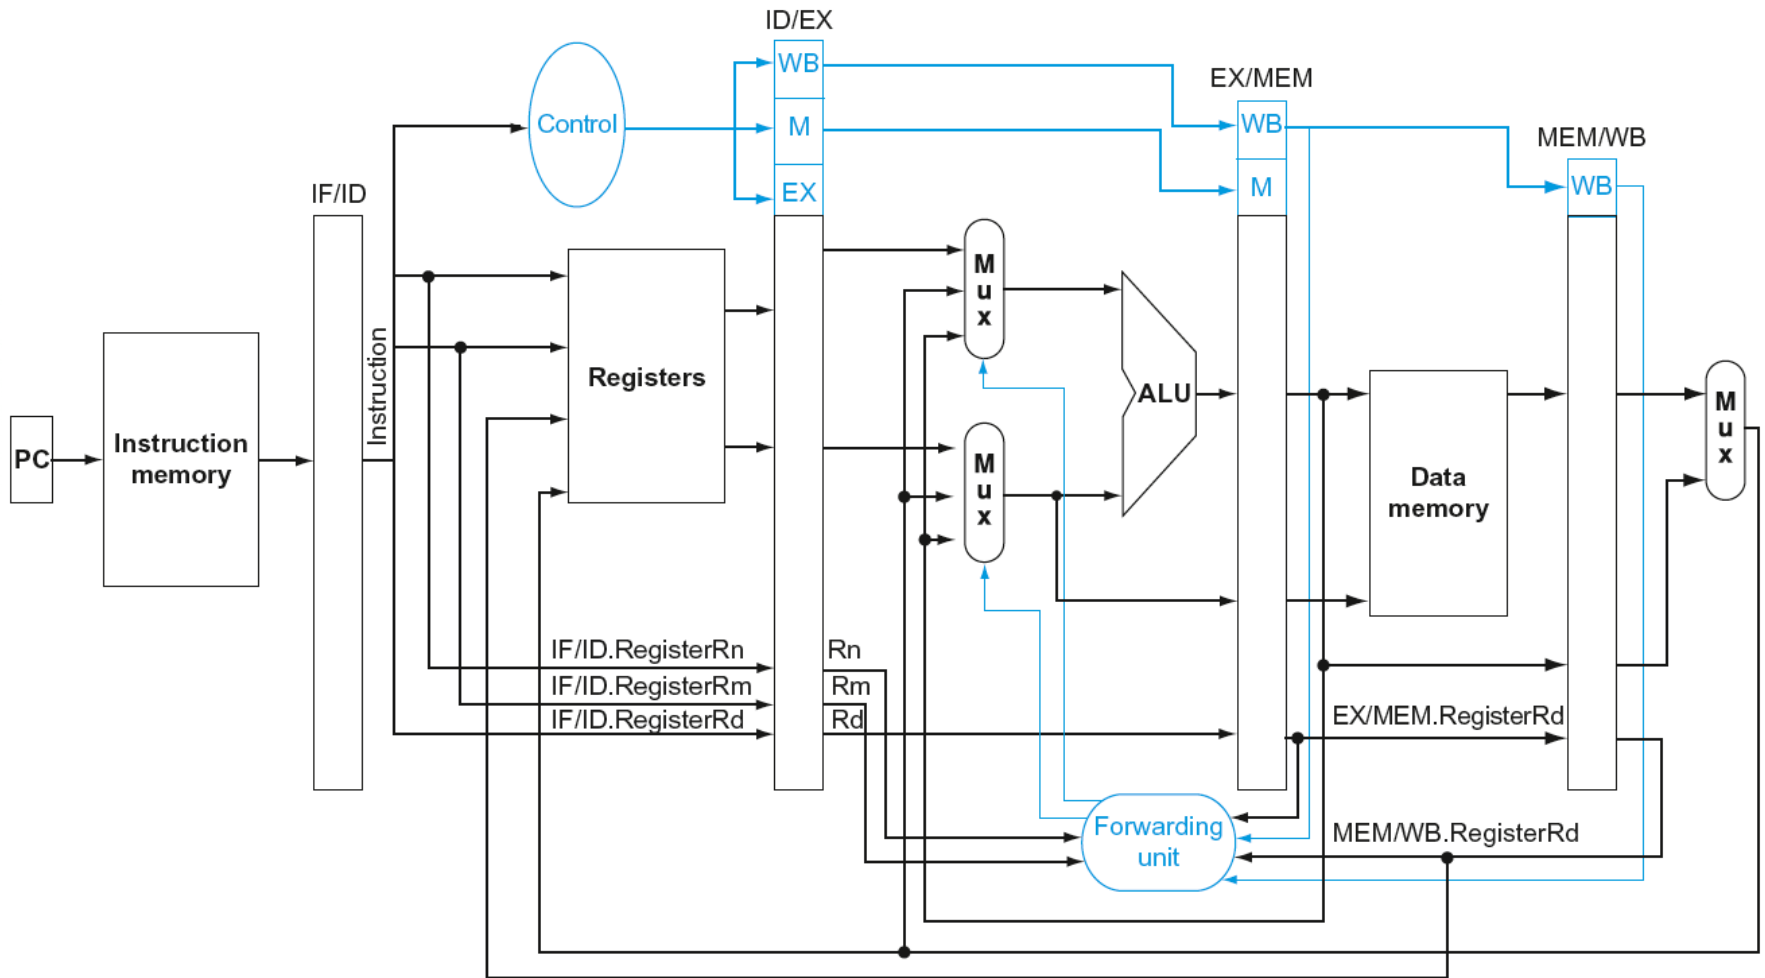
- Control signals derived from instruction
 - As in single-cycle implementation



Pipelined Control



Datapath with Forwarding



Fallacies

- Pipelining is easy (!)
 - The basic idea is easy
 - The devil is in the details
 - e.g., detecting data hazards
- Pipelining is independent of technology
 - So why haven't we always done pipelining?
 - More transistors make more advanced techniques feasible
 - Pipeline-related ISA design needs to take account of technology trends
 - e.g., predicated instructions

Pitfalls

- Poor ISA design can make pipelining harder
 - e.g., complex instruction sets (VAX, IA-32)
 - Significant overhead to make pipelining work
 - IA-32 micro-op approach
 - e.g., complex addressing modes
 - Register update side effects, memory indirection
 - e.g., delayed branches
 - Advanced pipelines have long delay slots

Concluding Remarks

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control