

PROBLEM 1 (2 points)

A fully connected neural network is used to classify a dataset that contains 10,000 color images of uppercase and lowercase letters (A-Z, a-z), where each image has 30x30x3 pixels. What are the sizes of the input layer and the output layer?

$$\text{Input layer} = 30 \times 30 \times 3 = \mathbf{2700}$$

$$\text{Output layer} = 26 + 26 = \mathbf{52}$$

PROBLEM 2 (4 points)

A labeled dataset has 5000 observations, 50 features and 20 classes. The feature matrix X is first processed by PCA. Three principal components are used to reconstruct the dataset \hat{X} . Combined with the class label, it is then used to train a Softmax Regression classifier.

What is the total number of *unused* principal components? **47**

What is the dimension of feature matrix X ? **5000x50**

What is the dimension of reconstructed matrix \hat{X} ? **5000x50**

What is the dimension of the matrix that contains *all* of the labels for training the Softmax Regression?

Possible answers depending on how you interpret the question:

- 5000x20 (20 one hot encoded labels)
- 5000x70 (50 features + 20 one-hot encoded labels)

PROBLEM 3 (8 points) – Choose one answer per question.

A). Which of the following is NOT true about Mercer's Theorem?

- (a) It can be used to find the mapping function ϕ
- (b) Computing the kernel is sufficient
- (c) It is not necessary to find the dimensionality of ϕ
- (d) Both (b) and (c)
- (e) None of the above

B). Which of the following is NOT true for SVM?

- (a) Training examples that are support vectors can be discarded
- (b) Training examples that are support vectors cannot be discarded
- (c) The decision boundary is placed in the middle of the "street"
- (d) The decision boundary is placed at the optimal location
- (e) None of the above


C). Your friend is training a classifier using a dataset whose underlying probability distribution is unknown. What should they use?

- (a) Linear Regression
- (b) Logistic Regression
- (c) Bayes Decision Theoretic
- (d) K-means
- (e) None of the above

D). How many eigenvalues and eigenvectors does matrix $X = \begin{bmatrix} 1 & 4 & 7 & 10 \\ 2 & 5 & 8 & 11 \\ 3 & 6 & 9 & 12 \end{bmatrix}$ have?

- (a) 3 eigenvalues, 4 eigenvectors
- (b) 4 eigenvalues, 3 eigenvectors
- (c) 3 eigenvalues, 3 eigenvectors
- (d) 4 eigenvalues, 4 eigenvectors
- (e) None of the above

E). What is the most likely AUC for the following ROC curve?



- (a) 0.95
- (b) 0.70
- (c) 0.30
- (d) 0.15
- (e) 0.05

F). Deep neural networks can have a problem where the gradients drop dramatically during back-propagation. What is this problem called?

- (a) Unknown gradients
- (b) Exploding gradients
- (c) Vanishing gradients
- (d) Descending gradients
- (e) None of the above

G). Which of the following is true about GANs?

- (a) During training the weights of the discriminator and the generator are optimized at the same time
- (b) During training the weights of the discriminator and the generator are “frozen” at the same time
- (c) The generator takes fake images as input and produces real images
- (d) GAN-generated images come from the output of the discriminator
- (e) None of the above

H). Considerable degradation in statistical significance caused by data sparsity in the hyperspace is not uncommon in high-dimensional data. What is this phenomenon known as?

- (a) Hebb’s phenomenon
- (b) Curse of the pharaohs
- (c) Curse of sparsity
- (d) Curse of dimensionality
- (e) None of the above

PROBLEM 4 (2 points)

List all the training examples generated from the sentence *Mary had two little lambs* by Word2Vec if the center word is *little* and the window size is 3 words.

{little, Mary}

{little, had}

{little, two}

{little, lambs}

PROBLEM 5 (3 points)

SVD is a powerful decomposition technique. It decomposes matrix A into the product of three matrices: $A = U.S.V^T$ where U is the left singular vectors of A (the eigenvectors of AA^T), S is the singular values of A (the square root of the eigenvalues of AA^T) and V is the right singular vectors of A (the eigenvectors of $A^T A$). If $A = \begin{bmatrix} 1 & 2 & 3 \\ 3 & 2 & 1 \end{bmatrix}$, determine its singular values.

$$\begin{bmatrix} 1 & 2 & 3 \\ 3 & 2 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 3 \\ 2 & 2 \\ 3 & 1 \end{bmatrix} = \begin{bmatrix} 14 & 10 \\ 10 & 14 \end{bmatrix} \rightarrow \begin{vmatrix} 14 - \lambda & 10 \\ 10 & 14 - \lambda \end{vmatrix} = 0$$

$$\rightarrow (14 - \lambda)(14 - \lambda) - 100 = 0 \rightarrow (14 - \lambda) = \pm 10$$

$$\rightarrow \lambda_1 = 24, \lambda_2 = 4$$

Singular values: $\sqrt{24} = 4.8990$ and $\sqrt{4} = 2$

PROBLEM 6 (3 points)

Suppose the principal components and the eigenvectors in PCA are given by matrix P and matrix V , respectively. Assume the mean of the dataset is 0.

$$P = \begin{bmatrix} 3 & 2 & 3 \\ 4 & 3 & 5 \\ 1 & 1 & 7 \end{bmatrix} \quad V = \begin{bmatrix} 2 & 0 & 3 \\ 5 & 1 & 2 \\ 4 & 4 & 0 \end{bmatrix}$$

What is the reconstructed dataset \hat{X} if only the first principal component is used?

$$\hat{X} = P.V^T$$

Where P is 3x1 and V is 3x1.

$$\hat{X} = \begin{bmatrix} 3 \\ 4 \\ 1 \end{bmatrix} \times [2 \quad 5 \quad 4] = \begin{bmatrix} 6 & 15 & 12 \\ 8 & 20 & 16 \\ 2 & 5 & 4 \end{bmatrix}$$

PROBLEM 7 (3 points)

A two-dimensional dataset is used to train an SVM classifier with a polynomial kernel $(\mathbf{x}^T \mathbf{z} + 1)^2$.

i	\mathbf{x}^T	y	Lagrange Multiplier	Support Vector
1	1, 0	-1	0.1	Yes
2	2, 4	-1	0.2	Yes
3	1, 1	-1	0.3	No
4	3, 1	+1	0.4	Yes
5	1, 3	+1	0.5	No
6	4, 4	+1	0.6	No

Determine the class label for the instance $(-1, 1)$. Assume the bias is 0.

Use only the support vectors: $M = 3$

$$h = \sum_{i=1}^M \alpha_i y_i K(\mathbf{x}_i^T \mathbf{z}) = \sum_{i=1}^M \alpha_i y_i (\mathbf{x}_i^T \mathbf{z} + 1)^2$$

For $\mathbf{z}^T = -1, 1$: $h = -(0.1)(-1 + 1)^2 - (0.2)(2 + 1)^2 + (0.4)(-2 + 1)^2 = -1.4 < 0$

→ **negative class**