

# COEN 241: HW - 1

## QEMU Installation:

1. Download Ubuntu ISO ([Download from here](#))
2. Install QEMU:  
> brew install qemu (stable version for Apple M1 chip: 6.2.0)
3. Create QEMU image:  
> qemu-img create ubuntu.img 20G -f qcow2
4. Install VM:  
> qemu-system-aarch64 \-accel hvf \-cpu cortex-a57 \-M virt,highmem=off  
  \-m 2G \-smp 2 \-drive  
    file=/opt/homebrew/Cellar/qemu/6.2.0/share/qemu/edk2-aarch64-code.fd,if  
    =pflash,format=raw,readonly=on \-drive  
    if=none,file=ubuntu.img,format=qcow2,id=hd0 \-device  
    virtio-blk-device,drive=hd0,serial="trial\_2" \-device  
    virtio-net-device,netdev=net0 \-netdev user,id=net0 \-vga none \-device  
    ramfb \-cdrom ubuntu-20.04.5-live-server-arm64.iso \-device usb-ehci  
    \-device usb-kbd \-device usb-mouse \-usb \-nographic
5. Run Image:  
> qemu-system-aarch64 \-accel hvf \-cpu cortex-a57 \-M virt,highmem=off  
  \-m 2G \-smp 2 \-drive  
    file=/opt/homebrew/Cellar/qemu/6.2.0/share/qemu/edk2-aarch64-code.fd,if  
    =pflash,format=raw,readonly=on \-drive  
    if=none,file=ubuntu.img,format=qcow2,id=hd0 \-device  
    virtio-blk-device,drive=hd0,serial="trial\_2" \-device  
    virtio-net-device,netdev=net0 \-netdev user,id=net0 \-vga none \-device  
    ramfb \-device usb-ehci \-device usb-kbd \-device usb-mouse \-usb  
    \-nographic

## Experimental Setup:

System: 2021 Macbook Air; Chip: Apple M1, Memory: 8GB RAM

Testing 3 configurations for both QEMU and Docker:

- 2 cores with 2GB memory allocation
- 4 cores with 4GB memory allocation
- 6 cores with 6GB memory allocation

## Docker Installation:

Rosetta is required for installation of Docker in order to ensure instruction set compatibility.

1. Run Rosetta:

```
> softwareupdate --install-rosetta
```

2. Install Docker Engine - Download the docker dmg image from [here](#).

3. Check and test the installation:

```
> docker run hello-world
```

### Output:

```
[benitarego@Benitas-MacBook-Air ~ % softwareupdate --install-rosetta
I have read and agree to the terms of the software license agreement. A list of
Apple SLAs may be found here: http://www.apple.com/legal/sla/
Type A and press return to agree: A
2023-01-30 17:47:16.292 softwareupdate[23379:879330] Package Authoring Error: 01
2-60306: Package reference com.apple.pkg.RosettaUpdateAuto is missing installKBy
tes attribute
Install of Rosetta 2 finished successfully
```

```
[benitarego@Benitas-MacBook-Air ~ % docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
7050e35b49f5: Pull complete
Digest: sha256:aa0cc8055b82dc2509bed2e19b275c8f463506616377219d9642221ab53cf9fe
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker!  
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.  
(arm64v8)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:  
\$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:  
<https://hub.docker.com/>

For more examples and ideas, visit:  
<https://docs.docker.com/get-started/>

## **Experimental Results:**

### **Configuration 1: 2 cores with 2GB memory allocation**

Tests:

1. Max-prime=2000, time=30 seconds
2. Max-prime=20000, time=30 seconds
3. Max-prime=200000, time=30 seconds

For Docker, we are using the zyclonite/sysbench image which comes with sysbench preinstalled.

To install sysbench:

```
> sudo apt install sysbench
```

Sysbench command:

```
> sysbench cpu --cpu-max-prime={some_value} --num-threads={some_value}  
--time={some_value} run
```

#### ***Case 1: Max-prime = 2000***

For QEMU:

```
> sysbench cpu --cpu-max-prime=2000 --num-threads=1 --time=30 run
```

OR

```
> source ~/QEMU_2000.sh
```

For Docker:

```
> docker run --rm zyclonite/sysbench --test(cpu) --cpu-max-prime=2000  
--time=30 run
```

OR

```
> source ~/DOCKER_2000.sh
```

### *QEMU Results:*

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 71182.82  
  
General statistics:  
    total time:          30.0005s  
    total number of events: 2135617  
  
Latency (ms):  
    min:                0.01  
    avg:                0.01  
    max:                2.28  
    95th percentile:    0.01  
    sum:               29720.01  
  
Threads fairness:  
    events (avg/stddev): 2135617.0000/0.00  
    execution time (avg/stddev): 29.7200/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 71130.71  
  
General statistics:  
    total time:          30.0004s  
    total number of events: 2134058  
  
Latency (ms):  
    min:                0.01  
    avg:                0.01  
    max:                3.44  
    95th percentile:    0.01  
    sum:               29717.67  
  
Threads fairness:  
    events (avg/stddev): 2134058.0000/0.00  
    execution time (avg/stddev): 29.7177/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 71465.38  
  
General statistics:  
    total time: 30.0003s  
    total number of events: 2144084  
  
Latency (ms):  
    min: 0.01  
    avg: 0.01  
    max: 5.68  
    95th percentile: 0.01  
    sum: 29710.15  
  
Threads fairness:  
    events (avg/stddev): 2144084.0000/0.00  
    execution time (avg/stddev): 29.7101/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 71295.74  
  
General statistics:  
    total time: 30.0002s  
    total number of events: 2138967  
  
Latency (ms):  
    min: 0.01  
    avg: 0.01  
    max: 3.29  
    95th percentile: 0.01  
    sum: 29693.98  
  
Threads fairness:  
    events (avg/stddev): 2138967.0000/0.00  
    execution time (avg/stddev): 29.6940/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 70708.97  
  
General statistics:  
    total time: 30.0010s  
    total number of events: 2121494  
  
Latency (ms):  
    min: 0.01  
    avg: 0.01  
    max: 12.53  
    95th percentile: 0.01  
    sum: 29706.58  
  
Threads fairness:  
    events (avg/stddev): 2121494.0000/0.00  
    execution time (avg/stddev): 29.7066/0.00
```

No. of Iterations	Events/sec
1	71182.82
2	71130.71
3	71465.38 - max value
4	71295.74
5	70708.97 - min value
Avg:	71156.72

## Docker Results:

```
benitarego@Benitas-MacBook-Air ~ % docker run --rm zyclonite/sysbench --test=cpu  
--cpu-max-prime=2000 --time=30 run  
  
WARNING: the --test option is deprecated. You can pass a script name or path on t  
he command line without any options.  
sysbench 1.0.20-f6f6117dc4 (using bundled LuaJIT 2.1.0-beta2)  
  
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 34105.65  
  
General statistics:  
    total time:          30.0002s  
    total number of events: 1023210  
  
Latency (ms):  
    min:                 0.03  
    avg:                 0.03  
    max:                11.76  
    95th percentile:     0.03  
    sum:                29837.75  
  
Threads fairness:  
    events (avg/stddev): 1023210.0000/0.00  
    execution time (avg/stddev): 29.8377/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 34182.79  
  
General statistics:  
    total time:          30.0001s  
    total number of events: 1025526  
  
Latency (ms):  
    min:                 0.03  
    avg:                 0.03  
    max:                1.75  
    95th percentile:     0.03  
    sum:                29850.79  
  
Threads fairness:  
    events (avg/stddev): 1025526.0000/0.00  
    execution time (avg/stddev): 29.8508/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 34227.73  
  
General statistics:  
    total time: 30.0002s  
    total number of events: 1026882  
  
Latency (ms):  
    min: 0.03  
    avg: 0.03  
    max: 0.94  
    95th percentile: 0.03  
    sum: 29847.86  
  
Threads fairness:  
    events (avg/stddev): 1026882.0000/0.00  
    execution time (avg/stddev): 29.8479/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 34201.29  
  
General statistics:  
    total time: 30.0003s  
    total number of events: 1026082  
  
Latency (ms):  
    min: 0.03  
    avg: 0.03  
    max: 2.54  
    95th percentile: 0.03  
    sum: 29803.42  
  
Threads fairness:  
    events (avg/stddev): 1026082.0000/0.00  
    execution time (avg/stddev): 29.8034/0.00
```

```

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 2000

Initializing worker threads...

Threads started!

CPU speed:
  events per second: 34194.07

General statistics:
  total time:          30.0002s
  total number of events: 1025878

Latency (ms):
  min:                0.03
  avg:                0.03
  max:               14.60
  95th percentile:    0.03
  sum:            29812.32

Threads fairness:
  events (avg/stddev):   1025878.0000/0.00
  execution time (avg/stddev): 29.8123/0.00

```

No. of Iterations	Events/sec
1	34105.65 - min value
2	34182.79
3	34227.73 - max value
4	34201.29
5	34194.07
Avg:	34182.31

Result: QEMU is faster than Docker in CPU performance.

### **Case 2: Max-prime = 20000**

For QEMU:

```
> sysbench cpu --cpu-max-prime=20000 --num-threads=1 --time=30 run
```

OR

```
> source ~/QEMU_20000.sh
```

For Docker:

```
> docker run --rm zyclonite/sysbench --test(cpu) --cpu-max-prime=20000
--time=30 run
```

OR

```
> source ~/DOCKER_20000.sh
```

*QEMU Results:*

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 20000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 2607.18  
  
General statistics:  
    total time:          30.0011s  
    total number of events: 78224  
  
Latency (ms):  
    min:                0.37  
    avg:                0.38  
    max:                13.37  
    95th percentile:    0.40  
    sum:               29966.70  
  
Threads fairness:  
    events (avg/stddev): 78224.0000/0.00  
    execution time (avg/stddev): 29.9667/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 20000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 2617.16  
  
General statistics:  
    total time:          30.0014s  
    total number of events: 78525  
  
Latency (ms):  
    min:                0.37  
    avg:                0.38  
    max:                9.23  
    95th percentile:    0.39  
    sum:               29919.85  
  
Threads fairness:  
    events (avg/stddev): 78525.0000/0.00  
    execution time (avg/stddev): 29.9199/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 20000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 2604.12  
  
General statistics:  
    total time: 30.0008s  
    total number of events: 78131  
  
Latency (ms):  
    min: 0.37  
    avg: 0.38  
    max: 11.34  
    95th percentile: 0.41  
    sum: 29908.75  
  
Threads fairness:  
    events (avg/stddev): 78131.0000/0.00  
    execution time (avg/stddev): 29.9087/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 20000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 2608.95  
  
General statistics:  
    total time: 30.0013s  
    total number of events: 78278  
  
Latency (ms):  
    min: 0.37  
    avg: 0.38  
    max: 12.70  
    95th percentile: 0.41  
    sum: 29936.29  
  
Threads fairness:  
    events (avg/stddev): 78278.0000/0.00  
    execution time (avg/stddev): 29.9363/0.00
```

```

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:
  events per second: 2612.34

General statistics:
  total time:          30.0007s
  total number of events: 78378

Latency (ms):
  min:                 0.37
  avg:                 0.38
  max:                 5.46
  95th percentile:    0.40
  sum:                29901.89

Threads fairness:
  events (avg/stddev): 78378.0000/0.00
  execution time (avg/stddev): 29.9019/0.00

```

No. of Iterations	Events/sec
1	2617.16 - max value
2	2607.18
3	2604.12 - min value
4	2608.95
5	2612.34
Avg:	2609.95

## Docker Results:

```
[benitarego@Benitas-MacBook-Air ~ % docker run --rm zyclonite/sysbench --test(cpu) --cpu-max-prime=20000 --time=30 run
WARNING: the --test option is deprecated. You can pass a script name or path on the command line without any options.
sysbench 1.0.20-f6f6117dc4 (using bundled LuaJIT 2.1.0-beta2)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 1986.20

General statistics:
total time: 30.0006s
total number of events: 59590

Latency (ms):
min: 0.49
avg: 0.50
max: 1.96
95th percentile: 0.51
sum: 29926.56

Threads fairness:
events (avg/stddev): 59590.0000/0.00
execution time (avg/stddev): 29.9266/0.00
```

```
Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 1966.42

General statistics:
total time: 30.0004s
total number of events: 58999

Latency (ms):
min: 0.49
avg: 0.51
max: 14.85
95th percentile: 0.54
sum: 29915.89

Threads fairness:
events (avg/stddev): 58999.0000/0.00
execution time (avg/stddev): 29.9159/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 20000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 1963.54  
  
General statistics:  
    total time: 30.0007s  
    total number of events: 58913  
  
Latency (ms):  
    min: 0.49  
    avg: 0.51  
    max: 10.40  
    95th percentile: 0.54  
    sum: 29893.68  
  
Threads fairness:  
    events (avg/stddev): 58913.0000/0.00  
    execution time (avg/stddev): 29.8937/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 20000  
  
Initializing worker threads...  
  
Threads started!  
  
WARNING: the --test option is deprecated. You can pass a script  
he command line without any options.  
CPU speed:  
    events per second: 1968.02  
  
General statistics:  
    total time: 30.0006s  
    total number of events: 59047  
  
Latency (ms):  
    min: 0.49  
    avg: 0.51  
    max: 5.49  
    95th percentile: 0.54  
    sum: 29892.35  
  
Threads fairness:  
    events (avg/stddev): 59047.0000/0.00  
    execution time (avg/stddev): 29.8923/0.00
```

```

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:
    events per second: 1982.38

General statistics:
    total time:          30.0006s
    total number of events: 59475

Latency (ms):
    min:                 0.49
    avg:                 0.50
    max:                 1.39
    95th percentile:     0.53
    sum:                29936.80

Threads fairness:
    events (avg/stddev): 59475.0000/0.00
    execution time (avg/stddev): 29.9368/0.00

```

No. of Iterations	Events/sec
1	1986.20 - max value
2	1966.42
3	1963.54 - min value
4	1968.02
5	1982.38
Avg:	1973.31

### Result:

We see that as we increase the max-prime argument value, we see the number of events per second decrease for our current experimental configuration. QEMU still remains faster than Docker desktop on the MacBook M1 chip.

### **Case 3: Max-prime = 200000**

For QEMU:

```
> sysbench cpu --cpu-max-prime=200000 --num-threads=1 --time=30 run
```

OR

```
> source ~/QEMU_200000.sh
```

For Docker:

```
> docker run --rm zyclonite/sysbench --test(cpu) --cpu-max-prime=200000  
--time=30 run
```

OR

```
> source ~/DOCKER_200000.sh
```

*QEMU Results:*

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 200000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 115.02  
  
General statistics:  
    total time: 30.0024s  
    total number of events: 3451  
  
Latency (ms):  
    min: 8.49  
    avg: 8.69  
    max: 64.80  
    95th percentile: 8.74  
    sum: 29972.45  
  
Threads fairness:  
    events (avg/stddev): 3451.0000/0.00  
    execution time (avg/stddev): 29.9724/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 200000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 115.27  
  
General statistics:  
    total time: 30.0073s  
    total number of events: 3459  
  
Latency (ms):  
    min: 8.49  
    avg: 8.67  
    max: 42.81  
    95th percentile: 8.74  
    sum: 29982.60  
  
Threads fairness:  
    events (avg/stddev): 3459.0000/0.00  
    execution time (avg/stddev): 29.9826/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 200000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 115.16  
  
General statistics:  
    total time: 30.0082s  
    total number of events: 3456  
  
Latency (ms):  
    min: 8.50  
    avg: 8.68  
    max: 43.27  
    95th percentile: 8.74  
    sum: 29984.52  
  
Threads fairness:  
    events (avg/stddev): 3456.0000/0.00  
    execution time (avg/stddev): 29.9845/0.00
```

```
Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 200000

Initializing worker threads...

Threads started!

CPU speed:
  events per second: 115.90

General statistics:
  total time:           30.0081s
  total number of events: 3478

Latency (ms):
  min:                 8.49
  avg:                 8.62
  max:                 22.35
  95th percentile:     8.74
  sum:                29984.33

Threads fairness:
  events (avg/stddev): 3478.0000/0.00
  execution time (avg/stddev): 29.9843/0.00
```

```
Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 200000

Initializing worker threads...

Threads started!

CPU speed:
  events per second: 115.31

General statistics:
  total time:           30.0037s
  total number of events: 3460

Latency (ms):
  min:                 8.51
  avg:                 8.66
  max:                 29.42
  95th percentile:     8.74
  sum:                29970.65

Threads fairness:
  events (avg/stddev): 3460.0000/0.00
  execution time (avg/stddev): 29.9706/0.00
```

No. of Iterations	Events/sec
1	115.02 - min value
2	115.27
3	115.16
4	115.90 - max value
5	115.31
Avg:	115.33

### Docker Results:

```
[benitarego@Benitas-MacBook-Air ~ % docker run --rm zyclonite/sysbench --test=cpu ]
--cpu-max-prime=200000 --time=30 run
WARNING: the --test option is deprecated. You can pass a script name or path on t
he command line without any options.
sysbench 1.0.20-f6f6117dc4 (using bundled LuaJIT 2.1.0-beta2)

Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 200000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 104.83

General statistics:
total time: 30.0078s
total number of events: 3146

Latency (ms):
min: 9.39
avg: 9.53
max: 39.84
95th percentile: 9.73
sum: 29985.25

Threads fairness:
events (avg/stddev): 3146.0000/0.00
execution time (avg/stddev): 29.9852/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 200000  
  
Initializing worker threads...  
  
Threads started!  
  
WARNING: the --test option is deprecated. You can pass a script  
he command line without any options.  
CPU speed:  
    events per second: 105.00  
  
General statistics:  
    total time: 30.0076s  
    total number of events: 3151  
  
Latency (ms):  
    min: 9.40  
    avg: 9.52  
    max: 31.79  
    95th percentile: 9.56  
    sum: 29983.65  
  
Threads fairness:  
    events (avg/stddev): 3151.0000/0.00  
    execution time (avg/stddev): 29.9837/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 200000  
  
Initializing worker threads...  
  
Threads started!  
  
WARNING: the --test option is deprecated. You can pass a script  
he command line without any options.  
CPU speed:  
    events per second: 104.94  
  
General statistics:  
    total time: 30.0063s  
    total number of events: 3149  
  
Latency (ms):  
    min: 9.38  
    avg: 9.52  
    max: 29.57  
    95th percentile: 9.73  
    sum: 29981.90  
  
Threads fairness:  
    events (avg/stddev): 3149.0000/0.00  
    execution time (avg/stddev): 29.9819/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 200000  
  
Initializing worker threads...  
  
WARNING: the --test option is deprecated. You can pass a script  
he command line without any options.  
Threads started!  
  
CPU speed:  
    events per second: 104.54  
  
General statistics:  
    total time: 30.0060s  
    total number of events: 3137  
  
Latency (ms):  
    min: 9.39  
    avg: 9.56  
    max: 32.46  
    95th percentile: 9.73  
    sum: 29984.69  
  
Threads fairness:  
    events (avg/stddev): 3137.0000/0.00  
    execution time (avg/stddev): 29.9847/0.00
```

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 200000  
  
Initializing worker threads...  
  
WARNING: the --test option is deprecated. You can pass a script  
he command line without any options.  
Threads started!  
  
CPU speed:  
    events per second: 104.50  
  
General statistics:  
    total time: 30.0071s  
    total number of events: 3136  
  
Latency (ms):  
    min: 9.39  
    avg: 9.56  
    max: 50.34  
    95th percentile: 9.91  
    sum: 29983.37  
  
Threads fairness:  
    events (avg/stddev): 3136.0000/0.00  
    execution time (avg/stddev): 29.9834/0.00
```

No. of Iterations	Events/sec
1	104.83
2	105.00 - max value
3	104.94
4	104.54
5	104.50 - min value
Avg:	104.76

**Result:**

We note that the number of events further decreases as we continue to increase the cpu-max-prime argument value. QEMU VM continues to be faster than Docker desktop for MacBook M1 Chip.

**Conclusion:**

QEMU VM is faster than Docker in terms of CPU performance on the M1 chip. Based on our experimental results, we can conclude that as we increase the cpu-max-prime argument value, our number of events per second continue to decrease, i.e, greater the value of cpu-max-prime argument, lower the number of events per second for sysbench CPU testing.

## File I/O testing QEMU vs Docker

### 1. For QEMU:

#### a. Sequential Rewrite

```
> sysbench --num-threads=16 --test=fileio --file-total-size=3G  
--time=30 --file-test-mode=seqrewr prepare  
> sysbench --num-threads=16 --test=fileio --file-total-size=3G  
--time=30 --file-test-mode=seqrewr run  
> sysbench --num-threads=16 --test=fileio --file-total-size=3G  
--time=30 --file-test-mode=seqrewr cleanup
```

```
[ubuntu@ubuntu:~$ sysbench --num-threads=16 --test=fileio --file-total-size=3G --  
time=30 --file-test-mode=seqrewr run  
WARNING: the --test option is deprecated. You can pass a script name or path on  
the command line without any options.  
WARNING: --num-threads is deprecated, use --threads instead  
sysbench 1.0.18 (using system LuaJIT 2.1.0-beta3)  
  
Running the test with following options:  
Number of threads: 16  
Initializing random number generator from current time  
  
Extra file open flags: (none)  
128 files, 24MiB each  
3GiB total file size  
Block size 16KiB  
Periodic FSYNC enabled, calling fsync() each 100 requests.  
Calling fsync() at the end of test, Enabled.  
Using synchronous I/O mode  
Doing sequential rewrite test  
Initializing worker threads...  
  
Threads started!  
  
File operations:  
  reads/s:          0.00  
  writes/s:        24741.96  
  fsyncs/s:        31737.76  
  
Throughput:  
  read, MiB/s:      0.00  
  written, MiB/s:   386.59  
  
General statistics:  
  total time:       30.0211s  
  total number of events: 1693588  
  
Latency (ms):  
  min:              0.00  
  avg:             0.28  
  max:             49.73  
  95th percentile:  1.23  
  sum:            476077.84  
  
Threads fairness:  
  events (avg/stddev): 105849.2500/1220.44  
  execution time (avg/stddev): 29.7549/0.01
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 0.00 writes/sec = 24741.96 fsyncs/sec = 31737.76 events/sec = 56413.26
2	reads/sec = 0.00 writes/sec = 25091.54 fsyncs/sec = 32183.12 events/sec = 57208.3

3	reads/sec = 0.00 writes/sec = 25099.53 fsyncs/sec = 32192.86 events/sec = 57225.82
4	reads/sec = 0.00 writes/sec = 25774.89 fsyncs/sec = 33056.79 events/sec = 58764.91
5	reads/sec = 0.00 writes/sec = 25716.55 fsyncs/sec = 32983.65 events/sec = 58633.30

b. Combined Random Rewrite

```
> sysbench --num-threads=16 --test=fileio --file-total-size=3G
--time=30 --file-test-mode=rndrw prepare
> sysbench --num-threads=16 --test=fileio --file-total-size=3G
--time=30 --file-test-mode=rndrw run
> sysbench --num-threads=16 --test=fileio --file-total-size=3G
--time=30 --file-test-mode=rndrw cleanup
```

```
[ubuntu@ubuntu:~$ sysbench --num-threads=16 --test=fileio --file-total-size=3G --
time=30 --file-test-mode=rndrw run
WARNING: the --test option is deprecated. You can pass a script name or path on
the command line without any options.
WARNING: --num-threads is deprecated, use --threads instead
sysbench 1.0.18 (using system LuaJIT 2.1.0-beta3)

Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Number of IO requests: 0
Read/Write ratio for combined random IO test: 1.50
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing random r/w test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          12792.27
  writes/s:         8527.94
  fsyncs/s:         27358.04

Throughput:
  read, MiB/s:      199.88
  written, MiB/s:   133.25

General statistics:
  total time:        30.0223s
  total number of events: 1459429

Latency (ms):
  min:                0.00
  avg:                0.33
  max:               20.68
  95th percentile:    1.08
  sum:            478341.26

Threads fairness:
  events (avg/stddev): 91214.3125/464.61
  execution time (avg/stddev): 29.8963/0.01
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 12792.27 writes/sec = 8527.94 fsyncs/sec = 27358.04 events/sec = 48611.5
2	reads/sec = 10980.66 writes/sec = 7320.22 fsyncs/sec = 23490.43 events/sec = 41725.19
3	reads/sec = 11286.93 writes/sec = 7524.45 fsyncs/sec = 24143.98 events/sec = 42888.69
4	reads/sec = 11229.25 writes/sec = 7486.00 fsyncs/sec = 24023.52 events/sec = 42671.85
5	reads/sec = 10204.11 writes/sec = 6802.46 fsyncs/sec = 21835.43 events/sec = 38775.19

2. For Docker:

a. Sequential Rewrite

```
> docker run --rm zyclonite/sysbench --num-threads=16 --test=fileio  
--file-total-size=3G --time=30 --file-test-mode=seqrewr prepare  
> docker run --rm zyclonite/sysbench --num-threads=16 --test=fileio  
--file-total-size=3G --time=30 --file-test-mode=seqrewr run  
> docker run --rm zyclonite/sysbench --num-threads=16 --test=fileio  
--file-total-size=3G --time=30 --file-test-mode=seqrewr cleanup
```

```

Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing sequential rewrite test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          0.00
  writes/s:        126972.76
  fsyncs/s:        162590.76

Throughput:
  read, MiB/s:      0.00
  written, MiB/s:   1983.95

General statistics:
  total time:       30.0030s
  total number of events: 8685809

Latency (ms):
  min:              0.00
  avg:              0.05
  max:             170.08
  95th percentile:  0.00
  sum:            462917.92

Threads fairness:
  events (avg/stddev):    542863.0625/8225.51
  execution time (avg/stddev): 28.9324/0.02

```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 0 writes/sec = 39389.37 fsyncs/sec = 50483.02 events/sec = 289498.02
2	reads/sec = 0 writes/sec = 121807.80 fsyncs/sec = 155980.56 events/sec = 277726.78
3	reads/sec = 0 writes/sec = 121085.32 fsyncs/sec = 155055.63 events/sec = 276075.38
4	reads/sec = 0 writes/sec = 122745.44 fsyncs/sec = 157181.65

	events/sec = 279866.05
5	reads/sec = 0 writes/sec = 125738.83 fsyncs/sec = 161010.54 events/sec = 286683.63

b. Combined Random Rewrite

```
> docker run --rm zyclonite/sysbench --num-threads=16 --test=fileio
--file-total-size=3G --time=30 --file-test-mode=rndrw prepare
> docker run --rm zyclonite/sysbench --num-threads=16 --test=fileio
--file-total-size=3G --time=30 --file-test-mode=rndrw run
> docker run --rm zyclonite/sysbench --num-threads=16 --test=fileio
--file-total-size=3G --time=30 --file-test-mode=rndrw cleanup
```

```
Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Number of IO requests: 0
Read/Write ratio for combined random IO test: 1.50
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing random r/w test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          31880.49
  writes/s:         21253.34
  fsyncs/s:         68076.70

Throughput:
  read, MiB/s:      498.13
  written, MiB/s:   332.08

General statistics:
  total time:        30.0026s
  total number of events: 3634615

Latency (ms):
  min:                0.00
  avg:                0.13
  max:                7.91
  95th percentile:    0.00
  sum:               478006.11

Threads fairness:
  events (avg/stddev): 227163.4375/1282.68
  execution time (avg/stddev): 29.8754/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 31800.49 writes/sec = 21253.34 fsyncs/sec = 68076.70 events/sec = 121143.33
2	reads/sec = 31654.73 writes/sec = 21102.80 fsyncs/sec = 67594.23 events/sec = 120284.84
3	reads/sec = 33374.73 writes/sec = 22249.51 fsyncs/sec = 71265.85 events/sec = 126823.17
4	reads/sec = 32671.59 writes/sec = 21780.76 fsyncs/sec = 69765.30 events/sec = 124150.73
5	reads/sec = 32702.74 writes/sec = 21801.47 fsyncs/sec = 69831.70 events/sec = 124268.86

**Conclusion:** Sequential rewrite operations are faster on QEMU than Docker. While for combined read write, docker is faster than QEMU.

## Configuration 2: 4 cores with 4GB memory allocation

Tests:

4. Max-prime=2000, time=30 seconds
5. Max-prime=20000, time=30 seconds
6. Max-prime=200000, time=30 seconds

### Case 1: Max-prime = 2000

QEMU Results:

```
Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 2000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 113926.26

General statistics:
total time: 30.0001s
total number of events: 3417869

Latency (ms):
min: 0.01
avg: 0.01
max: 6.30
95th percentile: 0.01
sum: 29596.91

Threads fairness:
events (avg/stddev): 3417869.0000/0.00
execution time (avg/stddev): 29.5969/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	113926.26 - Max value
2	113201.06
3	113018.63 - Min value
4	114854.19
5	113129.79
Avg:	113625.98

### Docker Results:

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 52698.93  
  
General statistics:  
    total time: 30.0001s  
    total number of events: 1581012  
  
Latency (ms):  
    min: 0.02  
    avg: 0.02  
    max: 2.89  
    95th percentile: 0.02  
    sum: 29734.11  
  
Threads fairness:  
    events (avg/stddev): 1581012.0000/0.00  
    execution time (avg/stddev): 29.7341/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	52698.93
2	52132.09
3	51893.90 - min value
4	56403.96 - max value
5	56134.25
Avg:	53852.62

### Result:

We note that the number of events further decreases as we continue to increase the cpu-max-prime argument value. QEMU VM continues to be faster than Docker desktop for MacBook M1 Chip.

## Case 2: Max-prime = 20000

QEMU Results:

```
Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:
    events per second: 3931.10

General statistics:
    total time:          30.0003s
    total number of events: 117937

Latency (ms):
    min:                 0.22
    avg:                 0.25
    max:                 7.69
    95th percentile:     0.28
    sum:                29945.74

Threads fairness:
    events (avg/stddev): 117937.0000/0.00
    execution time (avg/stddev): 29.9457/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	3931.10
2	3933.90
3	4281.01 - max value
4	4051.55 - min value
5	3972.59
Avg:	4034.03

### Docker Results:

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 20000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 3007.54  
  
General statistics:  
    total time: 30.0003s  
    total number of events: 90229  
  
Latency (ms):  
    min: 0.30  
    avg: 0.33  
    max: 10.71  
    95th percentile: 0.36  
    sum: 29962.54  
  
Threads fairness:  
    events (avg/stddev): 90229.0000/0.00  
    execution time (avg/stddev): 29.9625/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	3007.54
2	3247.12 - max value
3	3181.35
4	3164.53
5	2917.72 - min value
Avg:	3103.65

### Result:

We note that the number of events further decreases as we continue to increase the cpu-max-prime argument value. QEMU VM continues to be faster than Docker desktop for MacBook M1 Chip.

### **Case 3: Max-prime = 200000**

**QEMU Results:**

```
Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 200000

Initializing worker threads...

Threads started!

CPU speed:
  events per second: 188.73

General statistics:
  total time:           30.0049s
  total number of events: 5663

Latency (ms):
  min:                  5.14
  avg:                  5.30
  max:                  9.16
  95th percentile:      5.47
  sum:                 29999.23

Threads fairness:
  events (avg/stddev): 5663.0000/0.00
  execution time (avg/stddev): 29.9992/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	188.73 - max value
2	188.33
3	185.49 - min value
4	189.86
5	189.83
Avg:	188.44

### Docker Results:

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 169.27  
  
General statistics:  
    total time: 30.0055s  
    total number of events: 5079  
  
Latency (ms):  
    min: 5.68  
    avg: 5.91  
    max: 9.65  
    95th percentile: 7.17  
    sum: 30000.12  
  
Threads fairness:  
    events (avg/stddev): 5079.0000/0.00  
    execution time (avg/stddev): 30.0001/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	169.27 - min value
2	172.51
3	172.34
4	171.67
5	173.24 - max value
Avg:	171.80

### Result:

As we raise the value of the cpu-max-prime parameter, we observe that the number of events continues to decline. For the MacBook M1 Chip, QEMU VM is still quicker than Docker desktop.

## Conclusion:

When it comes to CPU performance on an M1 processor, QEMU VM outperforms Docker. According to our experimental findings, the number of events per second for sysbench CPU testing decreases as the cpu-max-prime argument value increases. In other words, the higher the value of the cpu-max-prime argument, the less events per second are produced.

## File I/O testing QEMU vs Docker

For QEMU:

Sequential Rewrite

```
Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing sequential rewrite test
Initializing worker threads...

threads started!

File operations:
  reads/s:          0.00
  writes/s:        38097.14
  fsyncs/s:        48828.63

Throughput:
  read, MiB/s:      0.00
  written, MiB/s:   595.27

General statistics:
  total time:       30.05168
  total number of events: 2610256

Latency (ms):
  min:                0.00
  avg:                0.18
  max:               35.97
  95th percentile:    0.58
  sum:            478937.83

Threads fairness:
  events (avg/stddev): 163141.0000/887.86
  execution time (avg/stddev): 29.9336/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 0.00 writes/sec = 38097.14 fsyncs/sec = 48828.63 events/sec = 87008.56
2	reads/sec = 0.00

	writes/sec = 39530.62 fsyncs/sec = 50664.21 events/sec = 90199.3
3	reads/sec = 0.00 writes/sec = 35886.20 fsyncs/sec = 33329.60 events/sec = 81888.26
4	reads/sec = 0.00 writes/sec = 33329.60 fsyncs/sec = 42727.80 events/sec = 76088.93
5	reads/sec = 0.00 writes/sec = 33704.58 fsyncs/sec = 43206.91 events/sec = 76908.83

### Combined Random Rewrite

```
Running the test with following options:
Number of threads: 16
Initialising random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Number of IO requests: 0
Read/Write ratio for combined random IO test: 1.50
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing random r/w test
Initialising worker threads...

Threads started!

File operations:
  reads/s:          14847.51
  writes/s:         9898.23
  fsyncs/s:        31739.38

Throughput:
  read, MiB/s:      231.99
  written, MiB/s:   154.66

General statistics:
  total time:           30.0248s
  total number of events: 1693939

Latency (ms):
  min:                  0.00
  avg:                  0.29
  max:                 37.23
  95th percentile:       0.86
  sum:                479034.62

Threads fairness:
  events (avg/stddev): 105871.1875/755.07
  execution time (avg/stddev): 29.9397/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 14847.51 writes/sec = 9898.23 fsyncs/sec = 31739.38 events/sec = 56464.43
2	reads/sec = 13910.0 writes/sec = 9273.0 fsyncs/sec = 29740.0 events/sec = 52915.76
3	reads/sec = 14165.54 writes/sec = 9443.58 fsyncs/sec = 30284.24 events/sec = 53880.36
4	reads/sec = 12349.31 writes/sec = 8232.87 fsyncs/sec = 26409.53 events/sec = 46994.56
5	reads/sec = 12001.06 writes/sec = 8000.63 fsyncs/sec = 25669.69 events/sec = 45638.56

For Docker:  
Sequential Rewrite

```
Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing sequential rewrite test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          0.00
  writes/s:        21687.04
  fsyncs/s:        27826.67

Throughput:
  read, MiB/s:      0.00
  written, MiB/s:   338.86

General statistics:
  total time:       30.0645s
  total number of events: 1486586

Latency (ms):
  min:              0.00
  avg:              0.32
  max:             37.16
  95th percentile:  1.01
  sum:            479122.85

Threads fairness:
  events (avg/stddev): 92911.2500/725.92
  execution time (avg/stddev): 29.9452/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 0.00 writes/sec = 21687.04 fsyncs/sec = 27826.67 events/sec = 49552.66
2	reads/sec = 0.00 writes/sec = 22690.64 fsyncs/sec = 29019.49 events/sec = 51806.5
3	reads/sec = 0.00 writes/sec = 25688.43 fsyncs/sec = 32946.24 events/sec = 58659.23
4	reads/sec = 0.00 writes/sec = 22382.72

	fsyncs/sec = 28716.33 events/sec = 51146.33
5	reads/sec = 0.00 writes/sec = 21880.69 fsyncs/sec = 28071.27 events/sec = 49996.26

### Combined Random Rewrite

```
Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
30iB total file size
Block size 16KiB
Number of IO requests: 0
Read/Write ratio for combined random IO test: 1.50
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling sync() at the end of test, Enabled.
Using synchronous I/O mode
Doing random r/w test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          13659.27
  writes/s:         9185.99
  fsyncs/s:        29204.47

Throughput:
  read, MiB/s:      213.43
  written, MiB/s:   142.28

General statistics:
  total time:       30.0398s
  total number of events: 1559140

Latency (ms):
  min:              0.00
  avg:              0.31
  max:              38.54
  95th percentile:  0.92
  sum:             479919.25

Threads fairness:
  events (avg/stddev):  97446.2588/836.57
  execution time (avg/stddev): 29.9387/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 13850.71 writes/sec = 9320.24 fsyncs/sec = 28826.54 events/sec = 53032.66
2	reads/sec = 14155.49 writes/sec = 8754.94 fsyncs/sec = 27129.15 events/sec = 50045.5

3	reads/sec = 14251.83 writes/sec = 9545.02 fsyncs/sec = 30457.31 events/sec = 54375.26
4	reads/sec = 14324.66 writes/sec = 9231.83 fsyncs/sec = 31432.93 events/sec = 55987.33
5	reads/sec = 14231.90 writes/sec = 9658.14 fsyncs/sec = 30324.53 events/sec = 55432.1

**Conclusion:**

On QEMU, sequential rewrite operations are quicker than Docker. Docker is quicker than QEMU for combined read and write operations.

### Configuration 3: 6 cores with 6GB memory allocation

Tests:

7. Max-prime=2000, time=30 seconds
8. Max-prime=20000, time=30 seconds
9. Max-prime=200000, time=30 seconds

#### Case 1: Max-prime = 2000

QEMU Results:

```
Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 2000

Initializing worker threads...

Threads started!

CPU speed:
events per second: 113197.63

General statistics:
total time: 30.0002s
total number of events: 3396025

Latency (ms):
min: 0.01
avg: 0.01
max: 2.88
95th percentile: 0.01
sum: 29571.95

Threads fairness:
events (avg/stddev): 3396025.0000/0.00
execution time (avg/stddev): 29.5720/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	113197.63 - min Value
2	113318.53
3	114254.59 - max value
4	113529.89
5	114251.46
Avg:	113529.89

### Docker Results:

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 2000  
  
Initializing worker threads...  
  
WARNING: the --test option is deprecated. You can pass a  
Threads started!  
  
CPU speed:  
    events per second: 54026.70  
  
General statistics:  
    total time: 30.0001s  
    total number of events: 1620842  
  
Latency (ms):  
    min: 0.02  
    avg: 0.02  
    max: 2.53  
    95th percentile: 0.02  
    sum: 29741.74  
  
Threads fairness:  
    events (avg/stddev): 1620842.0000/0.00  
    execution time (avg/stddev): 29.7417/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	54026.70
2	53403.88
3	56479.38 - max value
4	52534.09 - min value
5	52843.95
Avg:	53857.6

### Result:

We observe that as we raise the value of the cpu-max-prime parameter, the number of occurrences continues to decline. For the MacBook M1 Chip, QEMU virtual machine is still quicker than Docker desktop.

## Case 2: Max-prime = 20000

QEMU Results:

```
Prime numbers limit: 20000
Initializing worker threads...
Threads started!
CPU speed:
events per second: 4281.04
General statistics:
total time: 30.0001s
total number of events: 128434
Latency (ms):
min: 0.22
avg: 0.23
max: 3.59
95th percentile: 0.25
sum: 29955.70
Threads fairness:
events (avg/stddev): 128434.0000/0.00
execution time (avg/stddev): 29.9557/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	4281.04
2	3942.69 - min value
3	3945.40
4	4284.01 - max value
5	4054.45
Avg:	4101.51

### Docker Results:

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 20000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 3129.44  
  
General statistics:  
    total time: 30.0002s  
    total number of events: 93886  
  
Latency (ms):  
    min: 0.30  
    avg: 0.32  
    max: 3.62  
    95th percentile: 0.34  
    sum: 29964.98  
  
Threads fairness:  
    events (avg/stddev): 93886.0000/0.00  
    execution time (avg/stddev): 29.9650/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	3129.44
2	3264.53 - max value
3	3024.72 - min value
4	3141.35
5	3147.42
Avg:	3141.49

### Result:

As we keep raising the cpu-max-prime parameter value, we see that the number of events continues to drop. In terms of speed for the MacBook M1 Chip, QEMU VM continues to be quicker than Docker desktop.

### **Case 3: Max-prime = 200000**

**QEMU Results:**

```
Running the test with following options:
Number of threads: 1
Initializing random number generator from current time

Prime numbers limit: 200000

Initializing worker threads...

Threads started!

CPU speed:
events per second:    180.38

General statistics:
total time:          30.0028s
total number of events: 5412

Latency (ms):
min:                  5.15
avg:                 5.54
max:                 10.20
95th percentile:     7.04
sum:                29994.08

Threads fairness:
events (avg/stddev): 5412.0000/0.00
execution time (avg/stddev): 29.9941/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	180.38 - min value
2	182.33
3	181.83
4	184.49 - max value
5	180.86
Avg:	181.97

### Docker Results:

```
Running the test with following options:  
Number of threads: 1  
Initializing random number generator from current time  
  
Prime numbers limit: 200000  
  
Initializing worker threads...  
  
Threads started!  
  
CPU speed:  
    events per second: 171.39  
  
General statistics:  
    total time: 30.0003s  
    total number of events: 5142  
  
Latency (ms):  
    min: 5.69  
    avg: 5.83  
    max: 7.56  
    95th percentile: 6.09  
    sum: 29995.04  
  
Threads fairness:  
    events (avg/stddev): 5142.0000/0.00  
    execution time (avg/stddev): 29.9950/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Events/sec
1	171.39
2	174.24 - max value
3	172.67
4	172.47
5	171.34 - min value
Avg:	172.42

### Result:

As we keep raising the cpu-max-prime parameter value, we see that the number of events continues to drop. In terms of speed for the MacBook M1 Chip, QEMU VM continues to be quicker than Docker desktop.

## Conclusion:

When it comes to CPU performance on an M1 processor, QEMU VM outperforms Docker. According to our experimental findings, the number of events per second for sysbench CPU testing decreases as the cpu-max-prime argument value increases. In other words, the higher the value of the cpu-max-prime argument, the less events per second are produced.

## File I/O testing QEMU vs Docker

For QEMU:

Sequential Rewrite

```
Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing sequential rewrite test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          0.00
  writes/s:        25473.58
  fsyncs/s:        32673.72

Throughput:
  read, MiB/s:      0.00
  written, MiB/s:   398.02

General statistics:
  total time:           30.0533s
  total number of events: 1745550

Latency (ms):
  min:                  0.00
  avg:                  0.27
  max:                 39.67
  95th percentile:     0.81
  sum:                479121.05

Threads fairness:
  events (avg/stddev): 109096.8750/702.37
  execution time (avg/stddev): 29.9451/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 0.0 writes/sec = 25473.58 fsyncs/sec = 32673.72 events/sec = 58185
2	reads/sec = 0.0 writes/sec = 27687.54 fsyncs/sec = 33764.65 events/sec = 59124.32

3	reads/sec = 0.0 writes/sec = 24534.32 fsyncs/sec = 31986.87 events/sec = 56126.76
4	reads/sec = 0.0 writes/sec = 25765.32 fsyncs/sec = 32675.83 events/sec = 58234.21
5	reads/sec = 0.0 writes/sec = 28675.77 fsyncs/sec = 35687.32 events/sec = 59543.65

### Combined Random Rewrite

```
Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Number of IO requests: 0
Read/Write ratio for combined random IO test: 1.50
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing random r/w test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          9938.58
  writes/s:         6625.83
  fsyncs/s:        21268.99

Throughput:
  read, MiB/s:      155.29
  written, MiB/s:   103.53

General statistics:
  total time:       30.0701s
  total number of events: 1135621

Latency (ms):
  min:              0.00
  avg:              0.42
  max:             37.10
  95th percentile:  1.44
  sum:            479198.63

Threads fairness:
  events (avg/stddev): 70976.3125/665.20
  execution time (avg/stddev): 29.9499/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 9938.58 writes/sec = 6625.83 fsyncs/sec = 21268.98 events/sec = 37854.03

2	reads/sec = 9876.73 writes/sec = 6543.21 fsyncs/sec = 21124.43 events/sec = 35541.28
3	reads/sec = 9763.23 writes/sec = 6312.63 fsyncs/sec = 24569.69 events/sec = 34532.56
4	reads/sec = 9845.43 writes/sec = 6541.32 fsyncs/sec = 20267.32 events/sec = 35643.76
5	reads/sec = 9765.53 writes/sec = 7021.32 fsyncs/sec = 30832.11 events/sec = 39878.32

For Docker:

### Sequential Rewrite

```
Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing sequential rewrite test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          8.00
  writes/s:        13778.77
  fsyncs/s:       17701.47

Throughput:
  read, MiB/s:      0.00
  written, MiB/s:   215.29

General statistics:
  total time:      39.0893s
  total number of events: 945185

Latency (ms):
  min:              0.00
  avg:              0.51
  max:              56.15
  95th percentile:  1.42
  sum:             479385.45

Threads fairness:
  events (avg/stddev):    59074.0625/661.83
  execution time (avg/stddev): 29.9616/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 0.0 writes/sec = 13778.77 fsyncs/sec = 17701.47 events/sec = 31506.16
2	reads/sec = 0.0 writes/sec = 12456.32 fsyncs/sec = 16323.87 events/sec = 29876.36
3	reads/sec = 0.0 writes/sec = 13543.64 fsyncs/sec = 17634.49 events/sec = 30234.5
4	reads/sec = 0.0 writes/sec = 13534.11 fsyncs/sec = 17431.89 events/sec = 31232.21
5	reads/sec = 0.0 writes/sec = 14328.87 fsyncs/sec = 18785.62 events/sec = 32324.32

## Combined Random Rewrite

```
Running the test with following options:
Number of threads: 16
Initializing random number generator from current time

Extra file open flags: (none)
128 files, 24MiB each
3GiB total file size
Block size 16KiB
Number of IO requests: 0
Read/Write ratio for combined random IO test: 1.50
Periodic FSYNC enabled, calling fsync() each 100 requests.
Calling fsync() at the end of test, Enabled.
Using synchronous I/O mode
Doing random r/w test
Initializing worker threads...

Threads started!

File operations:
  reads/s:          9501.28
  writes/s:         6334.24
  fsyncs/s:        20336.56

Throughput:
  read, MiB/s:      148.46
  written, MiB/s:   98.97

General statistics:
  total time:       30.07725
  total number of events: 1085934

Latency (ms):
  min:              0.00
  avg:              0.44
  max:              40.28
  95th percentile:  1.39
  sum:             479277.20

Threads fairness:
  events (avg/stddev): 67870.8750/669.93
  execution time (avg/stddev): 29.9548/0.00
```

After running the commands in 5 iterations, we get the following results:

No. of Iterations	Output
1	reads/sec = 9501.28 writes/sec = 6334.24 fsyncs/sec = 20336.56 events/sec = 36197.8
2	reads/sec = 9674.98 writes/sec = 6984.73 fsyncs/sec = 20989.89 events/sec = 37019.21
3	reads/sec = 9594.32 writes/sec = 6434.43 fsyncs/sec = 20214.43 events/sec = 35788.33

4	reads/sec = 9287.90 writes/sec = 6021.14 fsyncs/sec = 18673.53 events/sec = 32764.12
5	reads/sec = 9323.32 writes/sec = 6213.41 fsyncs/sec = 20223.12 events/sec = 36097.87

### Conclusion:

Sequential rewrite operations are faster on QEMU than Docker. While for combined read write, docker is faster than QEMU.

### Performance Analysis:

#### 1. QEMU

##### a. 2 cores with 2GB memory allocation

- i. Sequential Rewrite  
read, MiB/s = 0  
written, MiB/s = 708.32
- ii. Combined Random Read  
read, MiB/s = 211.42  
written, MiB/s = 142.36

##### b. 4 cores with 4GB memory allocation

- i. Sequential Rewrite  
read, MiB/s = 0  
written, MiB/s = 599.31
- ii. Combined Random Read  
read, MiB/s = 237.59  
written, MiB/s = 158.33

##### c. 6 cores with 6GB memory allocation

- i. Sequential Rewrite  
read, MiB/s = 0  
written, MiB/s = 402.14
- ii. Combined Random Read  
read, MiB/s = 163.34  
written, MiB/s = 107.02

#### 2. Docker

##### a. 2 cores with 2GB memory allocation

- i. Sequential Rewrite  
read, MiB/s = 0

- ii. written, MiB/s = 560.21
- ii. Combined Random Read  
read, MiB/s = 243.22  
written, MiB/s = 168.56

**b. 4 cores with 4GB memory allocation**

- i. Sequential Rewrite  
read, MiB/s = 0  
written, MiB/s = 332.26
- ii. Combined Random Read  
read, MiB/s = 236.33  
written, MiB/s = 163.59

**c. 6 cores with 6GB memory allocation**

- i. Sequential Rewrite  
read, MiB/s = 0  
written, MiB/s = 208.21
- ii. Combined Random Read  
read, MiB/s = 145.35  
written, MiB/s = 95.74

**Details:**

Name: Benita Rego

Student ID: W1628656

GitHub Username: benitarego

Repository Name: COEN-241-CC

Folder: HW1

Link: <https://github.com/benitarego/COEN-241-CC/>

Commit ID: 4c1a7ad555761295fdb84928728f1cc5df7a740b