



COEN 241

Introduction to Cloud Computing

Lecture 2 - Virtualization





Lecture 1 Recap

- Cloud Computing Primer
- Course Overview
 - Course Objectives
 - Course Structure
 - TODOs
 - Logistics
 - Instructor Information
- Readings
 - Recommended: CCSA 1.1 - 1.3
 - Optional: CCSA 1.5 - 1.15




What is Cloud Computing?

Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., server, network, storage & applications) that can be rapidly provisioned and released with minimal management effort or service provider interaction.*



What is Cloud Computing?

On demand access to



over the internet



Cloud Computing Enables:

- The illusion of “Ubiquitous” & infinite “Shared Pool” of computing resources available “On-Demand”.
 - Eliminates the need to plan far ahead for provisioning.
- Rapidly and automatically provisioned computing resources.
- Minimal hardware / infrastructure management with increased reliability.



Cloud Computing Enables:

- Access to need-based performance.
- The elimination of an up-front commitment by Cloud users.
 - Allows companies to start small and expand only when needed.
- The ability to pay for use of computing resources on a short-term basis as needed (e.g., processors by the hour and storage by the day)
 - Rewards freeing resources when they are no longer useful.



Cloud Service Models



Cloud Deployment Models



**PUBLIC
CLOUD**



**PRIVATE
CLOUD**



**COMMUNITY
CLOUD**



**HYBRID
CLOUD**



Agenda for Today

- HW 0 Preview & Project Overview
- Multitenancy
- Virtualizations
- History of Virtualization
- Readings
 - Recommended: CCSA 2.1-2.2
 - Optional:
 - <http://www.firmcodes.com/memory-thrashing-in-operating-system/>
 - <https://www.ibm.com/cloud/learn/virtualization-a-complete-guide>
 - <http://www.leeandmelindavarian.com/Melinda/25paper.pdf>





HW 0 Preview

- Provided to give you some background context of Linux shell & git
- Keep the Shell tutorial for your reference
- Just need to submit short answers on Camino





Project TODO Overview

- Form Teams by Friday! **Everyone** must submit.
- Project Proposal by next Wednesday!
 - Brief description of the goal, motivation, technologies you will use, etc.
 - More details soon
- Project Presentation (last two (or three) classes of the quarter)
 - 20 minutes of presentation + demo
 - You will get a signup sheet after forming the groups
- Final Project (due finals week)
 - Conference style paper 6~8 pages





Multitenancy


Cloud Computing Enables:

- The illusion of “Ubiquitous” & infinite “**Shared Pool**” of computing resources available “On-Demand”.
 - Eliminates the need to plan far ahead for provisioning.
- Rapidly and automatically provisioned computing resources.
- Minimal hardware / infrastructure management with increased reliability.



The Need for a “Shared Pool”


- CPUs are getting more powerful (Higher speed, More cores)
- Often times, a single user does not fully utilize a single machine
 - Both in processing power and time
 - Average utilization is < 20%
- Often more cost efficient to “split” a large machine than having many small machines (due to fixed cost in wiring, networking and management)



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp


Multitenancy

- Definition: An architecture where multiple users share a single (software or hardware) resource




Multitenancy

- Benefits
 - Cost efficient
 - Easier to monitor and manage
 - Easily scalable
- Drawbacks
 - Generally more complex architecture
 - Security and isolation
 - Allocation



How to Enable Multitenancy?

- **Virtualization!**
 - A core part of this course
- Definition: Using **software** to create an abstraction layer over the **hardware** that allows the hardware elements of a single computer—processors, memory, storage and more—to be divided into **multiple virtual elements**.






Virtualization


Why can't we just have multiple processes?

- What are processes?
- How are they different from Threads?
- To make things simple, can't we just run a set of separate processes (or threads) for a given user?
 - We have multi-core systems now
 - Simple to manage
 - OS is now smart enough!
 - No need to learn virtualization! Bye bye
 - Why can't we do this?



Why can't we just have multiple processes?

- Hard to control resource contention
 - Need fine grained resource management for each process
- Hard to manage dependencies
 - All the users must be compatible with the shared OS
 - Cannot run different versions of the kernel
- Hard to guarantee performance!
 - Memory Thrashing
 - Unfair I/O access







Exercise: Memory Sharing

```
int main() {  
    int i;  
    char *p[2048];  
  
    for(i=0;i<2048;i++) {  
        p[i]=malloc(1024000); // allocate 1 MB  
        memset(p[i],'\0',1024000); // set as 0s  
    }  
  
    while(1) {  
        if (i>=2048)  
            i = 0;  
        memset(p[i],'\0',1024000);  
        i++;  
    }  
}
```

- What does this program do?
- What would happen if we run 10 processes with the same code given on the right?
- What happens if we have machines with following amounts of RAM?
 - 2GB
 - 10GB
 - 20GB

Exercise: Unfairness of I/O Resource Sharing



- P1 ■ 4 KB Synchronous Sequential Read (IO # = 1)
- P2 ■ 4 KB Asynchronous Random Read





Exercise: Unfairness of I/O Resource Sharing

- What happens if we increase the number of asynchronous IOs to 4?
- What happens if we increase the number of asynchronous IOs to 32?

Memory Thrashing



- Memory thrashing is a problem that arises when the memory being accessed is much more than the physical memory.
 - Unpredictive contention
 - Any solution?

Memory Management





<http://www.firmcodes.com/memory-thrashing-in-operating-system/>

Example: Unfairness of I/O Resource Sharing



Example: Unfairness of I/O Resource Sharing




Pros & Cons of Having Multiple Processes

- Pros:

- Multiple applications (from different users) can run “simultaneously”.
- Maximize resource utilization
e.g., multiple processes on a multi-core system


- Cons

- Contention
- Dependency/Compatibility Issues
 - The application’s version and the OS version should (strictly) match one another



Potential Solutions

- Contention
 - Strong regulation, such as limiting (via metering)
 - More advanced resource management: limit, share, and reservation for multiple resources, such as CPU, memory, block I/O, and network
- Dependency/compatibility
 - Provide processes with their own view of the system
 - With their own running environments (i.e., library versions)
- **Virtualization!**





What is Virtualization?

- The act of creating a virtual (rather than actual) version of something, including virtual computer hardware platforms, storage devices, and computer network resources.
- Virtualization is about adding a layer for manageability
 - Virtualization support is increasingly being added into various software and hardware



Types of Virtualization

- **Application Virtualization**

- E.g., JVM, PVM

- **OS Virtualization**

- Containers


- **System Virtualization**

- Virtual Machines

- Network Virtualization


- Storage Virtualization

- Many more...



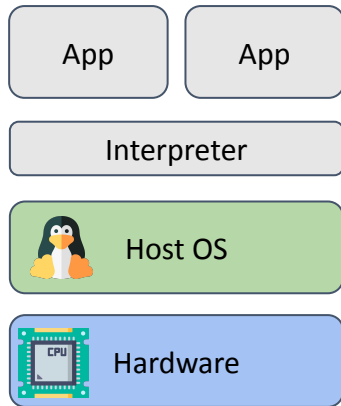
Application Virtualization

- Purpose: Provide a **platform-independent programming environment** that abstracts the underlying hardware or OS and allows a program to execute in the same way on any platform
- What is virtualized?
 - Processor
 - System calls
 - Memory management
- Each application is run as a **separate process**
- Examples: Java Virtual Machine



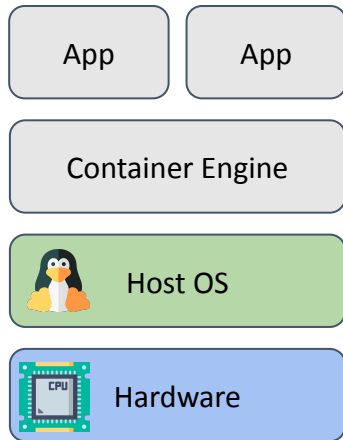
Application Virtualization

- Benefit
 - Portability
- Drawbacks
 - Language dependent, i.e, not flexible




OS Virtualization

- Virtualizing everything above the Host OS **Kernel**
 - Userspace
- We will revisit this concept!



Digression: What is a Kernel?

- **Operating System** is system program that provides interface between user and computer.
- **Kernel** is the core component of an operating system. It is also a system program and a part of of OS which converts user commands into machine language.






Digression: What is a Kernel?

	KERNEL	OPERATING SYSTEM
Definition	Kernel is part of an OS	Operating System is a system program.
Interface	Kernel is an interface between software and hardware	Operating System is an interface between user and hardware.
Type	Monolithic and Microkernels	Single and Multiprogramming batch system, Distributed, Realtime
Purpose	Manages kernel memory, processes, task and disk	In addition to the responsibilities of Kernel, OS provides added security, UI, window manager and etc.


System Virtualization

- What is virtualized?
 - CPU: vCPU
 - vMem
 - virtual network card
 - virtual hard disks
 - Even usb ports and more
- Creates a virtual computer system, called **virtual machine (VM)** consisting of the set of (virtual) hardware, upon which a full (guest) OS runs.
- **Hypervisor** manages the virtual machines



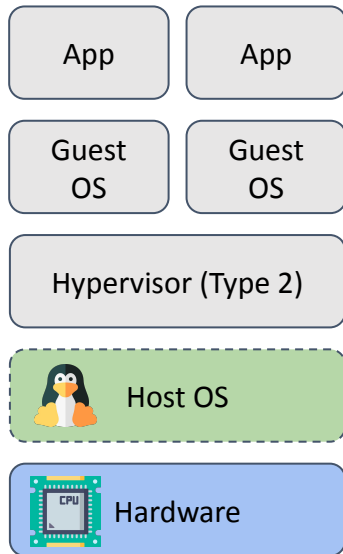
Hypervisor (Virtual Machine Monitor)

- A thin layer of software that's between the hardware and the operating system
- Virtualizes and manages all (or most) hardware resources for VMs
- Examples
 - Virtualbox
 - **QEMU**
 - VMWare
 - Parallels
 - Xen
 - KVM
- We will cover this more next lecture



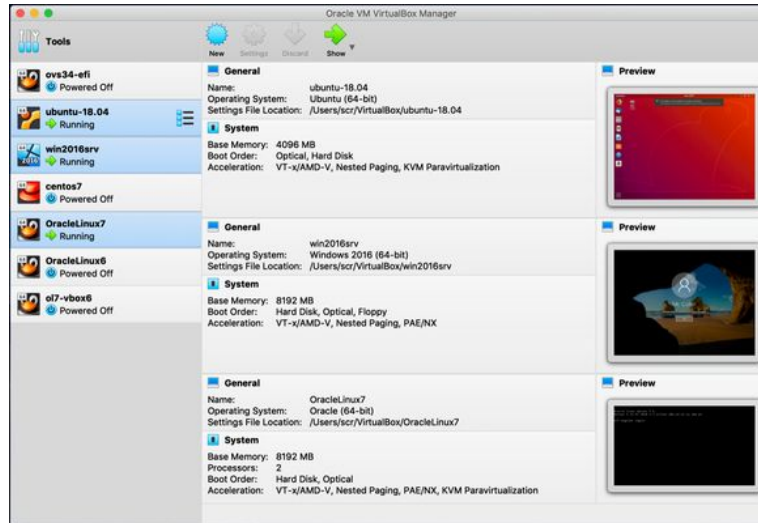
Virtual Machines

- A fully protected and isolated copy of the underlying physical machine's hardware (i.e., emulated by software)
- Must be created configured
- Portable and reconfigurable as needed
 - Open Virtualization Format (OVF)
 - Disk image



Virtual Machines

- Example of Virtual Machines
 - Virtual Disk Image file (VDI) for VirtualBox



App

App

Guest
OSGuest
OS

Hypervisor (Type 2)




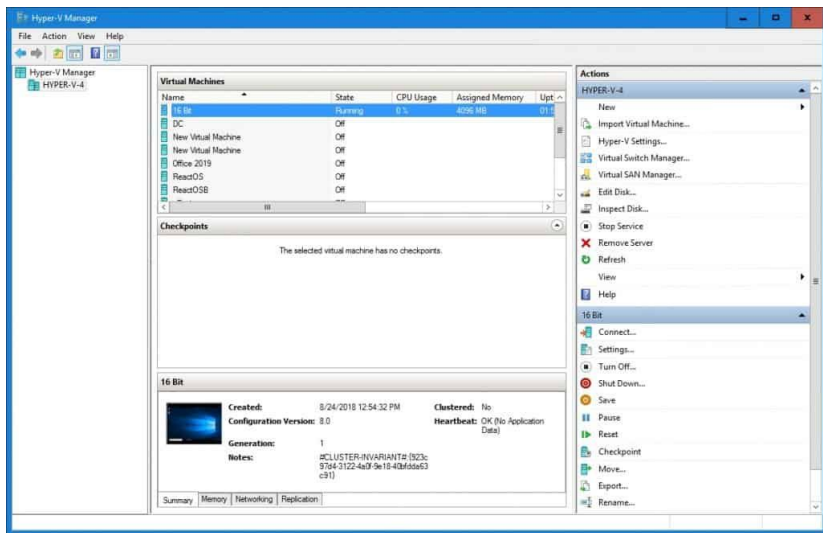
Host OS



Hardware

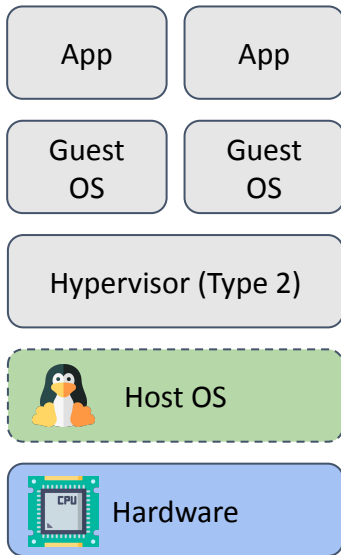
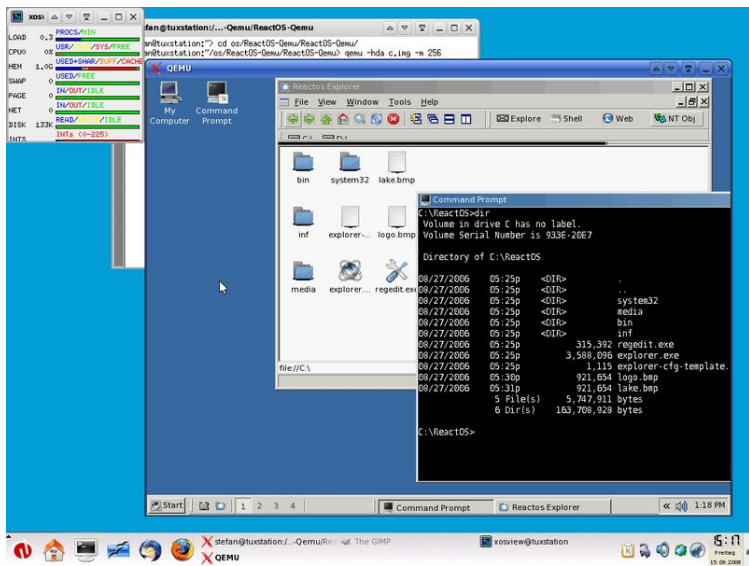
Virtual Machines

- Example of Virtual Machines
 - Virtual Hard Disk image (VHD) for Hyper-V



Virtual Machines

- Example of Virtual Machines
 - “QEMU Copy On Write” (qcow) disk image for QEMU



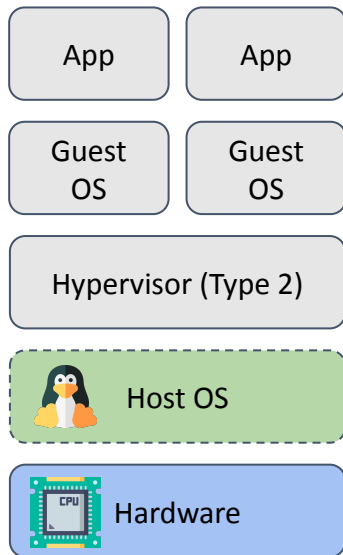
Pros & Cons of System Virtualization

- Pros

- Strong isolation (than containers)
- More levels of protection – VM's kernel and host kernel
- Full OS


- Cons

- Performance overhead
- Large memory & storage footprint
- More configurations needed



History of System Virtualization

- IBM VM/370 – A Hypervisor for IBM mainframe
 - Released in 1972
 - Build for IBM System/370 Mainframes
 - Model 145 priced from \$705,775 to \$1,783,000
 - Reimplementation of earlier CP/CMS OS.
 - Control Program/Cambridge Monitor System
 - Discontinued time-sharing OS of the late 1960s and early 1970s
 - Enables multiple OS environments
 - Full virtualization of the physical machine
- Feasible option when only few & costly machines are around



History of System Virtualization

- VMs became popular research idea in 1960s and 1970s
 - Entire conferences on virtual machine monitors
 - Hardware, Hypervisors & OS were often designed together
 -
- Interest died out in the 1980s and 1990s
 - Hardware got much cheaper
 - Operating systems got more powerful (e.g. multi-user)



History of System Virtualization

- Early x86 computers were underpowered
 - Windows 95 with few applications often maxed-out resources
 - Server work left to expensive server-class computers
- PC computing power increased
 - X86 servers began to appear
 - Mostly under-utilized due to no virtualization
- x86 virtualization surged mid-2000s thanks to the Xen hypervisor



History of System Virtualization

- Commercial virtual machines for x86 architecture
 - Connectx VirtualPC (now Microsoft)
 - VMware (1999-)
 - KVM (Linux)
- Research virtual machines for x86 architecture
 - **Xen (SOSP '03)**
 - Plex86
- Now widely used in all public cloud providers!





Agenda for Today

- HW 0 Preview
- Multitenancy
- Virtualizations
- Readings
 - Recommended: CCSA 2.1-2.2
 - Optional:
 - <http://www.firmcodes.com/memory-thrashing-in-operating-system/>
 - <https://www.ibm.com/cloud/learn/virtualization-a-complete-guide>





TODOs!

- HW 0
 - On Camino
- Start looking for teams ASAP!





Questions?

