



COEN 241

Introduction to Cloud Computing

Lecture 8 - Kubernetes





Lecture 7 Recap

- Kata Container
- Orchestration in Cloud
- Infrastructure as Code
- CoreOS

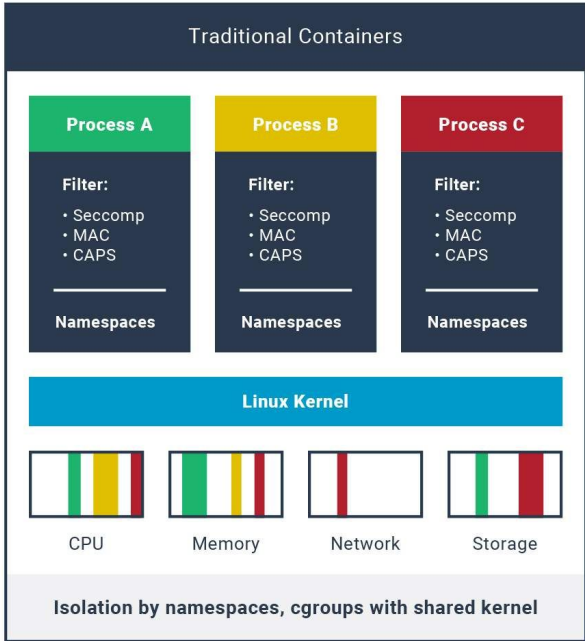
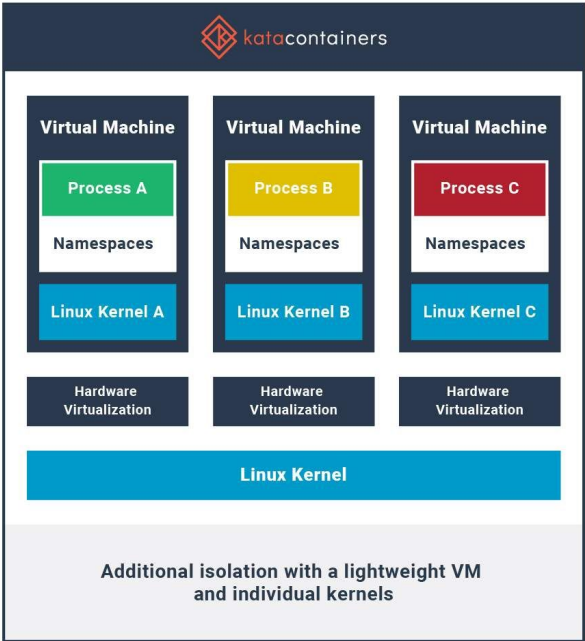


What is Kata Container?


- Released in 2017
 - From the merge of Intel's Clear Containers and Hyper's runV
 - "Wraps" containers into dedicated virtual machines
 - OCI runtime implementation can be plugged into the container engine
 - Docker
 - Can consume existing container images
- Kata is a container runtime
 - Can still be coupled with other Docker platforms



Kata Architecture



Kata Workflow



Orchestration

- Definition: Automatic management of computer systems
 - Deployment and configuration
 - Interconnection and coordination
 - Monitoring
 - Can also configure to hook up monitoring for management
 - E.g., Deploy more resources as load increases
- Growing set of very good (open-source) solutions:
 - **Machine focused:** e.g., Puppet, Terraform, Ansible, Salt, Chef...
 - **Cloud-based:** e.g., AWS CloudFormation, Terraform
 - **Container clusters:** e.g., Kubernetes



Automation vs Orchestration

- **Automation:** completing a **single task** or function without human intervention
- **Orchestration:** Managing a large-scale virtual environment or network by orchestrating the scheduling and integration of automated tasks between complex distributed systems and services
 - Simplifies interconnected workloads, repeatable processes, and operations.
- To simplify, **Automation** refers to a single task vs **Orchestration** arranges multiple tasks to optimize a workflow



Declarative Configuration Management

- Declarative tools specify the desired target state
 - E.g., Can I have a coffee on my desk at 9AM on Monday morning?
- The means to reach target state is up to the configurations
 - Can take corrective action to react to drift in machine's state
- State specification will be a domain specific language (DSL)
- Some example FOSS systems with large user communities
 - Puppet, Terraform, SaltStack



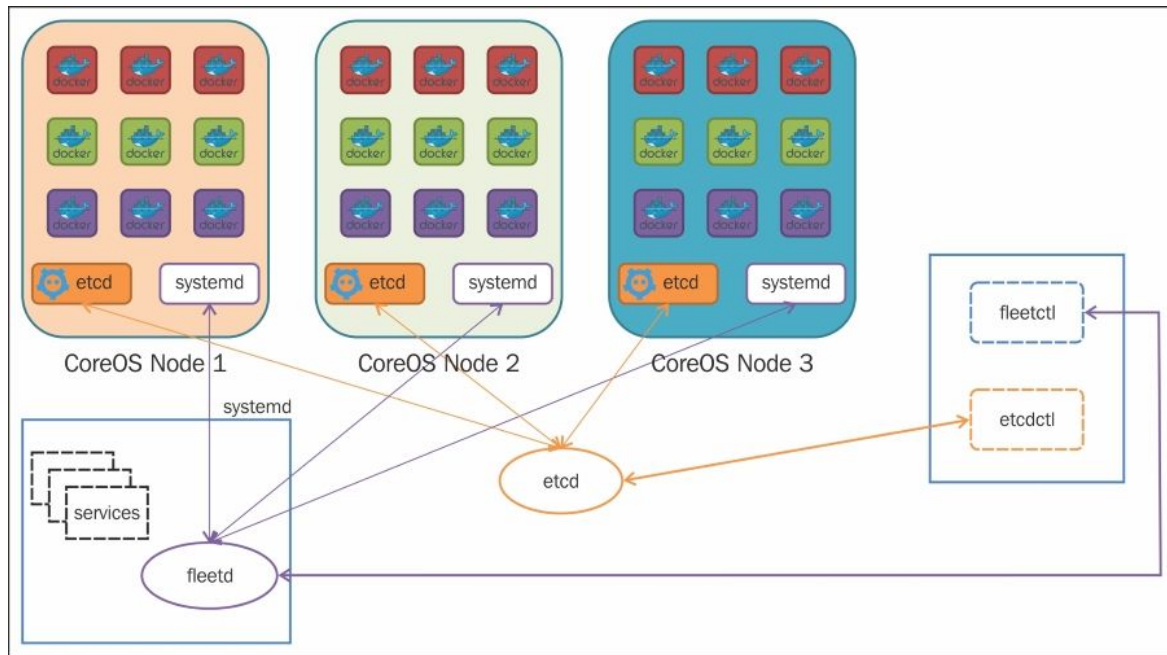


Imperative Configuration Management

- Also called 'procedural', i.e., specifying steps to run:
 - Usually written in chunks of code in configuration system authors' favorite PL
- Some example systems:
 - Ansible (Py), Chef (Ruby), Saltstack
- Can write imperative code to have declarative effect



Core OS Architecture





Agenda for Today

- Kubernetes
 - Concepts
 - Architecture
- Minikube
- Readings
 - Recommended: None
 - Optional:
 - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
 - <https://labs.play-with-k8s.com/>
 - <https://minikube.sigs.k8s.io/docs/handbook/>
 - <https://circleci.com/blog/what-is-yaml-a-beginner-s-guide/>





Kubernetes Overview

What is Kubernetes?

- An open-source container-orchestration system for automating computer application deployment, scaling, and management.
- It is a Greek word meaning helmsman or pilot.
- Often called K8s, which is derived by replacing the eight letters of “ubernet” with the digit 8.
- Project emerged from Google in 2014; v1.0 released 2015
 - Related to Borg—Google’s (secret) internal container scheduler
 - K8s is implemented in Go, in contrast to Borg’s C++



Why is Kubernetes Useful?

- Provides an easy way to scale your containers
- Avoids cloud vendor lock-in
- Allows easy multi-cloud adoption
- Better management of application
- Works with many different container technologies
- It is a BIG piece of software with lot of moving parts
 - We will give a short, high-level overview of it
 - There are numerous tutorials and classes dedicated to Kubernetes





Interacting with Kubernetes

- kubectl
- APIs



Some Kubernetes Terms

- **Pods:** tightly coupled **set of containers**; smallest unit of scheduling
- **Node:** A Pod always runs on a Node. A Node is **a worker machine** in Kubernetes and may be either **a virtual or a physical machine**.
- **Services:** An abstraction which defines a logical set of Pods and a policy by which to access them
 - Sometimes this pattern is called a micro-service architecture
- **Volumes:** Persistent storage that can be shared between containers
- **Namespaces:** Multiple virtual cluster in a same physical cluster



Kubernetes Pods

- Pods are the smallest deployable units of computing that you can create and manage in Kubernetes
- Group of one or more containers with **shared storage** and **network resources**, and a specification for how to run the containers
 - A group of Docker containers with shared namespaces and file system volumes
- A Pod's contents are always co-located and co-scheduled, and run in a shared context




Kubernetes Pods

- Pods are assigned an IP address, for networking
 - All containers within the pod share that address and its ports
- Pods provide app. storage (volumes) to containers
- Pods usually created by controllers, and not directly
- List and show detailed information of the Pods by running
 - `kubectl get pods`
 - `kubectl describe pods`



Kubernetes Pod Overview






Example Pod Definition (PodSpec)

```
apiVersion: v1
kind: Pod
metadata:
  name: multi-pod
  label: myApp
spec:
  volumes:
    - name: html
      emptyDir: {}
  containers:
    - name: 1st
      image: nginx
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
    - name: 2nd
      image: debian
      volumeMounts:
        - name: html
          mountPath: /html
      command: ["/bin/sh", "-c"]
      args:
        - while true; do
            date >> /html/index.html;
            sleep 1;
          Done
```

Kubernetes Nodes

- A set of worker machines that run containerized applications
- Every cluster has at least one worker node
- When you deploy K8s, you get a **cluster**
- `kubectl get nodes`



Kubernetes Services

- A **logical abstraction** for a deployed group of pods in a cluster
 - which all perform the same function
- Kubernetes way of configuring a proxy to forward traffic to a set of pods
- Since pods are ephemeral, a service enables a group of pods to be assigned a name and unique IP address (clusterIP)
- As long as the service is running on that IP address, it will not change.
- Services also define policies for the access to Pods



Types of Kubernetes Services

- **ClusterIP** (default): Exposes a service which is only accessible from within the cluster.
- **NodePort**: Exposes a service via a static port on each node's IP.
- **LoadBalancer**: Exposes the service via the cloud provider's load balancer.
- **ExternalName**: Maps a service to a predefined externalName field by returning a value for the CNAME record.

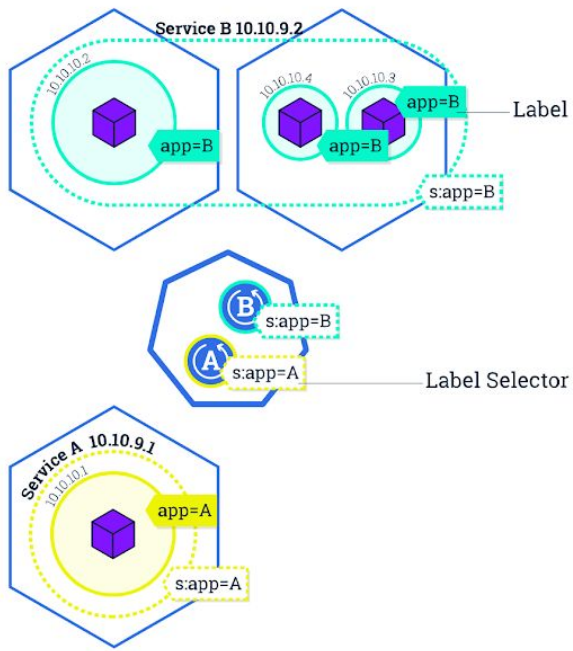


Example Service Definition

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```



Service Overview



Stateless vs Stateful Applications

- Stateless applications scale easily: just start more pods
 - E.g., web servers presenting read-only workloads
- Stateful apps are more difficult, e.g.:
 - Databases having primary and secondary instances
 - Machine Learning Applications
- Kubernetes StatefulSet controllers enables stateful applications
 - Provides durable storage
 - Persistent ID for each Pod
 - <https://kubernetes.io/docs/tutorials/stateful-application/basic-stateful-set/>

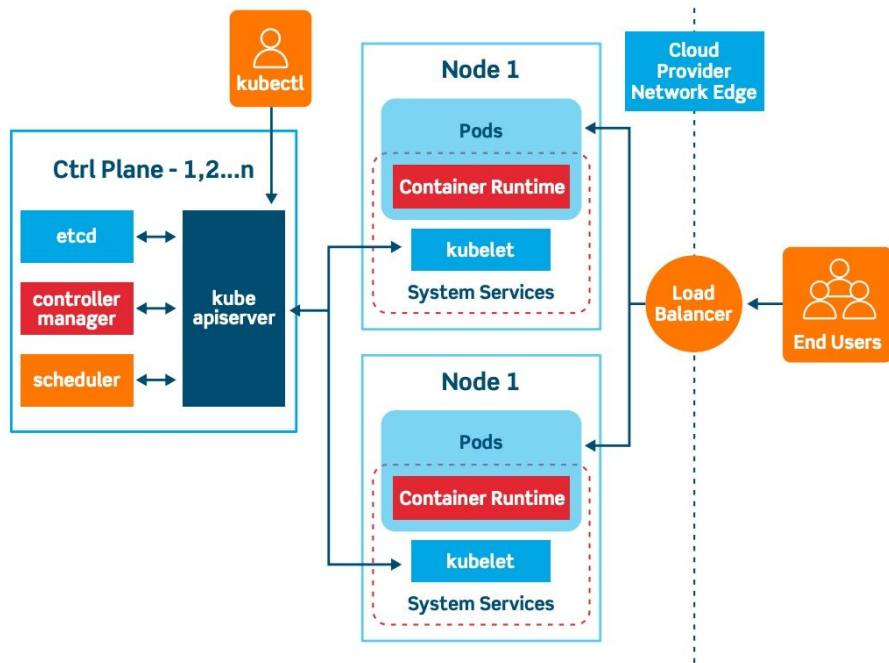




Kubernetes Architecture

Kubernetes Architecture Overview

- Kubernetes cluster consists of:
- Master Node(s) (control plane)
 - With a distributed storage (etcd)
- A number of cluster nodes



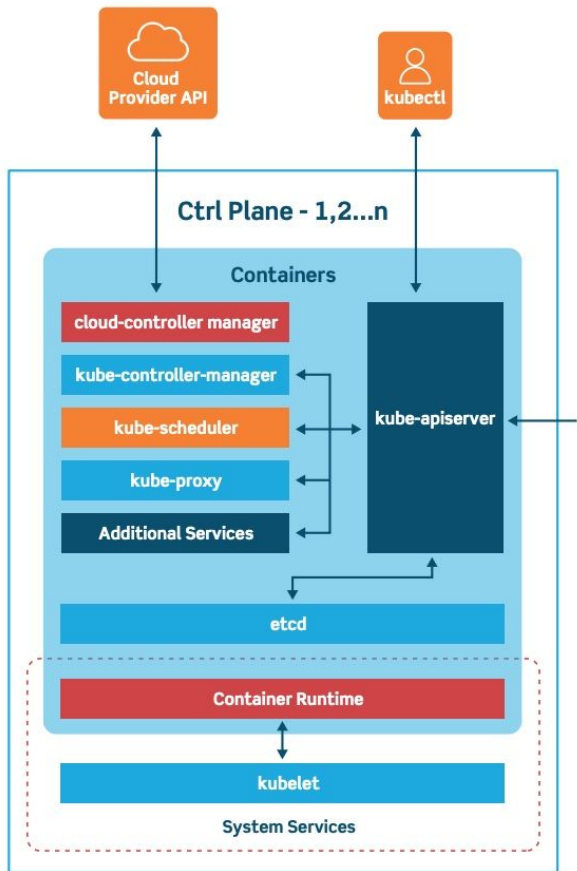
Kubernetes Master Node

- A logically centralized control point
- Maintains a record of all Kubernetes objects
 - Manages object states, responding to changes in the cluster
 - Works to make the actual state of system objects match the desired state
- Consists of the following components
 - **API server:** Allows Kubernetes cluster to be controlled from a user
 - **Controller manager:** Checks replication, status of nodes and state of cluster
 - **Scheduler:** Allocates pods waiting to run nodes
 - **etcd:** Consistent repository of configuration information



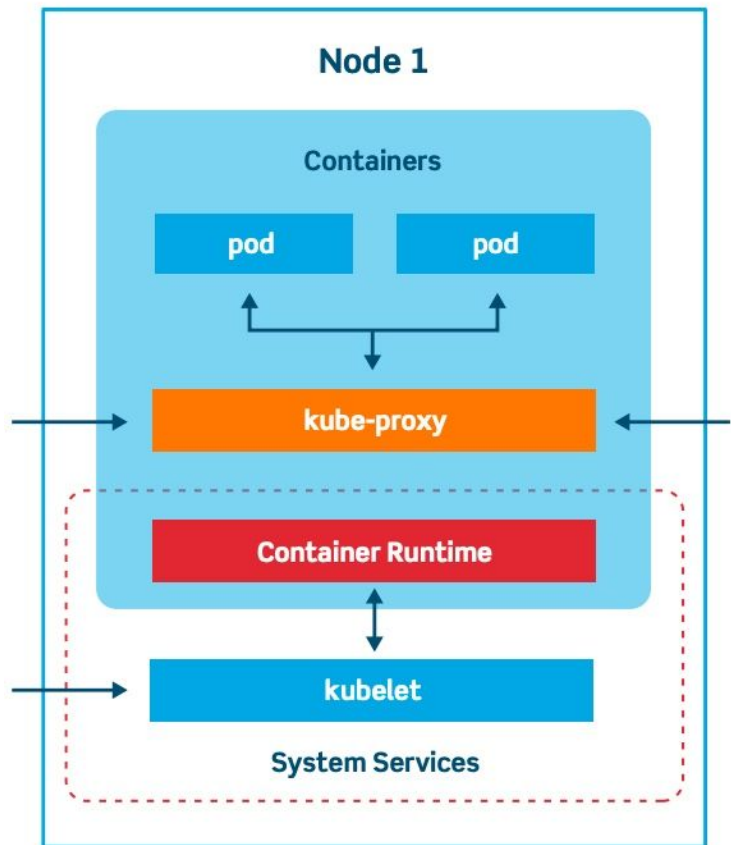
Kubernetes Master Node

- Can be one or more master nodes
- Can accept APIs from cloud providers
- Also provides proxy for network access
- **etcd** can be configured for multi-node distributed configuration manager



Kubernetes Cluster Node

- Nodes can run pods but also,
 - kube-proxy: provides network services; leveraging OS facilities
 - **kubelet**: checks on health of containers within a pod
 - cAdvisor: provides statistics about container resource use
- **Kubelet** is the primary and most important controller
 - Drives the container execution layer



Kubelet

- The kubelet is the primary "node agent" that runs on each node
- It can register the node with the kube apiserver using one of:
 - The hostname
 - A flag to override the hostname
 - Specific logic for a cloud provider
- The kubelet ingests a PodSpec and runs each pod





Extension/Alternatives to Kubernetes

K8s as a Service

- K8s can manage your containers, but how to set it up given a IaaS?
 - Need **VMs** running master and the nodes
- Cloud providers can provide a range of proprietary options
 - AWS
 - AWS Fargate provides a complete container service
 - AWS EKS provides control plane; you set up K8s nodes on EC2
 - Use EC2 to deploy all the components if you want full control
 - Azure and GCP has a similar services

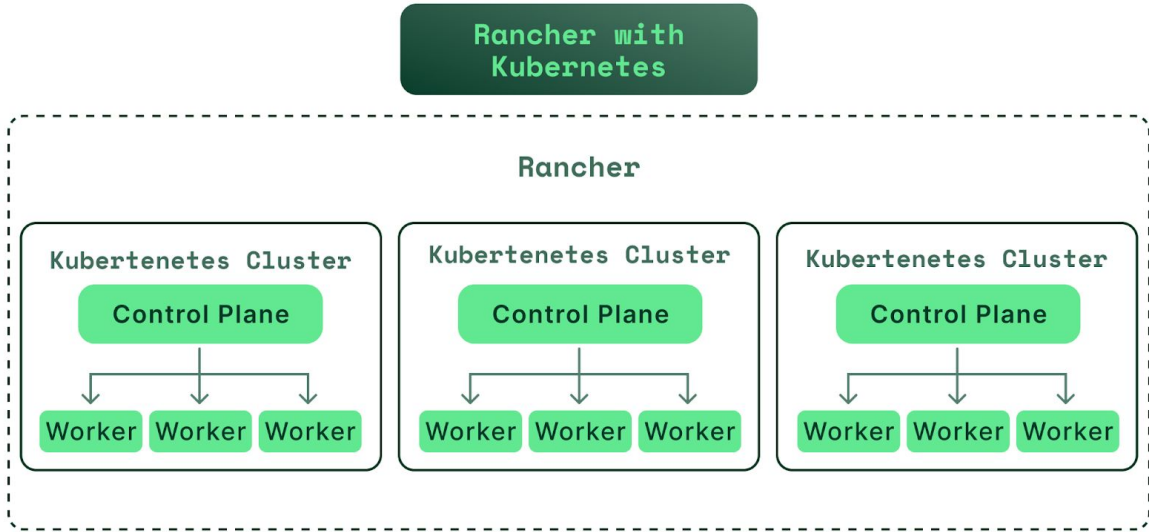


K8s in Multi-Cloud Application

- Quite difficult!
 - Very Manual
 - Or you can use a different service
- Tools like Rancher provide for **multi-K8s** cluster management
 - Can even import existing K8s cluster!
 - <https://rancher.com/>
 - <https://www.youtube.com/watch?v=LX0zVh-AYm4>
- Works because cloud providers' container services are very similar!
 - Benefit of learning the basics!



Rancher versus Kubernetes



Rancher versus Kubernetes

- **Kubernetes Features**
 - Kubernetes based platform is easily migratable across cloud providers
 - Easier to scale as compared to traditional applications hosted in virtual machines (VM)
 - Easy to control cluster density and autoscaling
 - In node failures, pods are automatically rescheduled to other nodes
 - Can be hard to manage multiple clusters
- **Rancher Features**
 - Easy visualization to manage multiple clusters
 - Can easily create new clusters or add existing ones to it
 - Add a concept of projects for better grouping of namespaces
 - Permissions can be configured per project across clusters



Terraform versus Rancher, Kubernetes, etc.

- Terraform is a stack below tools like Rancher, it:
 - Allows previous of the plans using a form of IAC
 - Can easily provision at level of particular instance type
 - Can work with Rancher and Kubernetes
- Rancher can also deploying Kubernetes to bare metal servers
 - Including monitoring & security management tools
- If you need specific type of infrastructure nodes it can get complicated
 - E.g., configuring a GPU node on Amazon for deep learning
 - Containers can use this hardware, but what about container managers



Why it's so hard to pick a “winner”?

- All are changing rapidly!
 - CoreOS Linux's original fleet functionality replaced by K8s
 - Rancher's original functionality are being replaced by Docker Swarm
- Then what to use? Consider your and **the developers' time**
 - Look to see whether a new tool can optimise your processes
 - when taking into account the cost of transition
- Must have **IAC** and continuous integration pipelines
 - IAC is a must for quickly adapting it!

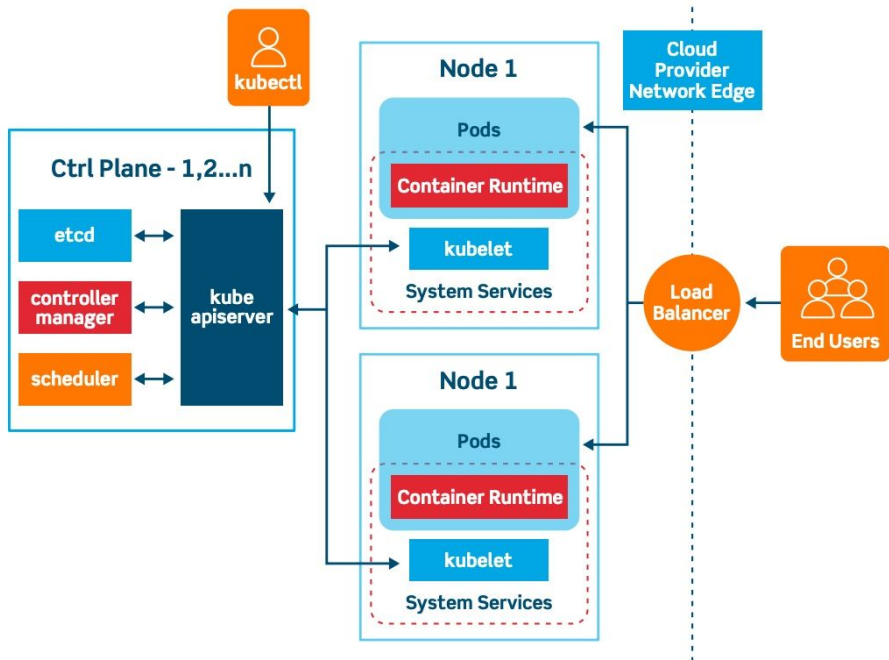




Minikube

Production Kubernetes Cluster

- Requires
 - Multiple Master Nodes
 - Multiple Worker Nodes
- Hard to try Kubernetes locally!
- Minikube to the rescue



Minikube: A local Kubernetes setup

- <https://minikube.sigs.k8s.io/docs/start/>
- Running both master and worker processes **on the same node**
- Minimum requirement: 2 CPUs, 2GB of RAM, 20GB of free disk space
- Use **kubectl** to interact with the Minikube cluster
- **Kubectl**: A CLI for Kubernetes cluster to talk to the master
 - Not only Minikube, but any type of Kubernetes cluster
 - <https://kubernetes.io/docs/reference/kubectl/overview/>





TODOs!

- Kubernetes Quiz!
 - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
 - Up to Module 4 Only!
- HW 2 will be out next class
- Study for Midterms!



Agenda for Today

- Kubernetes
 - Concepts
 - Architecture
- Minikube
- Readings
 - Recommended: None
 - Optional:
 - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
 - <https://labs.play-with-k8s.com/>
 - <https://minikube.sigs.k8s.io/docs/handbook/>
 - <https://circleci.com/blog/what-is-yaml-a-beginner-s-guide/>





Questions?

