



COEN 241

Introduction to Cloud Computing

Lecture 12 - Software Defined Networking





Midterm Feedback

- It would be awesome if you can fill out this survey!
 - <https://forms.gle/HQ7xxptrNdkYy49p6>





Lecture 11 Recap

- Networking 101
- Data Center Networking



OSI Model : Building Block of Networking

- Open Systems Interconnection provides a standard for different computer systems to be able to communicate with each other
- Can be seen as a universal language for computer networking
- It's based on the concept of splitting up a communication system into **seven abstract layers**, each one stacked upon the last
- Each layer has different **protocols** implementations
 - **Protocols:** An established set of rules that determine how data is transmitted between different devices in the same network
 - Example: HTTP Status Codes: 200, 401, 500
HTTP Commands: GET, POST




OSI Model : Building Block of Networking



Data Carried at Different Layers

- Data for applications using IP is augmented at each layer below
 - Ethernet frames' total size larger than IP packets' total size for a given data
- UDP shown here transports datagrams (chunks of data)
- TCP instead transports coherent streams of data
 - Includes retransmissions and congestion control



Type of Traffic and Workloads in a Data Center

- **Traffic to/from Internet**
 - The node-to-node internal network use is small
- **Consider scale-out applications within a DC**
 - Interacting servers do so with distributed effect & high volume
 - Application dependent
 - MapReduce, Storage Replication
- **New 'interesting' traffic patterns: (i.e., non-unicast)**
 - Anycast: traffic reaches any 'nearby' applicable host
 - Multicast or 1-N: traffic is being sent to multiple hosts simultaneously




Data Center Networks

- Once users' packet reaches the data center, it enters the DC network
- Data centers have a structured physical layout
 - Physical servers are grouped into racks, which usually provide:
 - Intra-rack network switches; power distribution; wiring management
 - Racks grouped into aisles or zones—for maintenance / access
- DC networking is hierarchically organized:
 - Core router receiving data from the internet
 - Aggregation layer receiving data from the core
 - Top-of-rack switches installed within racks
 - **Virtual networking within physical servers themselves**




Data Center Networks



Data Center Network Topology (Fat-Tree)

- The three-tier, fat-tree topology mentioned is now outdated
- Multiple issues with the three-tier design
 - Upper-level routers: highly specialized and expensive
 - Designed to sustain high bandwidth in all directions
 - Core routers' prices in the order of \$100,000 a piece
 - High energy usage
 - Not well suited for intra-DC traffic (why?)




Data Center Network Topology (Clos)


- Modern trend is toward using commodity switching kit
 - Create an aggregation switch from a set of cheaper switches
 - Employs a particular addressing scheme and routing algorithm
 - To build and configure by the DC operators
- Often related to a topology known as a Clos network
 - Have ingress, middle and egress stages built from switches
 - The structure can be used recursively (expand middle)
 - Simpler and flatter




Data Center Network Topology (Clos)



Data Center Network Topology (Clos)



Data Center Switch Architecture

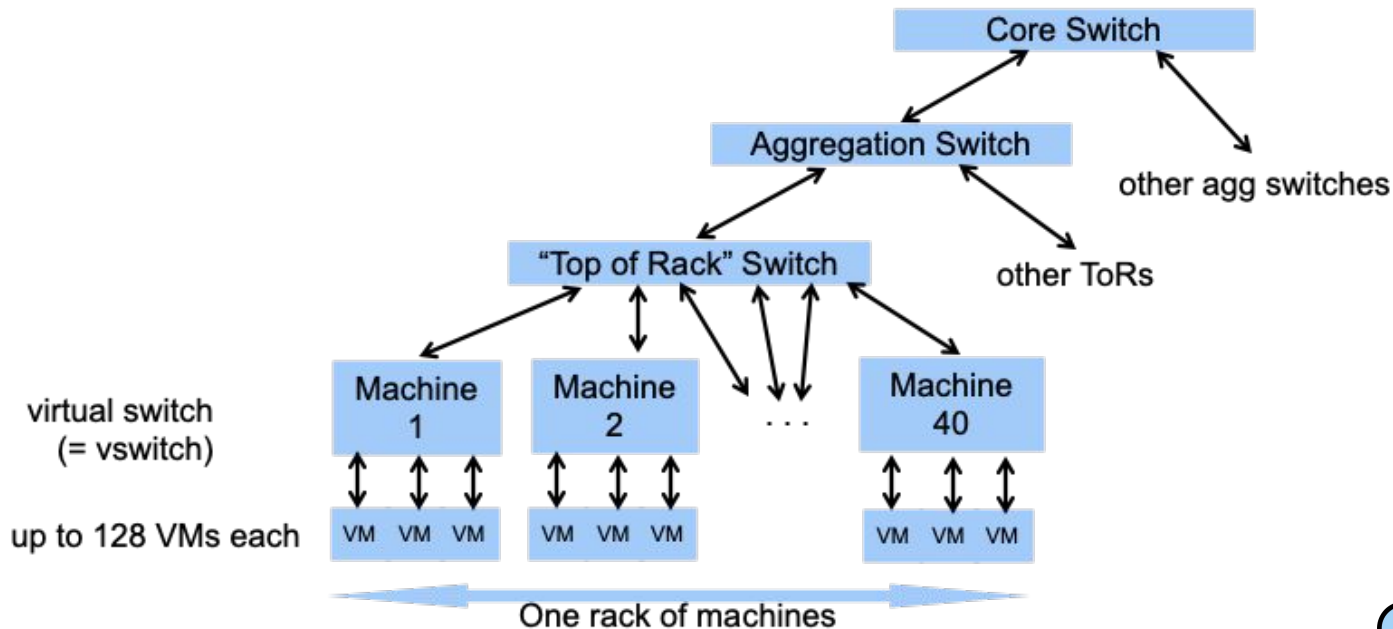


Virtual Switch

- Switch within a server that switches packets between VMS
 - Another big part of Data Center networking and **Software-Defined Networking**
- **Data Plane**
 - Software that is responsible for actual packet switching
 - Implemented in Intel DPDK and/or in Kernel modules
- **Control Plane**
 - Software that is responsible for making routing decisions
- Examples:
 - Open vSwitch (OvS)



Virtual Switch: DC Network After VMs





Agenda for Today

- Software Defined Networking
- OpenFlow
- Readings
 - Recommended: <http://yuba.stanford.edu/~casado/ethane-sigcomm07.pdf>
 - Optional:
 - <http://ccr.sigcomm.org/online/files/p69-v38n2n-mckeown.pdf>





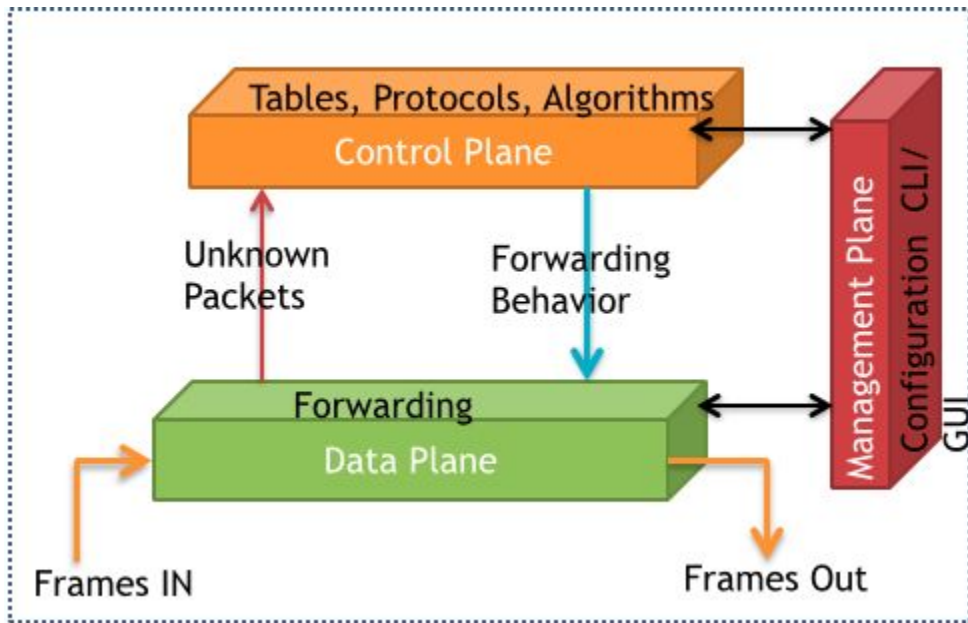
“Planes” in Networking

The Networking “Planes”

- **Data plane:** processing and delivery of packets with local forwarding state
 - Forwarding state + packet header with forwarding decision
 - Tasks: Filtering, buffering, scheduling
- **Control plane:** computing the forwarding state in routers
 - Determines how and where packets are forwarded
 - Tasks: Routing, (automatic) traffic engineering, failure detection/recovery, ...
- **Management plane:** configuring and tuning the network
 - Task: (Manual) traffic engineering, ACL config, device provisioning, ...




The Networking “Planes”



Data Plane


- Also known as the forwarding plane
- Streaming algorithms on packets
 - Matching on some header bits
 - Perform some actions
- Example
 - IP Forwarding
 - Firewall
 - etc...




Destination	Mask	Next Hop
30.0.0.0	255.0.0.0	40.0.0.7
40.0.0.0	255.0.0.0	deliver direct
128.1.0.0	255.255.0.0	deliver direct
192.4.10.0	255.255.255.0	128.1.0.9

Control Plane

- Compute paths the packets will follow
 - Populate forwarding tables
 - Implement routing protocols, which are often calculated in a distributed way
- Example: Link-state routing (OSPF, IS-IS), BGP
 - Flood the entire topology to all nodes
 - Each node computes shortest paths
 - Dijkstra's algorithm

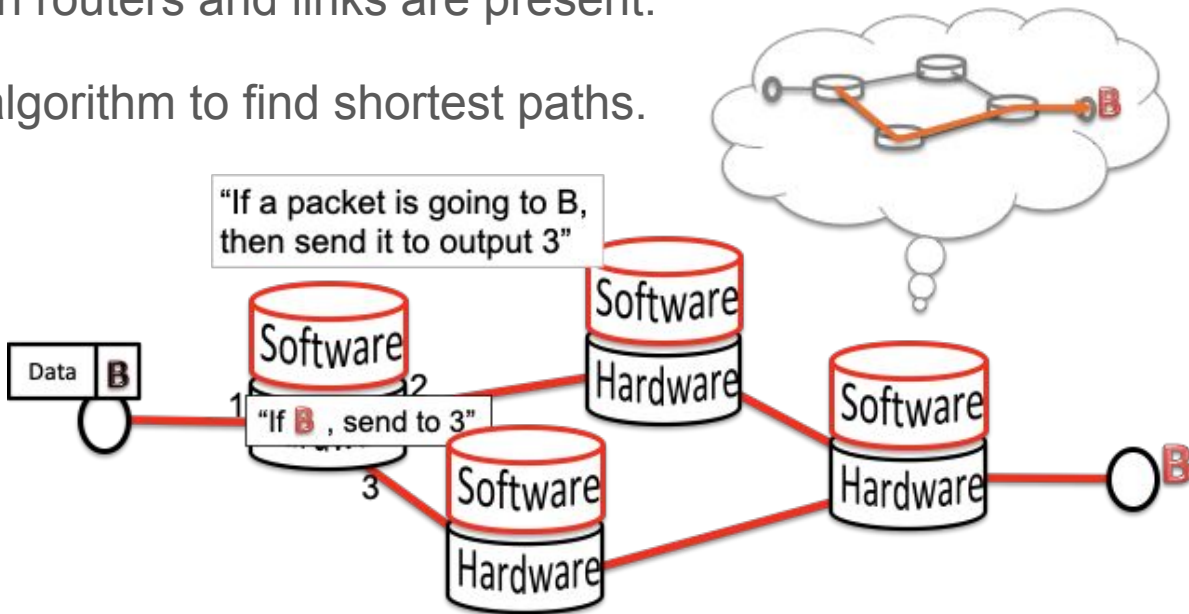


Routing Protocol Example




Routing Protocol Example

- Figure out which routers and links are present.
- Run Dijkstra's algorithm to find shortest paths.



Management Plane

- Traffic Engineering: Setting the weights for each links
 - To be used for the routing protocols
- How to set the right weights?
 - Inversely proportional to link capacity?
 - Proportional to propagation delay?
 - Network-wide optimization based on traffic?

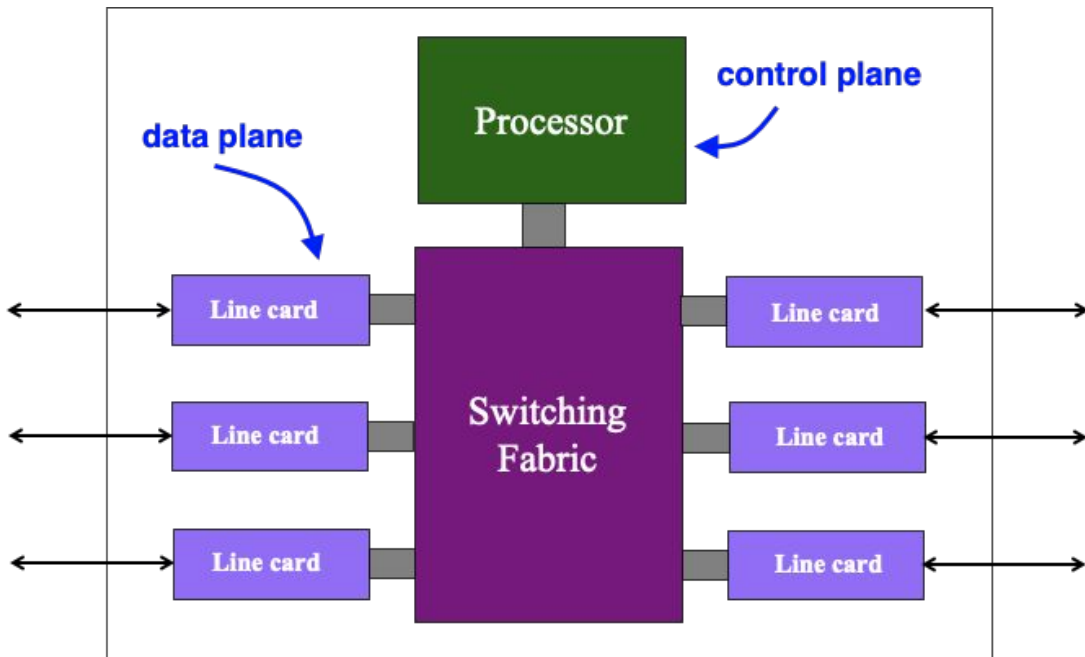




Timescale & Location of Each “Planes”

	Data Plane	Control Plane	Management Plane
Time-scale	Packet (nsec)	Event (10 msec to sec)	Human (min to hours)
Location	Linecard hardware	Router software	Humans or software scripts

Timescale & Location of Each “Planes”





Intro to Software Defined Networking

Challenges in Networking

- Many task-specific protocols and control mechanisms
 - No modularity, limited functionality
- Indirect control mechanisms
 - Must invert protocol behavior, “coax” it to do what you want
 - E.g. Changing weights instead of the actual paths for traffic engineering
- Uncoordinated control mechanisms due to distributed nature
 - Cannot control which router updates first
- Complex interactions between protocols and mechanisms
 - Routing, addressing, access control, QoS



Challenges in Networking

- Therefore, Computer networks are:
 1. Hard to reason about
 2. Hard to evolve as a whole
 3. Expensive to build, operate and manage



An Example: Inter-domain Routing

- Today's inter-domain routing protocols, e.g., BGP, artificially limit routes
 - Routing depend only on **destination IP address blocks**
 - Can only influence **immediate neighbors**
 - Very difficult to incorporate other information
- Application-specific peering
 - Route video traffic one way, and non-video another
- Blocking denial-of-service traffic
 - Dropping unwanted traffic further upstream
- Inbound traffic engineering
 - Splitting incoming traffic over multiple peering links





An Innovation from Stanford

- In 2006, **OpenFlow** was proposed, which provides an open protocol to program the flow-table in different switches and routers.
- In 2007, Nicira was founded by Martin Casado, Nick McKeown and Scott Shenker. This company focuses on software defined networking and network virtualization. It was acquired by VMware in 2012
- In 2007, a SIGCOMM paper named Ethane was published, presenting the initial ideas of **Software Defined Networking (SDN)**.

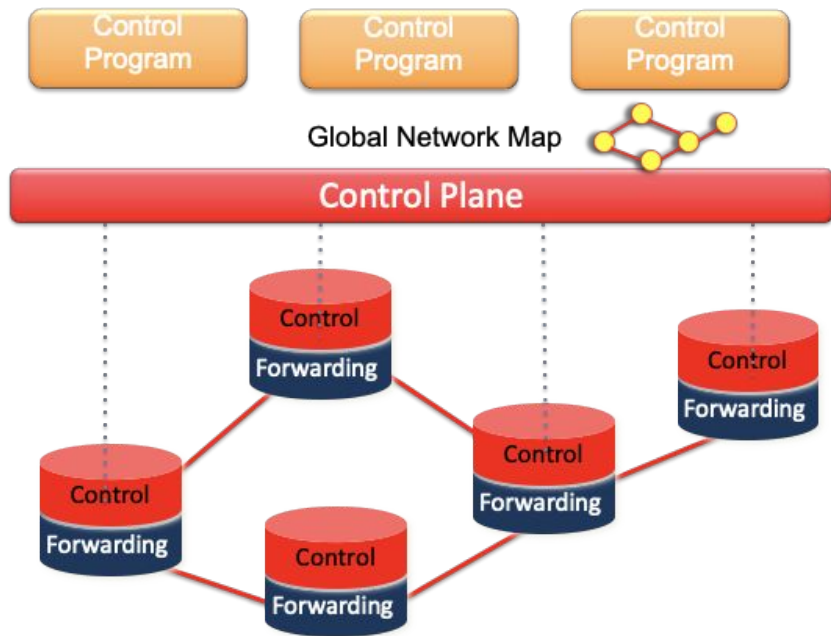


Software Defined Networking Definition

- A network in which the **control plane** is **physically separate** from the **data plane** and a **single (logically centralized) control plane** controls several forwarding devices (data planes).
- In a simple way, having a central, physically separate controller for many routers and switches in the network.




Software Defined Networking Overview



Helpful Analogy: Mainframe Computers



Vertically integrated
Closed, proprietary
Slow innovation
Small industry




Horizontal
Open interfaces
Rapid innovation
Huge industry



Helpful Analogy: Switches and Routers



Vertically integrated
Closed, proprietary
Slow innovation



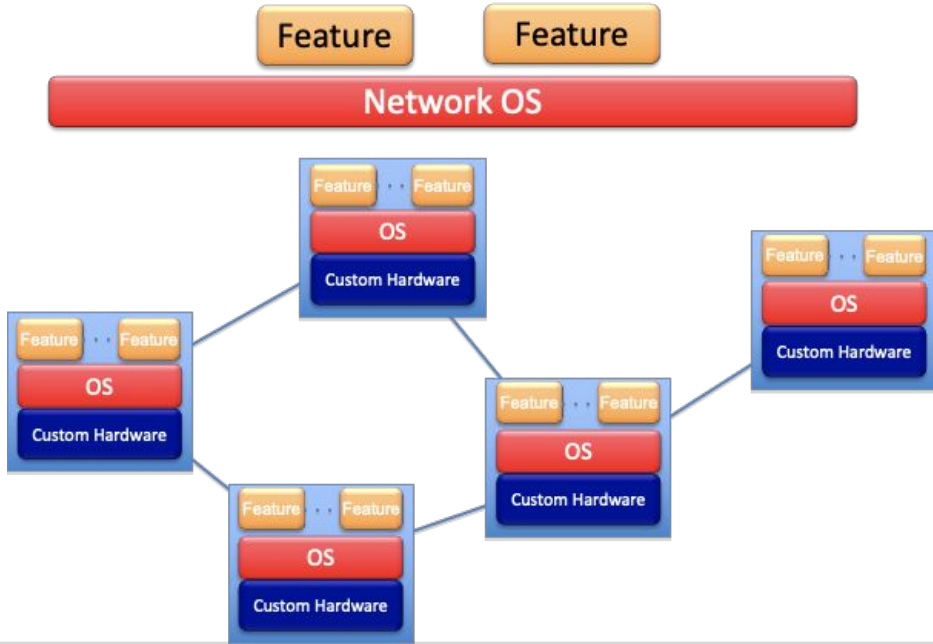
Horizontal
Open interfaces
Rapid innovation

Why is this Better?

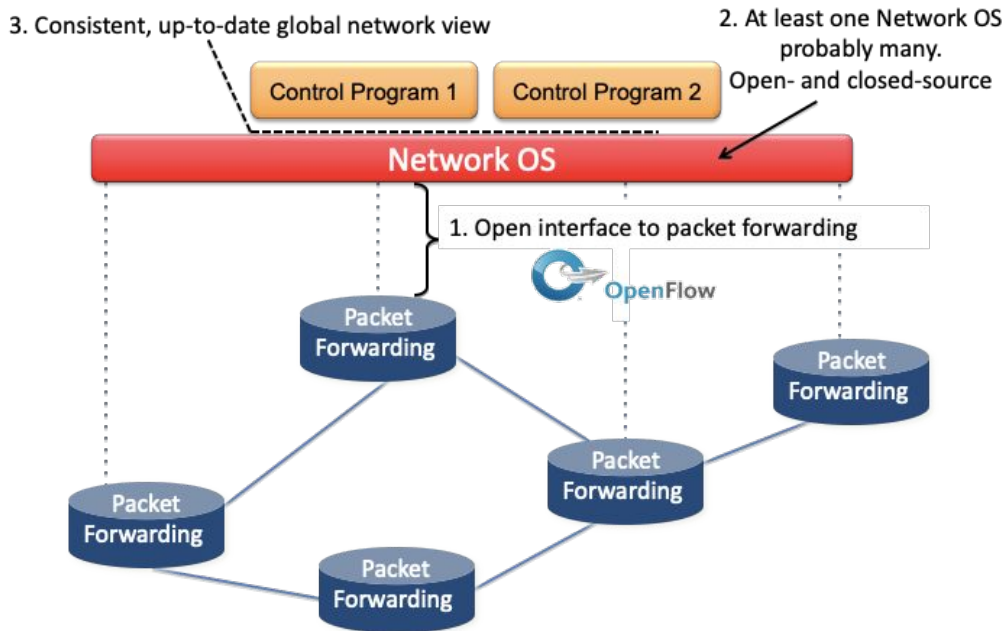
- **Simpler management**
 - No need to “invert” control-plane operations
- **Faster pace of innovation**
 - Less dependence on vendors and standards
- **Easier interoperability**
 - Compatibility only in “wire” protocols
- **Simpler, cheaper equipment**
 - Minimal software
 - Minimal hardware



How SDN Changes the Network



Components of SDN



Forwarding Abstraction (with Open Interface)

- A standard way of defining forwarding state and communicating the state to the hardware
 - Flexible
 - Behavior specified by control plane
 - Built from basic set of forwarding primitives
 - Minimal
 - Streamlined for speed and low-power
 - Control program not vendor-specific
- **OpenFlow** is an example of such an abstraction



Network Operating System

- Network OS: A distributed system that creates a consistent, up-to-date network view
 - Runs on servers (controllers) in the network
 - NOX, ONIX, FBOSS, Floodlight, Trema, OpenDaylight, HyperFlow, Kandoo, Beehive, Beacon, Maestro, ... + more
- Uses forwarding abstraction to:
 - Get state information from forwarding elements
 - Give control directives to forwarding elements



Control Program

- Control program operates on view of network
 - Input: global network view (graph/database)
 - Output: configuration of each network device
- Control program is not a distributed system
 - Abstraction hides details of distributed state




Does SDN Really Simplify the Network?

- Some are skeptical that simplicity can be sustained in the long run
 - Basic paradigm in networks (layers) is in fact a simple model.
 - The ever-changing performance requirements and functionality goals have forced more complexity into network design
- Abstraction doesn't really eliminate complexity
 - NOS, Hypervisor are still complicated pieces of code
- However, SDN achieved the following:
 - Simpler interface for control program (user-specific)
 - Pushed complexity into reusable code (SDN platform)
- Analogous to how compilers are written



Why is SDN so Important?

- Rise of merchant switching silicon
 - Democratized switching
 - Vendors eager to unseat incumbents
- Cloud / Data centers
 - Operators face real network management problems
 - Extremely cost conscious; desire a lot of control
- The right balance between vision & pragmatism
 - OpenFlow compatible with existing hardware
- A “killer app”: Network virtualization



Why is SDN so Important?

- Consider a multi-tenant datacenter
 - Want to allow each tenant to specify virtual topology
 - This defines their individual policies and requirements
- Data center's network hypervisor compiles these virtual topologies into set of switch configurations
 - Takes 1000s of individual tenant virtual topologies
 - Computes configurations to implement all simultaneously
- This is what people are paying money for....
 - Enabled by SDN's ability to virtualize the networks



Practical Challenges in Implementation

- **Scalability**
 - Decision elements responsible for many routers
- **Reliability**
 - Surviving failures of decision elements and routers
- **Response time**
 - Delays between decision elements and routers
- **Consistency**
 - Ensuring multiple decision elements behave consistently
- **Security**
 - Network vulnerable to attacks on decision elements
- **Interoperability**
 - Legacy routers and neighboring domains





OpenFlow

OpenFlow

- In 2006, OpenFlow was proposed, which provides an open protocol to program the flow-table in different switches and routers.
- People can try new routing protocols and security models by a software controller
- <http://ccr.sigcomm.org/online/files/p69-v38n2n-mckeown.pdf>




OpenFlow

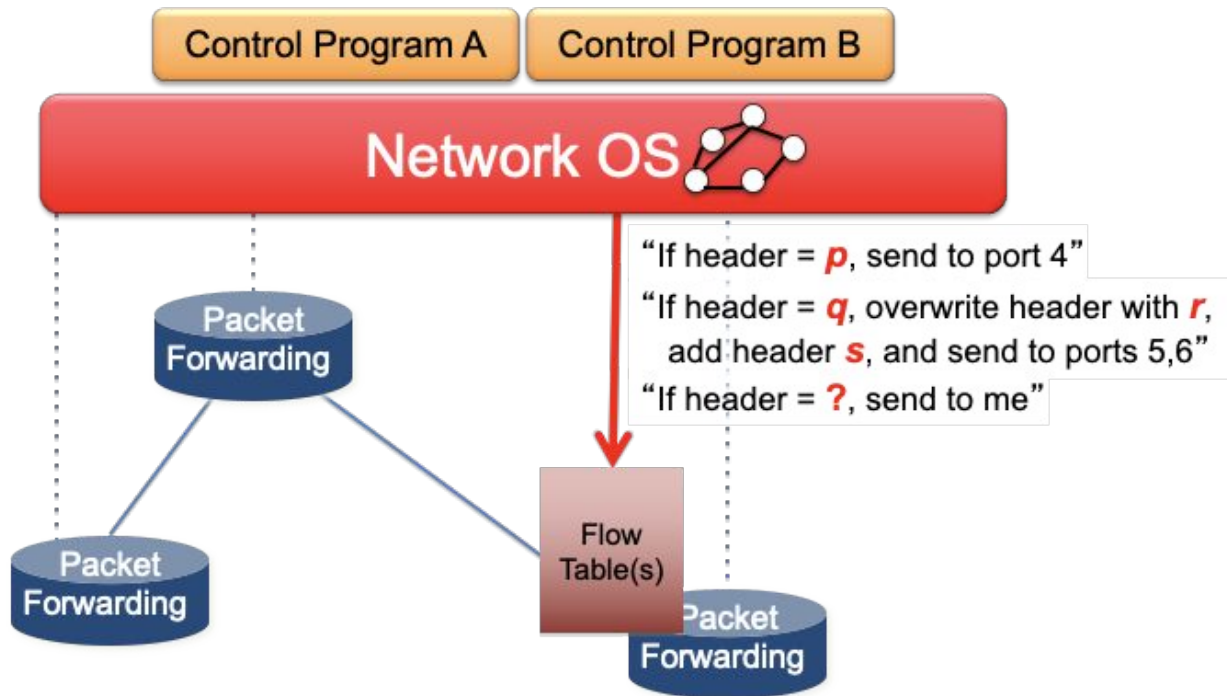
- OpenFlow allows remote management of switch rules
 - OF switches use dedicated network link to a controller
 - Controller is often a 'normal' server, e.g., running Linux
 - Typically for first-time packet forwarding, call out to controller
 - Controller provides resulting packet matching rules & actions
 - Establishes flow to potentially be used for subsequent packets
- OF can be easily implemented within existing switches
 - OF controller can co-exist well with existing control programs



OpenFlow



OpenFlow





TODOs!

- Final Project
- Quiz 3 is out at the end of class
- HW 3 is out





Agenda for Today

- Software Defined Networking
- OpenFlow
- Readings
 - Recommended: <http://yuba.stanford.edu/~casado/ethane-sigcomm07.pdf>
 - Optional:
 - <http://ccr.sigcomm.org/online/files/p69-v38n2n-mckeown.pdf>





Questions?

