



Advanced Operating Systems: Three Easy Pieces

1. Virtualization:

1.1 The CPU

Outline

- **Process and thread:**

- ☐ Abstraction
- ☐ Interlude (intervening)

- **Ensuring Processes cannot harm each other** –
System Calls

- **Scheduling:** separation of policy and mechanism

- **Multiprocessor** scheduling

- **Multi-level Feedback Queue:** MLFQ

- 
- The abstraction
 - Interlude

Process and Threads



The Abstraction: The Process

How to provide the illusion of many CPUs?

■ CPU virtualizing

- ❑ **The OS can promote** the illusion that many virtual CPUs exist.
- ❑ **Time sharing:** Running one process, then stopping it and running another
 - The potential cost (context switch) is a **performance issue**.

A Process

A process is a **running program / program in execution.**

■ A process Comprises of:

❑ **Memory (address space)**

- Instructions
- Data section

❑ **Registers**

- Program Counter (PC)
- Stack Pointer (SP)
-

❑ **Meta data - Data structures** to help the OS manage the process such as **process table entry**, **user_area**, etc.

Process API

- **These APIs are available on any modern OS:**
 - ❑ **Create**
 - Create a new process to run a program
 - ❑ **Destroy**
 - Halt a runaway process
 - ❑ **Wait**
 - Wait for a process to stop running
 - ❑ **Miscellaneous Control (Suspend/Resume)**
 - Some kind of method to suspend a process and then resume it
 - ❑ **Status**
 - Get some status info about a process


Process Creation

1. **Load** a program code into memory, into the address space of the process:
 - ❑ Programs initially reside on disk in *executable format* (*a.out*).
 - ❑ OS performs the loading process **lazily** (i.e., partial load):
 - Loading pieces of code or data (dynamic loader) only as they are needed during program execution.
2. **The program's** run-time **stack** is allocated.
 - ❑ **Use the stack for** *local variables, function parameters, and return address*.
 - ❑ **Initialize the stack with arguments** → as an example, `argc` and the `argv` array of `main()` function

Process Creation (Cont.)

3. **The program's heap** is created.
 - ❑ Used for explicitly requested dynamically allocated memory.
 - ❑ Program request such space by calling **malloc()** and free it by calling **free()**.
4. **The OS** do some other initialization tasks:
 - ❑ **input/output (I/O) setup**
 - Each process by default has three open file descriptors, i.e., Standard **input**, **output**, and **error**
5. **Start the program** running at the entry point, namely **main()**.
 - ❑ **The OS transfers control** of the CPU to the newly-created process.

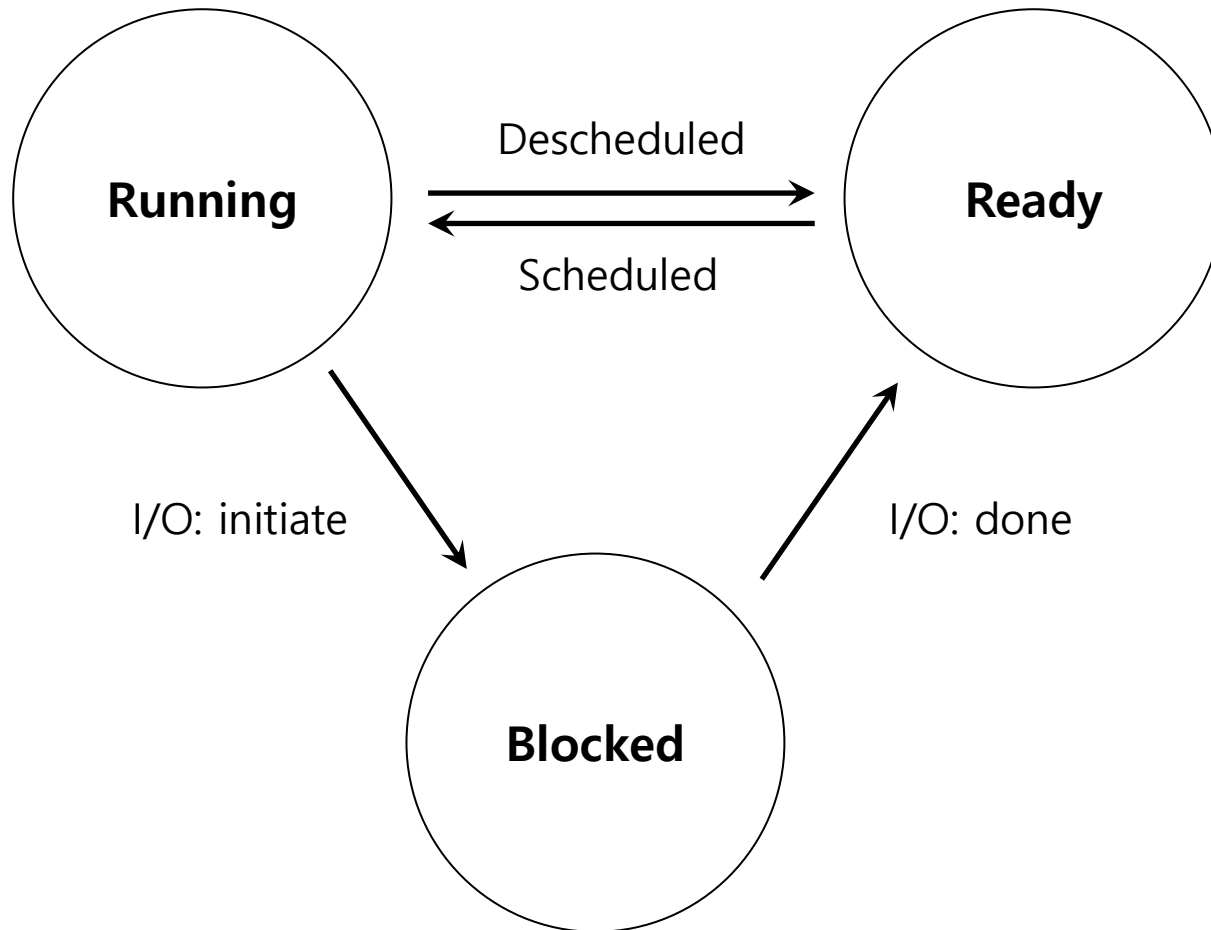
Loading: From Program on Disk to Process in Memory



Process States

- **A process can be in one of three states:**
 - **Running**
 - A process is running on a processor.
 - **Ready**
 - A process is ready to run but for some reason the OS has chosen not to run it at this given moment.
 - **Blocked**
 - A process has performed some kind of I/O operation.
 - When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor.

Process State Transition



Data structures

- **The OS** has some key data structures that track various relevant pieces of information.
 - **Process list/queue**
 - **Ready** processes queue
 - **Blocked** processes queue
 - **Running:** Current running process (comes only from the Ready queue)
 - **Register context**
- **PCB(Process Control Block): Unix Proc[] table**
 - A C-structure that contains information about each process such as pid, ppid, p_signal, address of u_area, etc.

Example) The xv6 kernel Proc Structure

```
// the general-purpose registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;           // Index pointer register
    int esp;           // Stack pointer register
    int ebx;           // Called the base register
    int ecx;           // Called the counter register
    int edx;           // Called the data register
    int esi;           // Source index register
    int edi;           // Destination index register
    int ebp;           // Stack base pointer register
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO (being created), SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

Example) The xv6 kernel Proc Structure (Cont.)

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char    *mem;           // Start of process memory
    uint    sz;             // Size of process memory
    char    *kstack;        // Bottom of kernel stack
                           // for this process
    enum    proc_state state; // Process state
    int     pid;            // Process ID
    struct  proc *parent;    // Parent process
    void    *chan;          // If non-zero, sleeping on chan
    int     killed;         // If non-zero, have been killed
    struct  file *ofile[NOFILE]; // Open files
    struct  inode *cwd;      // Current directory
    struct  context context; // Switch here to run process
    struct  trapframe *tf;   // Trap frame for the
    ...      // current interrupt
};
```



Interlude: Process API

The fork() System Call

■ Create a new process

- The newly-created process has its own copy of the **address space, registers, and PC.**

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();          // rc is a pid
    if (rc < 0) {              // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {     // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {                  // parent goes down this path (main)
        printf("hello, I am parent of %d (pid:%d)\n",
               rc, (int) getpid());
    }
    return 0;
}
```

Calling fork() example (Cont.)

Result (Not deterministic)

No guarantee if it is the parent or the child will resume execution first?

```
prompt> ./p1  
hello world (pid:29146)  
hello, I am parent of 29147 (pid:29146)  
hello, I am child (pid:29147)  
prompt>
```

or

```
prompt> ./p1  
hello world (pid:29146)  
hello, I am child (pid:29147)  
hello, I am parent of 29147 (pid:29146)  
prompt>
```

The wait() System Call

- This parent won't return until the child has run and exited.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) { // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else { // parent goes down this path (main)
        int wc = wait(NULL); // parent wait for child exit,
                            // wc = child pid that terminated
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
               rc, wc, (int) getpid());
    }
    return 0;
}
```

My child pid Pid of the terminated child My (parent) pid

The wait() System Call (Cont.)

Result (Deterministic)

Parent will wait for the child to exit before it resumes execution...

```
prompt> ./p2
hello world (pid:29266)
hello, I am child (pid:29267)
hello, I am parent of 29267 (wc:29267) (pid:29266)
prompt>
```

The exec() System Call

- Run a program that is different from the calling program:

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {                                // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {                        // child (new process) run different prog
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc");                // program: "wc" (word count)
        myargs[1] = strdup("p3.c");             // argument: file to count content
        myargs[2] = NULL;                       // marks end of array
        ...
    }
```

% wc p3.c

The exec() System Call (Cont.)

(Cont.)

```
...
    execvp(myargs[0], myargs); // runs word count => wc filename
    printf("this shouldn't print out");
} else {
    int wc = wait(NULL);
    // parent goes down this path (main)
    // wc = child pid that exited
    // rc = pid of the child
    // In this cscenario: wc == rc
    printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
           rc, wc, (int) getpid());
}
return 0;
}
```

Result

```
prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
29 107 1030 p3.c
<#lines    #words    #characters    file name>                # output of WC
hello, I am parent of 29384 (wc:29384) (pid:29383)
prompt>
```

All of the above with redirection

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <fcntl.h>
#include <sys/wait.h>

int
main(int argc, char *argv[]){
    int rc = fork();
    if (rc < 0) {                // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child: redirect standard output to a file
        close(STDOUT_FILENO);
        open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU);
        ...
    }
```

All of the above with redirection (Cont.)

(Cont.)

```
...
// now exec "wc"...
char *myargs[3];
myargs[0] = strdup("wc");           // program: "wc" (word count)
myargs[1] = strdup("p4.c");         // argument: file to count
myargs[2] = NULL;                   // marks end of array
execvp(myargs[0], myargs);          // runs word count
} else {                             // parent goes down this path (main)
    int wc = wait(NULL);
}
return 0;
}
```

Result

```
prompt> ./p4
prompt> cat p4.output                # Output of WC
32 109 846 p4.c
<#lines  #words  #characters  file name>
prompt>
```


Operating System Roles?

■ What is a resource and its abstraction?

1. **CPU:** process and/or thread
2. **Memory:** address space
3. **Device/Disk:** files

Physical Resource	Abstraction
CPU	Process / Thread
Memory	Address Space
Disk	Files



Ensuring Processes cannot harm each other – System Calls



Mechanism: Limited Direct Execution

How to efficiently virtualize the CPU with control?

- The OS needs to share the physical CPU between processes by **time sharing**.
- **Issues:**
 - **Performance:** How can we implement virtualization without adding excessive overhead to the system?
 - **Control:** How can we run processes efficiently while retaining control over the CPU?

Direct Execution

- Just run the program directly on the CPU:

OS	Program
<ol style="list-style-type: none">1. Create entry in the process list2. Allocate memory for program3. Load program into memory4. Set up stack with <code>argc / argv</code>5. Clear registers6. Execute call <code>main()</code> <ol style="list-style-type: none">9. Free memory of process10. Remove from the process list	<ol style="list-style-type: none">7. Run <code>main()</code>8. Execute <code>return</code> from <code>main()</code>

**Without *limits* on running programs,
the OS wouldn't be in control of anything and
thus would be "just a library"**

Problem 1: Restricted Operation

- **What if a process** wishes to perform some kind of restricted operation such as ...
 - ❑ Issuing an I/O request to a disk
 - ❑ Gaining access to more system resources such as CPU or memory
- **Solution:** Using **protected control** transfer:
 - ❑ **User mode:** Applications do not have direct access to hardware resources.
 - ❑ **Kernel mode:** The OS has access to the full resources of the machine

System Call

- **Allows the kernel** to **carefully expose** certain key pieces of system-level functionality (i.e., provide concrete system-level services through set of **system calls**) to user program, such as:
 - ❑ Accessing the file system
 - ❑ Creating and destroying processes
 - ❑ Communicating with other processes
 - ❑ Allocating more memory
 - ❑ Performing any I/O

System Call (Cont.)

■ Trap instruction

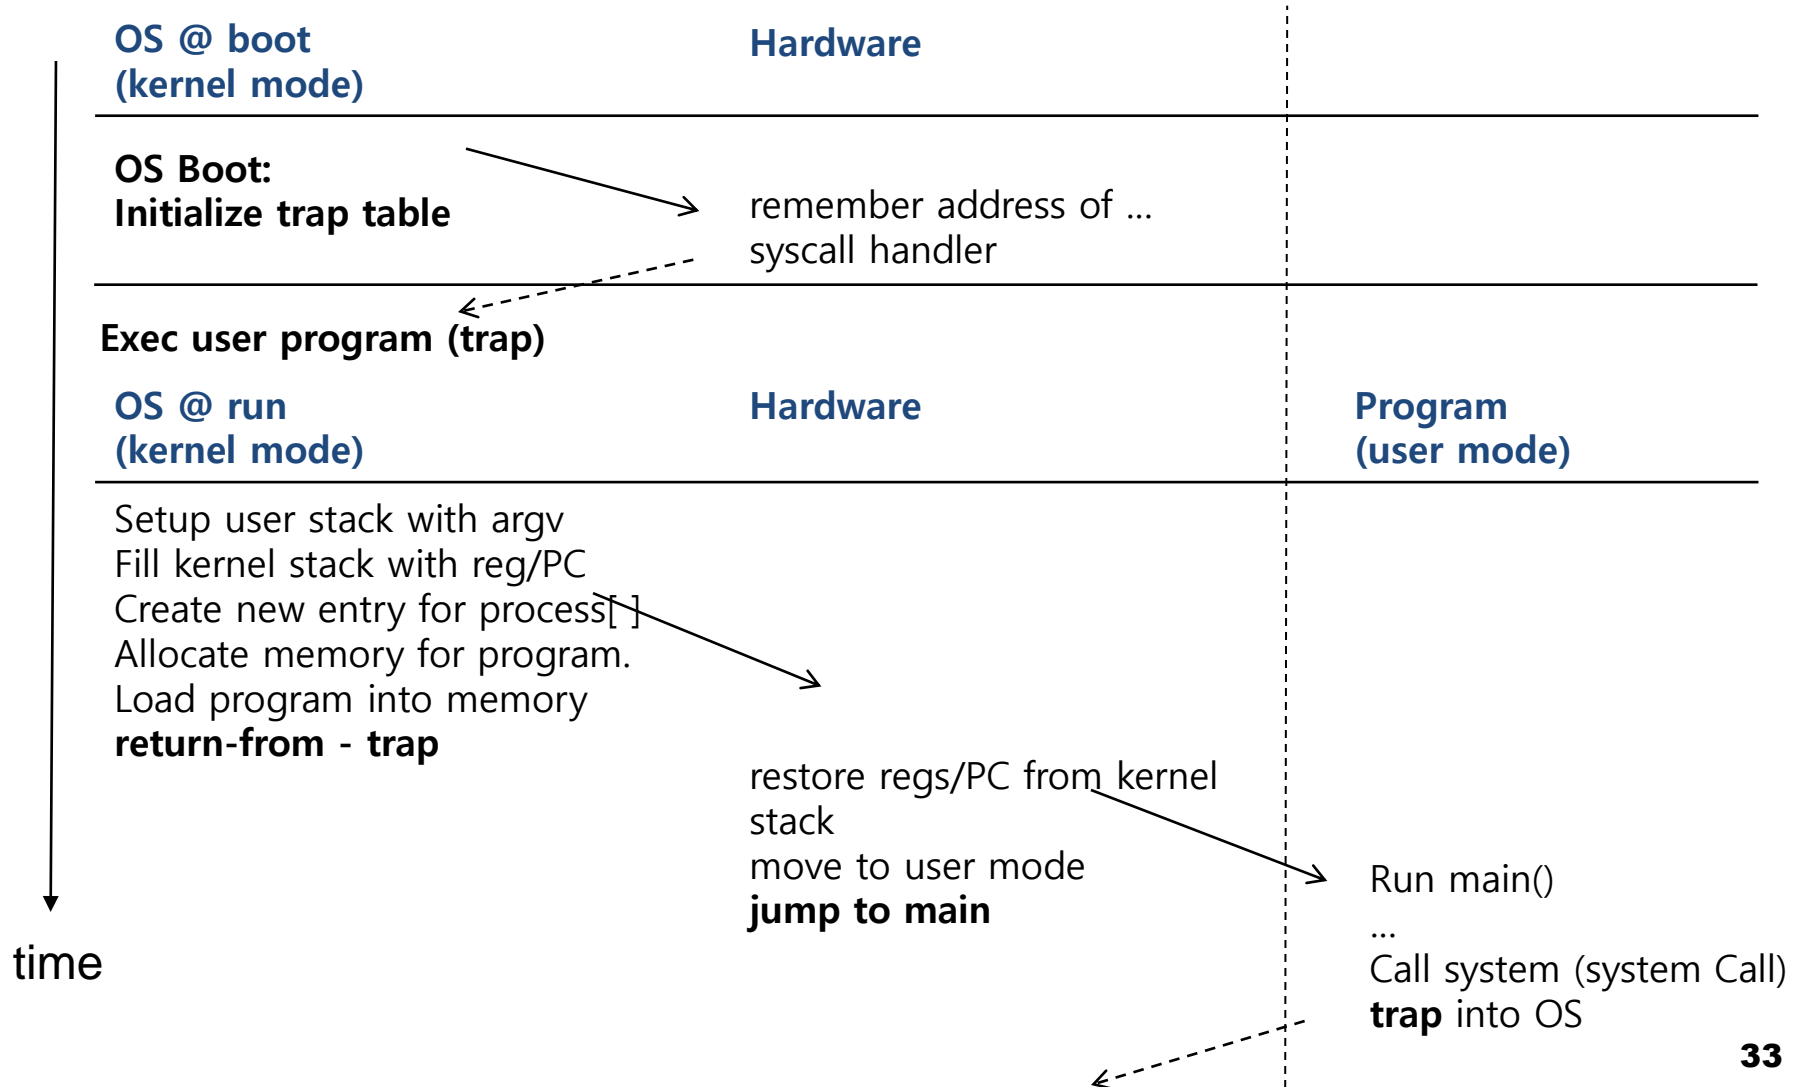
- ❑ **Set up kernel stack** (create stack frame in kernel stack including user-level return address)
- ❑ **Jump into the kernel**
- ❑ **Raise the privilege level** to kernel mode

<Execute the user service request in the kernel>

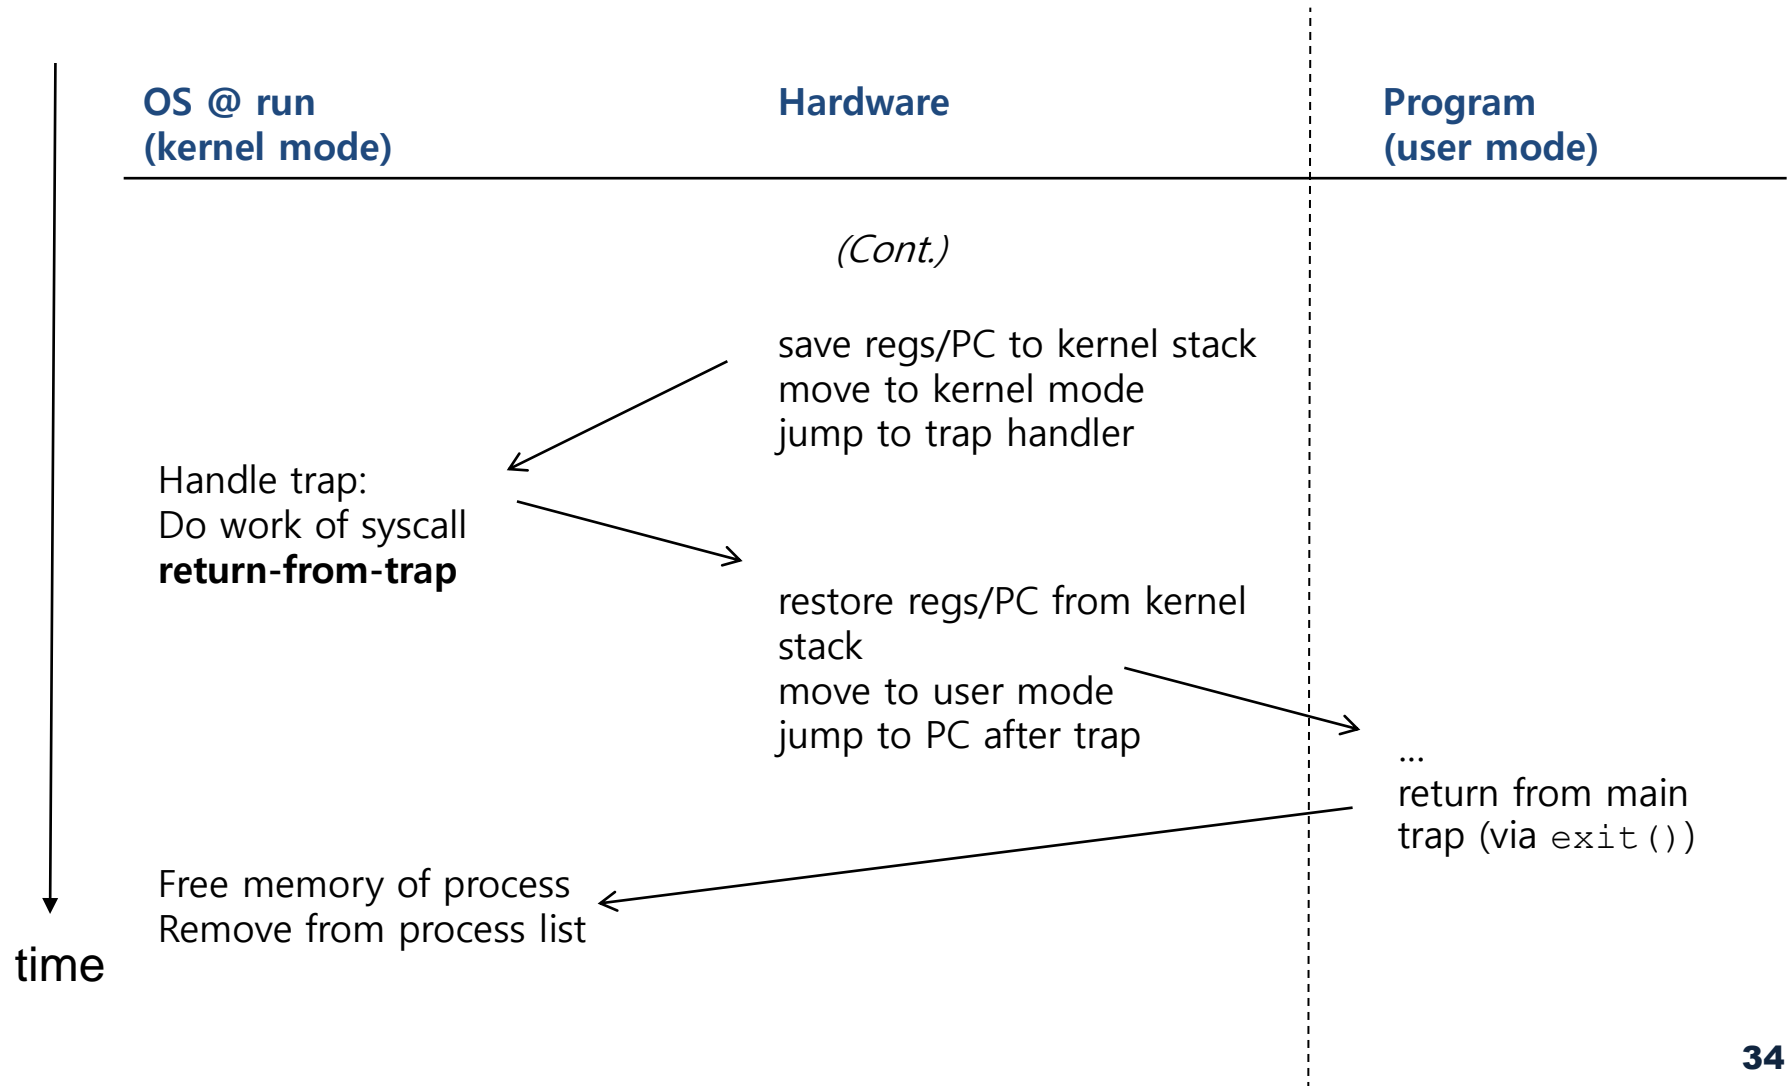
■ Return-from-trap instruction

- ❑ **Return into the calling user program** (extract return address from the kernel stack frame)
- ❑ **Reduce the privilege level** back to user mode

Limited Direction Execution Protocol: Execute (a.out) User program



Limited Direction Execution Protocol: Execute System Call



Problem 2:

Switching Between Processes

- How can the OS **regain control** of the CPU so that it can switch between *processes*?
 - ❑ **A cooperative Approach:** Wait for system calls to be issued by the user
 - ❑ **A Non-Cooperative Approach:** The OS takes control unilaterally

A cooperative Approach: Wait for system calls

- **Processes periodically give up the CPU** by making **system calls** such as **yield**:
 - ❑ **The OS decides** to run some other task.
 - ❑ **Application also transfer control to the OS** when they do something illegal.
 - Divide by zero (exception)
 - Try to access memory that it shouldn't be able to access (exception)
 - ❑ **Example for Wait for system calls**: Early versions of the Macintosh OS, The old Xerox Alto system

A process gets stuck in an infinite loop – not giving up the CPU?
→ **Reboot the machine!**

A Non-Cooperative Approach: OS Takes Control

■ A timer interrupt

- ❑ **During the boot sequence**, the OS starts the timer.
- ❑ **The timer** raise an interrupt every so many milliseconds repeatedly
- ❑ **When the interrupt is raised:**
 - The currently running process is halted.
 - Save enough of the state of the program to be able to resume in the future
 - A pre-configured interrupt handler in the OS runs.

A timer interrupt gives OS the ability to run and schedule the CPU.

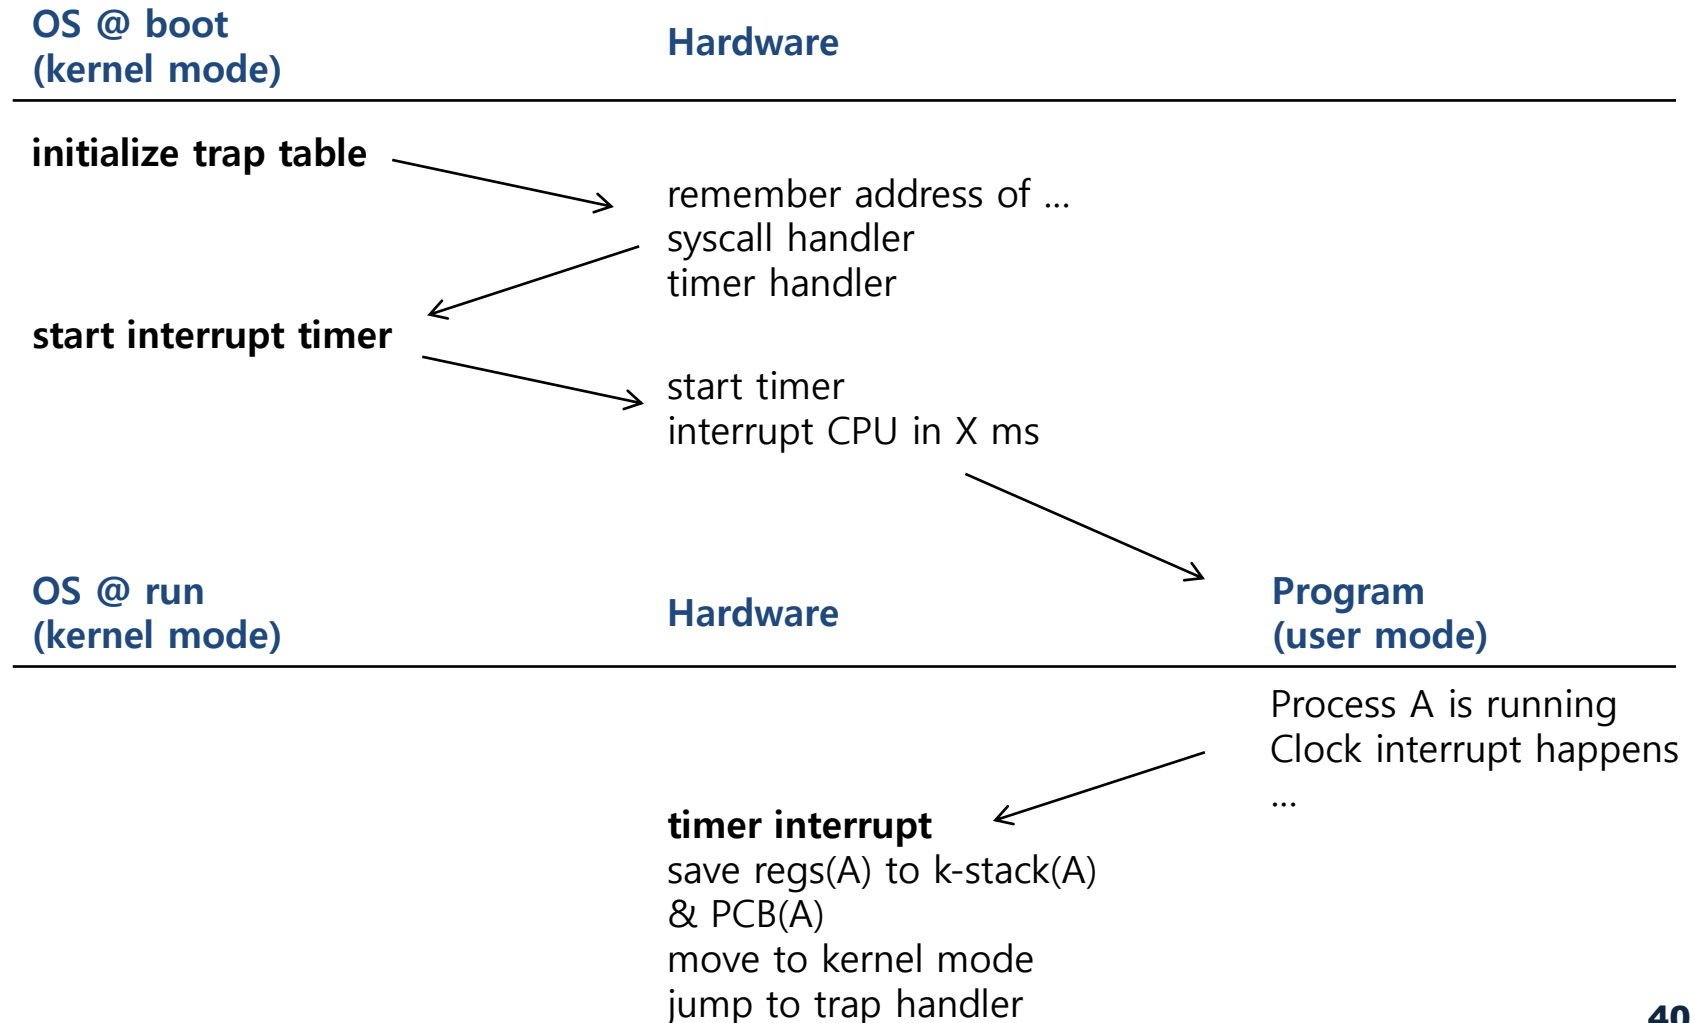
Saving and Restoring Context

- Scheduler makes a decision (on return from the clock interrupt but before resuming the current user-level process), either to:
 - ❑ **Whether to continue/resume** running the **current process**, or switch to a **different one** (how)?
 - ❑ **If the decision is made to switch**, the OS executes context switch.

Context Switch

- **A low-level piece of assembly code:**
 - ❑ **Save few register values** for the current process into its **kernel stack** and its **PCB**:
 - General purpose registers
 - PC
 - kernel stack pointer (SP)
 - ❑ **Restore few register values** for the soon-to-be-executing process from its **kernel stack** and its **PCB**
 - ❑ **Switch to the kernel stack** for the soon-to-be-executing process

Limited Direction Execution Protocol (Timer interrupt)

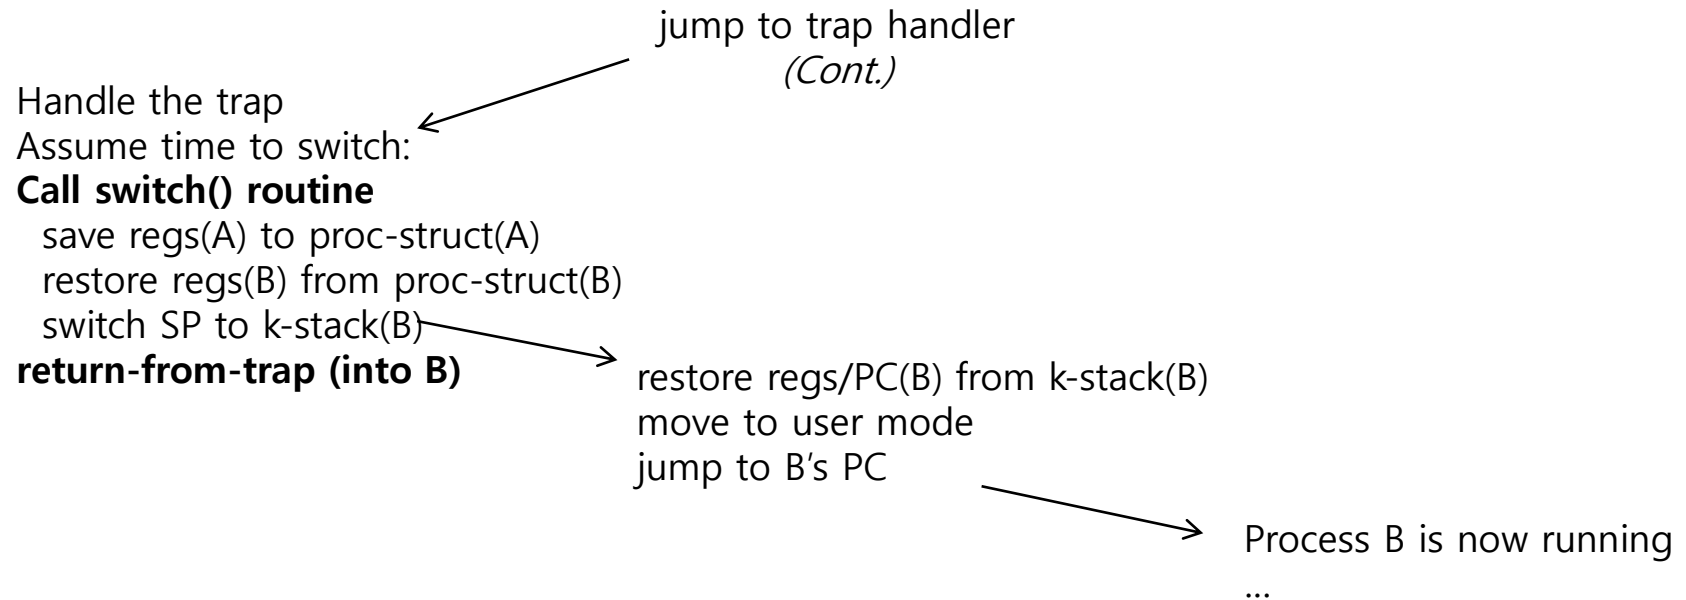


Limited Direction Execution Protocol (Timer interrupt)

OS @ run
(kernel mode)

Hardware

Program
(user mode)



The xv6 Context Switch Code

```
1 # void switch(struct context **old, struct context *new);
2 #
3 # Save current register context in old
4 # and then load register context from new.
5 .globl switch
6 switch:
7     # Save old registers (Save)
8     movl 4(%esp), %eax           # put old ptr into eax
9     popl 0(%eax)                # save the old IP
10    movl %esp, 4(%eax)           # and stack
11    movl %ebx, 8(%eax)           # and other registers
12    movl %ecx, 12(%eax)
13    movl %edx, 16(%eax)
14    movl %esi, 20(%eax)
15    movl %edi, 24(%eax)
16    movl %ebp, 28(%eax)
17
18    # Load new registers (restore)
19    movl 4(%esp), %eax           # put new ptr into eax
20    movl 28(%eax), %ebp          # restore other registers
21    movl 24(%eax), %edi
22    movl 20(%eax), %esi
23    movl 16(%eax), %edx
24    movl 12(%eax), %ecx
25    movl 8(%eax), %ebx
26    movl 4(%eax), %esp          # stack is switched here
27    pushl 0(%eax)               # return addr put in place
28    ret                         # finally return into new ctxt
```

Worried About Concurrency?

- **What happens if**, during interrupt or trap handling, another interrupt occurs?
- **OS handles these situations:**
 - **Disable interrupts** during interrupt processing ← single CPU
 - Use a number of sophisticated **locking** schemes to protect concurrent access to internal data structures ← SMP



Scheduling: separation of Policy and Mechanism



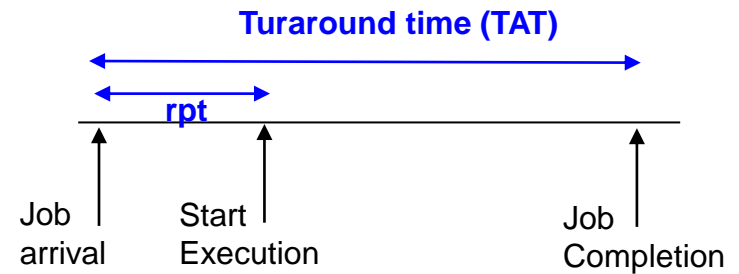
Scheduling: Introduction

Scheduling: Introduction

■ Workload assumptions (for simplicity):

1. **Early UNIX** assigns **fixed time slice** to a process.
Linux uses Fair Scheduler where each process is assigned a **proportion time** slice (time slice depends on the load and is not fixed value).
2. **All jobs arrive at** the same time.
3. **All jobs only use** the **CPU** (i.e., they perform no I/O).
4. **The run-time** of each job is known/tracked.

Scheduling Metrics



■ Performance metric: Response time

- The time at which **the job start execution** minus the time at which **the job arrived** in the system.

$$T_{\text{response time}} = T_{\text{first run time}} - T_{\text{arrival time}}$$

■ Performance metric: Turnaround time

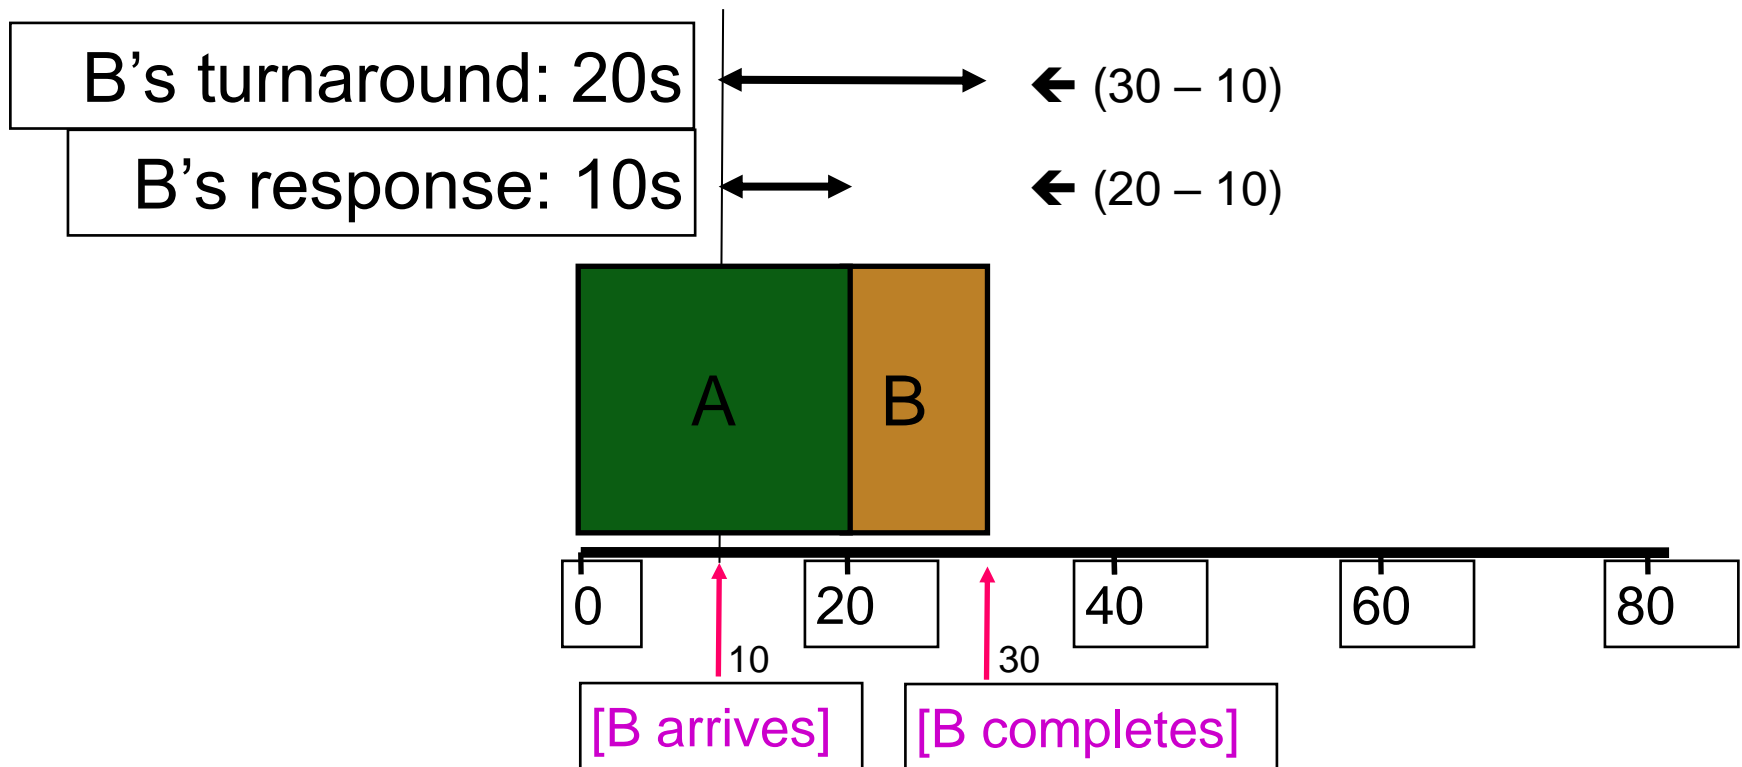
- The time at which **the job completes** minus the time at which **the job arrived** in the system.

$$T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$$

■ Waiting time metric: Waiting time

$$T_{\text{waiting}} = T_{\text{turnaround}} - T_{\text{service-time}}$$

Scheduling Metrics



■ Another metric is fairness:

- ❑ **Performance** (RDBMS) and **fairness** (OS) are often at odds in scheduling.


First In, First Out (FIFO)

■ First Come, First Served (FCFS)

□ **Very simple** and easy to implement

■ Example:

1. **Assume each job** runs for 10 seconds.
2. **A arrived** just before B which arrived just before C.




Avg TAT = 20
sec , while
service time is
10 sec for A, B
and C

$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Why FIFO is not that great? – Convoy effect

- **Let's relax assumption-1 (pg 46):** Each job **no longer** runs for the same amount of time.
- **Example:**
 - ❑ **A runs** for 100 seconds, B and C run for 10 each.
 - ❑ **A arrived** just before B which arrived just before C.



Avg TAT = 110 sec , while service time is 10 sec for B and C

$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec}$$


Shortest Job First (SJF)

- **Run the shortest job first**, then the next shortest, and so on:

- ☐ **Non-preemptive scheduler**

- **Example:**

- ☐ A runs for 100 seconds, B and C run for 10 each.
- ☐ A arrived just before B which arrived just before C.



Avg TAT = 50 sec, while service time is 10 sec for B and C and 100 sec for A


$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

SJF with Late Arrivals from B and C

- **Let's relax assumption 2 (pg 46):** Jobs can arrive at any time.

- **Example:**

- A arrives at $t=0$ and needs to run for 100 seconds.
- B and C arrive at $t=10$ and each need to run for 10 seconds.



Avg TAT = 103 sec, while service time is 10 sec for B and C

$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$

Shortest Time-to-Completion First (STCF)

- **Add preemption to SJF:**

- ☐ Also known as **Preemptive Shortest Job First (PSJF)**


- **A new job enters the system:**

- ☐ Determine between the remaining time of current jobs and the new arrived job
- ☐ Schedule the job which has the least time left

Shortest Time-to-Completion First (STCF)

■ Example:

- A arrives at $t=0$ and needs to run for 100 seconds.
- B and C arrive at $t=10$ and each need to run for 10 seconds



Avg TAT = 50
sec instead of
103 sec while
service time is
10 sec for B and
C

$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$

New scheduling metric: Response time

- The time from when the job arrives to the first time it is scheduled to run.

$$T_{response} = T_{first\ run} - T_{arrival}$$

- STCF (Shortest To Complete First) and related disciplines are not particularly good for response time.

How can we build a scheduler that is
sensitive to response time?

Round Robin (RR) Scheduling


■ Time slicing Scheduling

- ❑ **Run a job** for a time slice (say 1 sec) and then switch to the next job in the **run queue** until the jobs are finished.
 - **Time slice** is sometimes called a scheduling quantum.
- ❑ **It repeatedly** does so until all jobs are finished.
- ❑ **The length of a time slice** must be *a multiple of* the timer-interrupt period (say time slice = 20 msec).

RR is fair, but performs poorly on metrics
such as turnaround time

RR Scheduling Example


- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.



$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5sec$$

$$T_{TAT} = (5+10+15)/3 = 5\ sec$$

SJF (Bad for Response Time)



$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1sec$$

$$T_{TAT} = (13+14+15)/3 = 14\ sec$$

RR with a time-slice of 1sec (Good for Response Time and Not for Turn-Around Time)

The length of the time slice is critical

■ The shorter time slice (RR)

- ❑ **Better** response time ↑
- ❑ **The cost of context switching** will dominate overall performance ↓

■ The longer time slice (RR)

- ❑ **Amortize** the cost of switching ↑
- ❑ **Worse** response time ↓

Deciding on the length of the time slice presents
a **trade-off** to a system designer

Incorporating I/O with RR


- **Let's relax assumption 3 (pg. 46):** All programs perform I/O
- **Example:**
 - A and B need 50ms of CPU time each.
 - A runs CPU for 10ms periodically and performs an I/O between every two CPU runs, i.e., A runs 5 CPU periods and 4 I/O
 - I/Os each take 10ms
 - B simply uses the CPU for 50ms and performs no I/O
 - The scheduler runs A first, then B after

Incorporating I/O with RR (Cont.)

$$\text{RPT} = (0 + 90)/2 = 45$$

$$\text{TAT} = (90 + 140)/2 = 115$$

$$\text{CPU Util} = 100/140 = 71.5\%$$



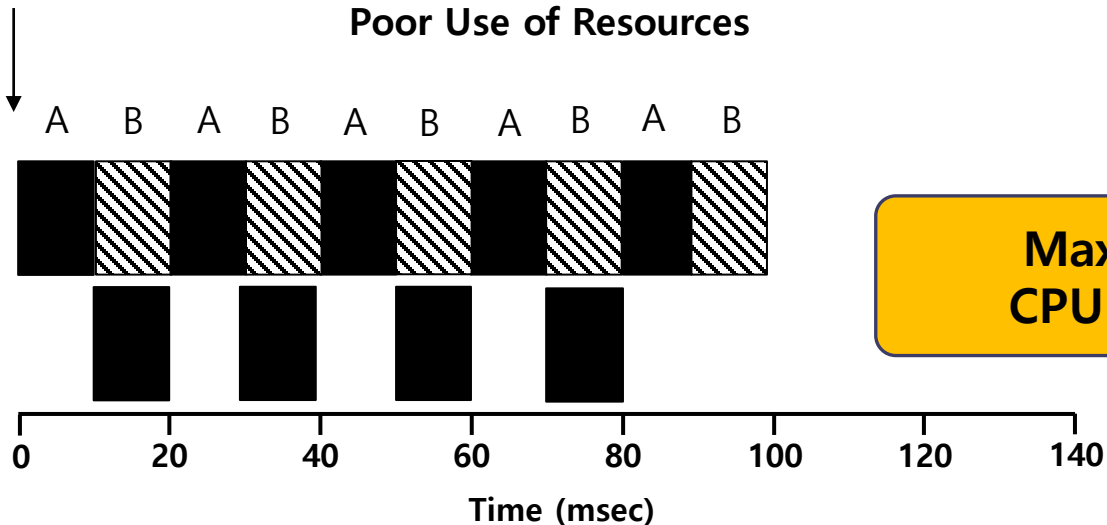
A, B arrives

Poor Use of Resources

$$\text{RPT} = (0 + 10)/2 = 5$$

$$\text{TAT} = (90 + 100)/2 = 95$$

$$\text{CPU Util} = 100/100 = 100\%$$



Maximize the CPU utilization

Overlap Allows Better Use of Resources

Incorporating I/O with RR (Cont.)

■ When a job initiates an I/O request:

- ❑ The job is blocked waiting for I/O completion.
- ❑ The OS scheduler should schedule another job on the CPU.

■ When the I/O completes:

- ❑ An interrupt is raised.
- ❑ The OS moves the process from blocked queue back to the end of the ready queue.



Scheduling: Proportional Share

Proportional Share Scheduler

■ Fair-share scheduler:

- ❑ **Guarantee** that each job obtain *a certain percentage* of CPU time.
- ❑ **Not optimized for** turnaround time or response time

Basic Concept

■ Tickets:

- ❑ Represent the share of a resource that a process should receive
- ❑ The percent of your tickets to total tickets represents your share of the system resource in question.

■ Example:

- ❑ **There are two processes, A and B and 100 total tickets:**
 - Process A has 75 tickets → receive 75% of the CPU
 - Process B has 25 tickets → receive 25% of the CPU

Lottery scheduling

■ The scheduler picks a winning ticket:

- Load the state of that *winning process* and run it.

■ Example:

- There are 100 tickets

- Process A has 75 tickets: 0 ~ 74 (out of 100)

- Process B has 25 tickets: 75 ~ 99 (out of 100)

Generate random number between 0 - 99

Scheduler's winning tickets: 63 85 70 39 76 17 29 41 36 39 10 99 68 83 63

Resulting scheduler: A B A A B A A A A A A B A B A

A:B = 11:4

The longer these two jobs compete,
The more likely they are to achieve the desired percentages.

Ticket Mechanisms

■ Ticket currency:

- ❑ A user allocates tickets among their own jobs in whatever currency they would like.
- ❑ The system converts the currency into the correct global value.
- ❑ **Example:**
 - There are 200 tickets (Global currency)
 - User (Process type A) has 100 tickets & 2 processes (A1, A2)
 - User (Process type B) has 100 tickets & 1 process (B1)

User A $\rightarrow 50$ (A's currency) to A1 $\rightarrow 50$ (global currency) $\rightarrow 25\%$
 $\rightarrow 50$ (A's currency) to A2 $\rightarrow 50$ (global currency) $\rightarrow 25\%$

User B $\rightarrow 100$ (B's currency) to B1 $\rightarrow 100$ (global currency) $\rightarrow 50\%$

Ticket Mechanisms (Cont.)

■ Ticket transfer:


- ❑ A process can temporarily hand off *its tickets* to another process.

■ Ticket inflation:

- ❑ A process can temporarily raise or lower the number of tickets it owns.
- ❑ If any one process needs *more CPU time*, it can boost its tickets.

Implementation

- **Example:** There are three processes, A, B, and C.
 - **Keep the processes in a list:**



```
1 // counter: used to track if we've found the winner yet
2 int counter = 0;
3
4 // winner: use some call to a random number generator to
5 // get a value, between 0 and the total # of tickets (400)
6 int winner = getrandom(0, totaltickets);
7
8 // current: use this to walk through the list of jobs
9 node_t *current = head;
10
11 // loop until the sum of ticket values is > the winner
12 while (current) {
13     counter = counter + current->tickets;
14     if (counter > winner)
15         break; // found the winner (i.e., current)
16     current = current->next;
17 }
18 // 'current' is the winner: schedule it...
```

Implementation (Cont.)

■ U: unfairness metric

- The time the first job completes divided by the time that the second job completes. ← For equal service time $TAT_1/TAT_2 = 1$

■ Example: Assume two jobs arrive at the same time:

- **There are two jobs, each jobs has runtime 10:**
 - First job finishes at time 10
 - Second job finishes at time 20 // it means strict sequential execution
- $U = \frac{10}{20} = 0.5$
- U will be close to 1 when both jobs finish at nearly the same time.


Unfairness metric is between → 0.5 – 1.0

TAT is “Turn Around Time”

Lottery Fairness Study

■ There are two jobs:

- Each jobs has the same number of tickets (100).



When the job length (service time) is not very long, average unfairness can be **quite severe (number/small number)**.

Stride Scheduling

Ticket
Stride = 1/ticket
Pass = sum strides

Run process with
smallest pass

- **Stride** of each process (opposite or inverse of # of tickets)
 - **Stride** = (A large number) / (the number of tickets of the process)
 - **Example:** Assume a large number = 10,000
 - Process A has 100 tickets → stride of A is $(10,000/100) = 100$
 - Process B has 50 tickets → stride of B is $(10,000/50) = 200$
- Select process with lowest “**Pass**” value (measure for how long you ran = sum of strides process ran so far) to run (take it out of the queue) a quantum time; “Pass” is a measure that enables processes with high tickets (i.e., low stride) to run more frequently
- Run the selected process
- When done with the quantum, **increment process “Pass” value by its “stride” value** and put it back in the queue

```
current = remove_min(queue);           // pick client with minimum Pass
schedule(current);                     // use resource for quantum
current->pass += current->stride;       // compute new pass for current proc using stride
insert(queue, current);                // put back into the queue
```

A pseudo code implementation

Stride Scheduling Example

Stride = 10,000 / tickets

	Pass(A) tickets = 100 (stride=100)	Pass(B) Tickets = 50 (stride=200)	Pass(C) Tickets = 250 (stride=40)	Who Runs?
	0	0	0	A
	100	0	0	B
	100	200	0	C
	100	200	40	C
	100	200	80	C
	100	200	120	A
	200	200	120	C
	200	200	160	C
	200	200	200	...

Pass

A:B:C = 2:1:5
Perfect !

If new job enters with pass value 0,
It will **monopolize** the CPU!



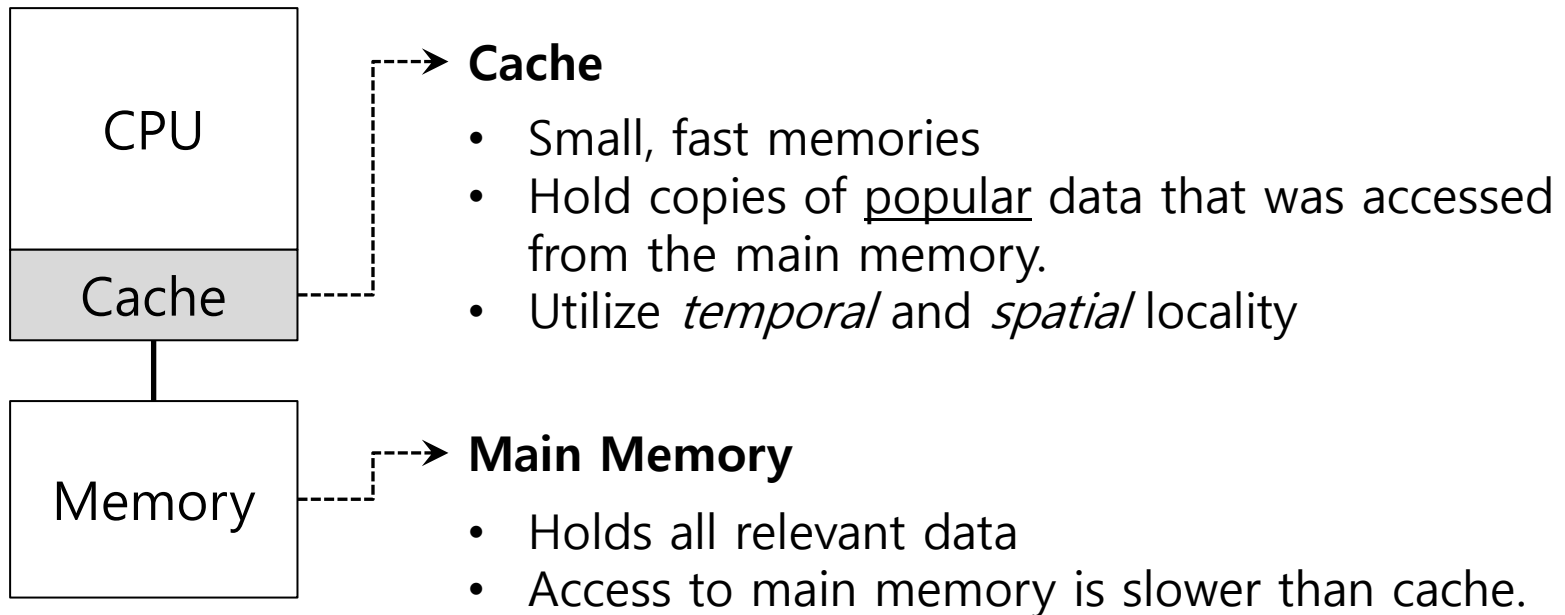
Multiprocessor (SMP) Scheduling

Multiprocessor Scheduling

- The rise of the **multicore processor** is the source of multiprocessor-scheduling proliferation.
 - **Multicore**: Multiple CPU/Cores are packed onto a single chip and sharing memory.
- Adding more CPUs does not make that single application run faster → You'll have to rewrite application to run in parallel, using **threads**.

How to schedule jobs on **Multiple CPUs**?

Single CPU with cache




By keeping data in cache, the system can make slow memory
appear to be a fast one

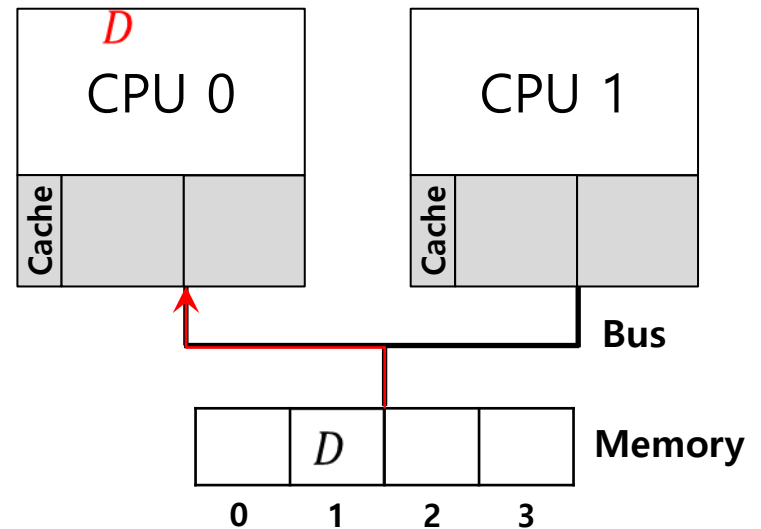
Cache coherence

- Consistency of shared resource data stored in multiple caches.

0. Two CPUs with caches sharing memory

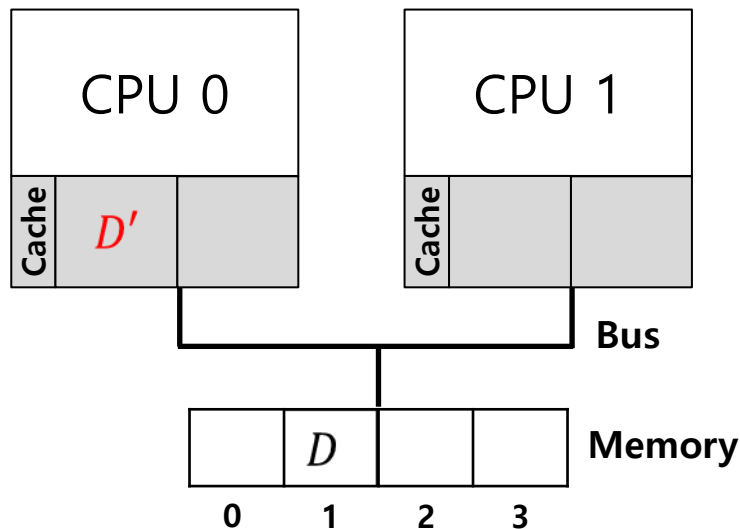


1. CPU0 reads a data at address 1.

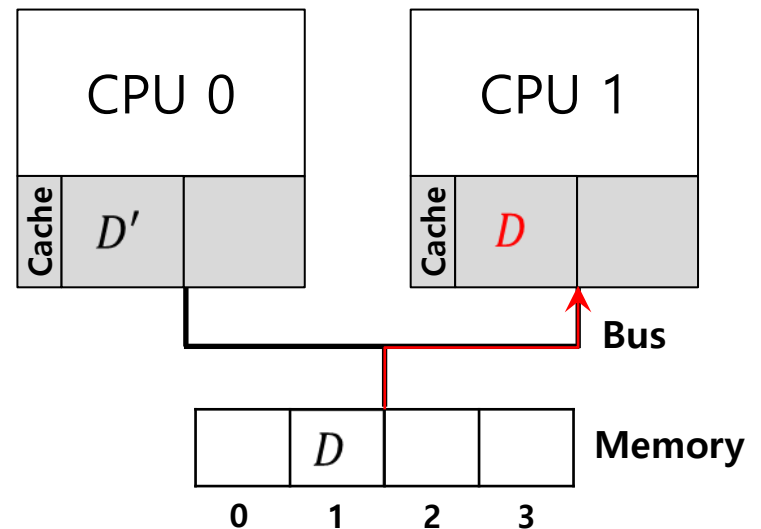


Cache coherence (Cont.)

2. D is updated and CPU1 is scheduled.



3. CPU1 re-reads the value at address 1

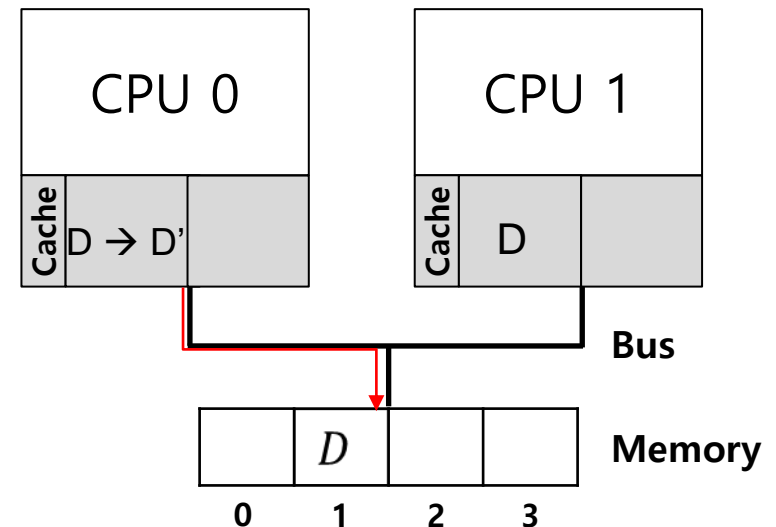


CPU1 gets the **old value D** instead of the correct value D' .

Cache coherence solution

■ Bus snooping:

- ❑ **Each cache** pays attention to main memory updates by **observing the bus**.



- ❑ **When a CPU (CPU1)** sees an update for a data item (CPU 0 updating RAM) that it (CPU 1) holds in its cache, it will notice the change and either invalidate its copy or update it.

Don't forget synchronization


- When accessing shared data across CPUs, **mutual exclusion** primitives should be used to guarantee correctness.

```
1      typedef struct __Node_t {
2          int value;
3          struct __Node_t *next;
4      } Node_t;                                // linked list
5
6      int List_Pop() {
7          Node_t *tmp = head;                  // remember old head ...
8          int value = head->value;              // ... and its value
9          head = head->next;                    // advance head to next pointer
10         free(tmp);                            // free old head
11         return value;                         // return value at head
12     }
```

Simple **Delete** to 1st entry in a linked-list Code –
Incorrect code (concurrency)

Don't forget synchronization (Cont.)

■ Solution



```
1      pthread_mutex_t m;                                // lock for the whole linked list
2
3      typedef struct __Node_t {
4          int value;
5          struct __Node_t *next;
6      } Node_t;                                          // linked list
7
8      int List_Pop() {
9          lock(&m)
10         Node_t *tmp = head;                            // remember old head ...
11         int value = head->value;                        // ... and its value
12         head = head->next;                              // advance head to next pointer
13         free(tmp);                                     // free old head
14         unlock(&m)
15         return value;                                  // return value at head
16     }
```

Simple List Delete Code with lock


Cache Affinity

- **Keep the process** on the same CPU during execution if at all possible:
 - ❑ **A process typically** builds up a fair bit of state in the cache of a CPU.
 - ❑ **The next time** the process run, it will run faster if some of its state is *already present* in the cache on that CPU.

A multiprocessor scheduler should consider **cache affinity** when making its scheduling decision.

Single queue Multiprocessor Scheduling (SQMS)

- **Put all jobs** that need to be scheduled into **a single queue**:
 - Each CPU simply picks the next job from the globally shared queue.
 - **Cons:**
 - Some form of **locking** have to be inserted → **Lack of scalability**
 - **Cache affinity, i.e., lack of**
 - **Example:**





- **Possible job scheduler across CPUs:**

P.S. Each Process switches between different CPUs constantly after every quantum

CPU0	A	E	D	C	B	... (repeat) ...
CPU1	B	A	E	D	C	... (repeat) ...
CPU2	C	B	A	E	D	... (repeat) ...
CPU3	D	C	B	A	E	... (repeat) ...

Scheduling Example with Cache affinity

- ❑ Preserving affinity for most (i.e., except E)




- Jobs A through D are not moved across processors.
- Only job E Migrating from CPU to CPU.
- ❑ Implementing such a scheme can be **complex**.

Multi-queue Multiprocessor Scheduling (MQMS)

- **MQMS** consists of multiple scheduling queues:
 - ❑ Each queue per CPU will follow a particular scheduling discipline.
 - ❑ When a job enters the system, it is placed in **exactly one** scheduling queue.
 - ❑ Avoid/Minimize the problems of information sharing and synchronization.

MQMS Example

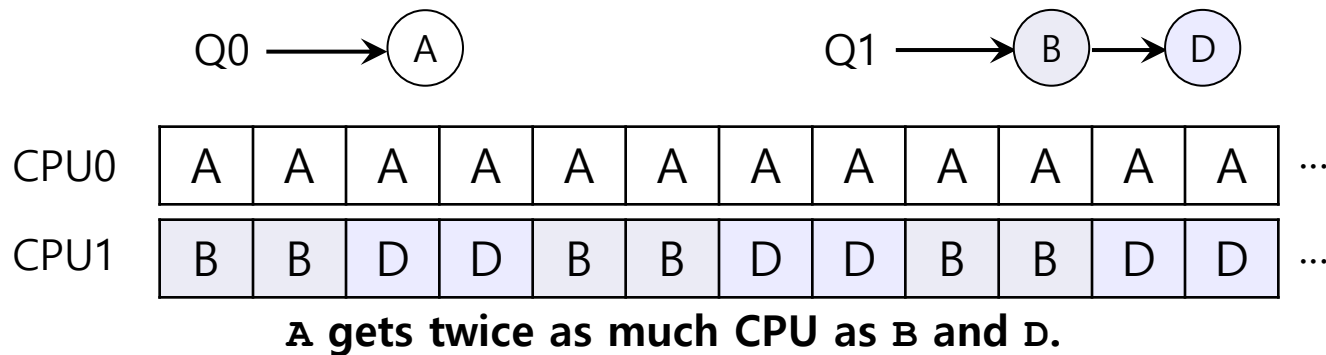
- With designated processes on every processor and use round robin between processes on each CPU, the system might produce a schedule that looks like this:




MQMS provides more **scalability** and **cache affinity**.

Load Imbalance issue of MQMS

■ After job C in Q0 finishes:




■ After job A in Q0 finishes:




How to deal with load imbalance?

- **The answer** is to move jobs (**Migration**) – rebalance the workload:

- **Example:**



The OS moves one of B or D to CPU 0



Or




How to deal with load imbalance? (Cont.)

■ A more tricky-case:



■ A possible migration pattern:

□ Keep switching jobs:



Work Stealing

■ Move jobs between queues:

□ Implementation:

- A source queue that is low on jobs is picked.
- The source queue occasionally peeks at another target queue.
- If the target queue is more-full than the source queue, the source will “**steal**” one or more jobs from the target queue.

□ Cons:

- *High overhead, no CPU affinity, and trouble scaling*

Linux Multiprocessor Schedulers

■ **O(1):**

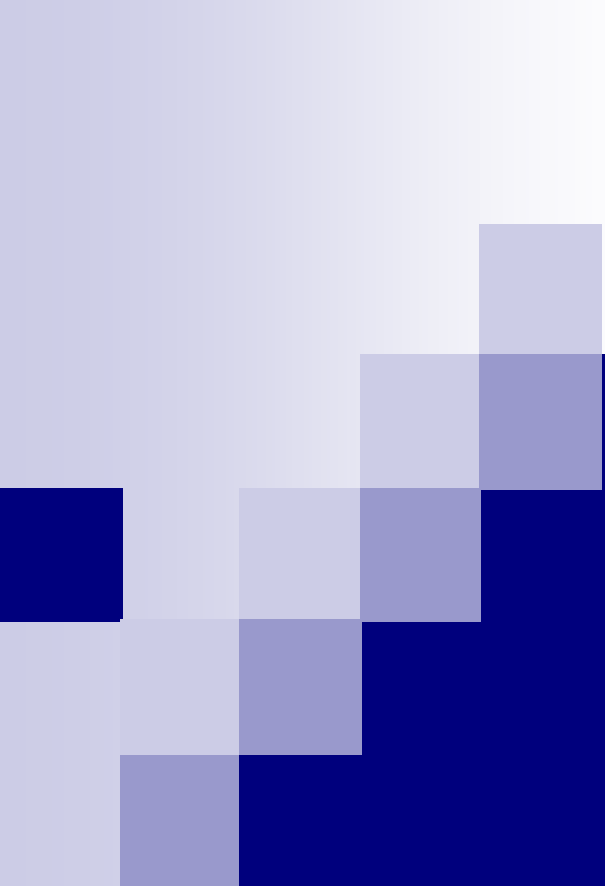
- ❑ A Priority-based scheduler
- ❑ Use Multiple queues
- ❑ Change a process's priority over time
- ❑ Schedule those with highest priority
- ❑ Interactivity is a particular focus

■ **Completely Fair Scheduler (CFS – Linux 2.6.23):**

- ❑ Deterministic proportional-share approach
- ❑ Multiple queues
- ❑ Goal is to maximize CPU utilization in interactive environment; **instead of using a queue use Red/Black tree to order jobs execution**

Linux Multiprocessor Schedulers (Cont.)

- **BF (Brain F) Scheduler (BFS):**
 - ❑ **A single queue** approach
 - ❑ **Proportional-share**
 - ❑ **Based on** Earliest Eligible Virtual Deadline First (EEVDF) – leverage “nice” support
 - ❑ **BFS** has been retired in favor of MuQSS (**M**ultiple **Q**ueue **S**kiplist **S**cheduler - a re-written implementation of the same concept)



Multi-level Feedback Queue Scheduling – Advanced Scheduling


MLFQ: Multi-level Feedback Queue

- **Goal:** general-purpose scheduling
- **Must support** two job types with distinct goals
 - “**interactive**” programs care about **response time**
 - “**batch**” programs care about **turnaround time**
- **Approach:** multiple levels of round-robin;
each level has higher priority than lower levels and preempts them

MLFQ: Priorities

Rule 1: If $\text{priority}(X) > \text{Priority}(Y)$, X runs

Rule 2: If $\text{priority}(X) == \text{Priority}(Y)$, X & Y run in RR



“Multi-level”



How to know how to set priority?

Q1

Approach 1: nice



Approach 2: history “feedback”

MLFQ: Rules

Rule 1: If $\text{priority}(X) > \text{Priority}(Y)$, X runs

Rule 2: If $\text{priority}(X) == \text{Priority}(Y)$, X & Y run in RR

More rules:

Rule 3: Processes start at top priority

Rule 4: If job uses whole slice, demote process
(longer time slices result in lower priorities)

Rule 5: If job do not use the whole time slice (i.e.,
issuing system call or I/O), promote the
process to the next higher queue



END