



Advanced Operating Systems: Three Easy Pieces

Datacenter OS (DC/OS)

Outline

- **The Datacenter needs an OS**
 - Motivation, why, what is needed..
- **What is Apache Mesos**
- **Twitter DC/OS based on Mesos**
 - Master Node
 - Agent Node
 - Frameworks
 - Others



The Datacenter needs an OS

The Datacenter is the new Computer

- Datacenter is running today's most popular consumer applications:
 - Facebook, Google, iCloud, etc.
- Needed (Datacenter) for big data in business & Science
- Datacenter is widely accessible through cloud computing/Internet

Claim: **this new “Datacenter / computer” needs an operating system**

Why Datacenter needs an OS

- **Growing diversity of applications:**

- ❑ **Computing frameworks:** MapReduce, Dryad¹, Pregel², Percolator, Dremel³
- ❑ **Storage systems:** GFS, BigTable, Dynamo, HDFS, etc.

- **Growing diversity of users:**

- ❑ 200+ Hive users at Facebook

- Same reasons regular computers needed one!



1 **Dryad:** [https://en.wikipedia.org/wiki/Dryad_\(programming\)](https://en.wikipedia.org/wiki/Dryad_(programming))
Microsoft: General purpose data parallel applications

2 **Pregel:** https://kowshik.github.io/JPPregel/pregel_paper.pdf
Google: System for large-scale Graph Processing

3. **Dremel:** <https://static.googleusercontent.com/media/research.google.com/en//pubs/archive/36632.pdf>
Google: Interactive Analysis of Web-scale Dataset

What Operating System Provide

Resource Sharing

time-sharing, virtual memory, ...

Data Sharing

files, pipes, IPC, ...

Programming Abstractions

libraries, languages

Debugging & Monitoring

ptrace, DTrace, top, ...

What Operating System Provide

Resource Sharing

time-sharing, virtual memory

Most importantly: **an ecosystem**

...enabling independently developed software to interoperate seamlessly

Debugging & Monitoring

ptrace, DTrace, top, ...

Data
files

ing
ons
ages

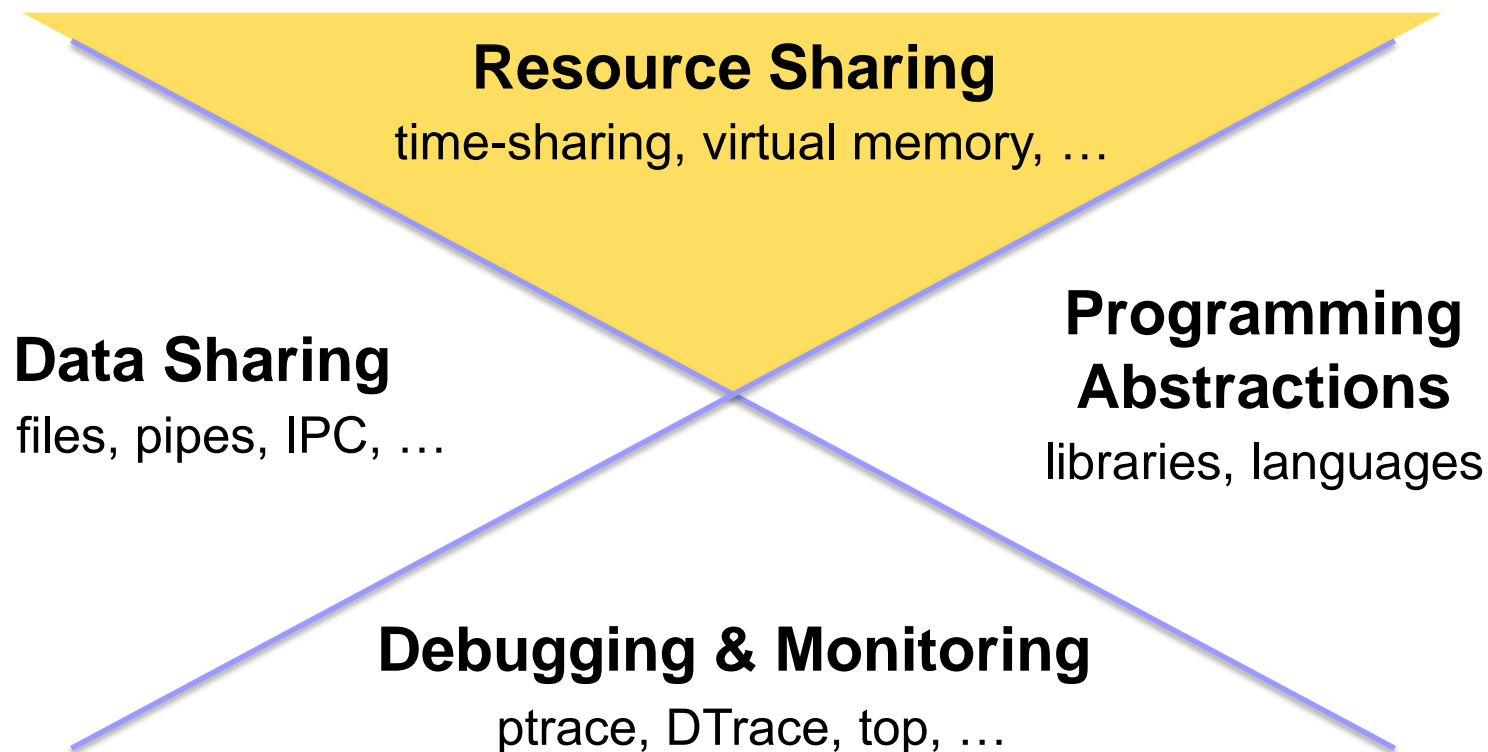
Today's Datacenter Operating System

- Platforms like Hadoop are well-aware of these issues:
 - ❑ **Inter-user resource sharing**, but at the level of MapReduce jobs (though this is changing - YARN)
 - ❑ **InputFormat API for storage systems** (but what happens with the next hot platform after Hadoop?)
- **Other examples:** Amazon services, Google stack

The **problems** motivating a datacenter OS are well recognized, but solutions are **narrowly targeted**

Can researchers take a longer-term view?

Tomorrow's Datacenter OS



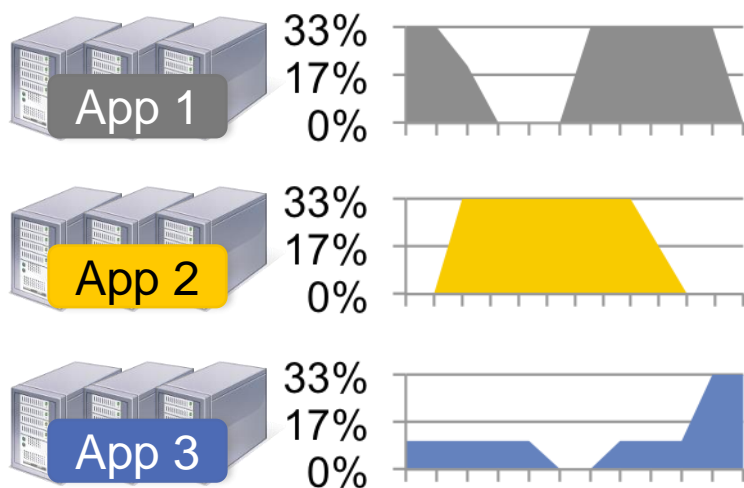
Resource Sharing

“To solve these interaction-problems we would like to have a computer made simultaneously available to many users in a manner somewhat like a **telephone exchange. Each user would be able to use a console at his own pace and without concern for the activity of others using the system.”**


– Fernando J. Corbató, 1962

Resource Sharing

- **Today**, cluster apps are built to run independently and assume they own a fixed set of nodes
- **Result:** inefficient static partitioning
- **What's the right interface for dynamic sharing?**



Static
partitioning



Dynamic
partitioning



Memory Management

- **Memory is an increasingly important resource:**
 - In-memory iterative processing (AllegroGraph “Graph DB based on W3C RDF (Resource Description Framework)”, Pregel, Spark, etc.)
 - HDFS cache for MapReduce cluster could serve 90% of jobs at Facebook (HotOS ‘11)
- **What are the right memory management algorithms for a parallel analytics cluster?**

Programming and Debugging

- Although there are new programming models for applications, **system programming remains hard**:
 - ❑ Can we identify useful common abstractions? (Chubby¹ “lock service for distributed system / Big Table”, Sinfonia², Mesos are some examples)
 - ❑ How much can languages (e.g., Go, Erlang) help?
- **Debugging is very hard**:
 - ❑ Magpie³, X-Trace, Dapper⁴ are some steps here
- Can a clean-slate design of the stack help?

- 1 **Chubby**:
<https://static.googleusercontent.com/media/research.google.com/en//archive/chubby-osdi06.pdf>
- 2 **Sinfonia**: new paradigm to building scalable distributed systems:
<http://www.sosp2007.org/papers/sosp064-aguilera.pdf>
- 3 **Magpie**: online modelling & performance-aware systems:
https://www.usenix.org/legacy/publications/library/proceedings/hotos03/tech/full_papers/barham/barham_html/paper.html
- 4 **Dapper** is a large scale Distributed Systems tracing infrastructure:
<https://research.google.com/pubs/pub36356.html>

What is Needed

- **Focus on paradigms, not only on performance:**
 - Industry is spending a lot of time on performance
- **Explore clean-slate* approaches:**
 - Much datacenter software is written from scratch
 - People using Erlang, Scala, functional models (MR)
- **Bring cluster computing to non-experts:**
 - Most impactful (**datacenter as the new workstation**)
 - **Hard to make a Google-scale stack usable without a Google-scale ops team**

* Absence of existing constraints.



What is Apache Mesos



1. Mesos Overview


1. Overview

- **Mesos**¹ is a platform for Fine-Grained resource sharing in the data center – **Cluster Manager**.
- While there is rapid innovation in cluster computing frameworks, there is no single framework that is optimal for all applications.
- **An Alternative** is to run multiple frameworks in a single cluster, **without static partitioning**, to maximize utilization and to share data between frameworks while providing isolation between the different frameworks ← **Mesos**

1 **Mesos** (Cluster Manager): <http://mesos.apache.org/>


1. Static vs. Dynamic Partitioning

- **Mesos** is a common resource sharing layer over which diverse frameworks can run




1. Static vs. Dynamic Partitioning

Today: static partitioning



Mesos: dynamic sharing



1. Other Benefits

- **Run multiple instances of the *same* framework:**
 - ❑ Isolate production and experimental (testing) jobs
 - ❑ Run multiple versions of the same framework concurrently with complete isolation from each other
- **Build *specialized frameworks* targeting different particular problem domains:**
 - ❑ Better performance than general-purpose abstractions + more flexibility



2. Mesos Goals


2. Mesos Goals

- **High utilization** of resources
- **Support diverse frameworks** (current & future)
- **Scalability** to 10,000's of nodes
- **Reliability** in face of failures
- **Fine-grain sharing**, i.e., one node can run multiple frameworks concurrently
- **Support the “Resource Offers” Model**: simple application-controlled scheduling mechanism

Resulting design: Small microkernel-like core that
pushes scheduling logic to frameworks


2. Fine-Grained Sharing

Coarse-Grained Sharing (HPC):



Storage System (e.g., HDFS)

Fine-Grained Sharing (Mesos):



Storage System (e.g., HDFS)

+ Improved utilization, responsiveness, data locality

2. Resource Offers Model

- **Option:** Global scheduler

- Frameworks express needs in a **specification language**, and a global scheduler matches them to available resources in the cluster

+ Can make optimal decisions

- **Complex:** language must support all framework needs:

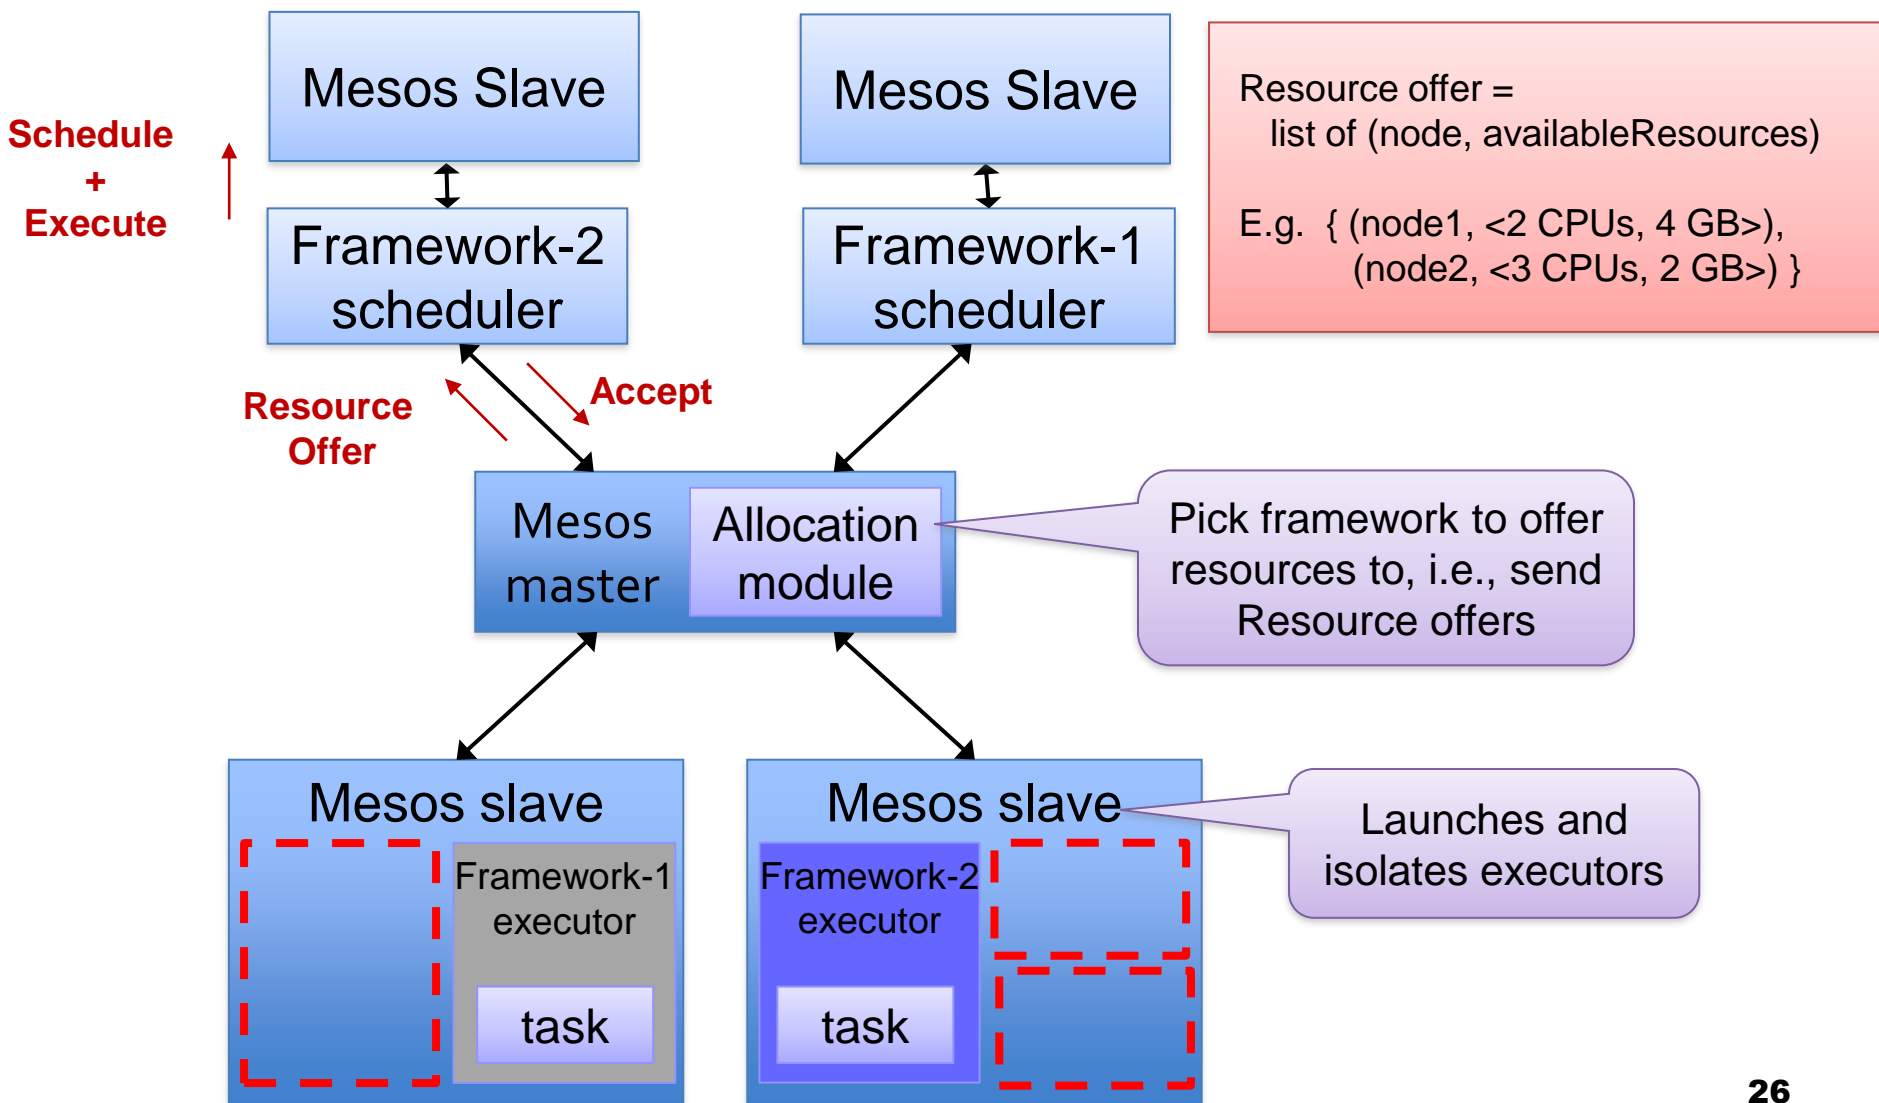
- Difficult to scale and be robust
- Future frameworks may have unanticipated needs

2. Mesos Resource Offers

■ Mesos: Resource offers

- Offer available resources to frameworks, let them pick which resources (up to what is being offered) to use and which tasks a framework will launch accordingly
- + Keeps Mesos simple; makes it possible to support future frameworks
- Decentralized decisions might not be optimal

2. Mesos Architecture





3. Results

3. Mesos vs. Static Partitioning

- Compared performance with statically partitioned cluster where each framework gets 25% of nodes

Framework	Speedup on Mesos
Facebook Hadoop Mix	1.14 X
Large <u>Hadoop</u> Mix	2.10 X
Spark	1.26 X
Torque / MPI	0.96 X


- Overall performance improvement = $(1.14 + 2.1 + 1.26 + 0.96)/4$
= 1.365 → 36.5% improvement

3. Mesos Scalability

- **Mesos only performs** *inter-framework* scheduling (e.g., fair sharing across framework), which is easier than intra-framework (within a framework) scheduling

Result*:

Scaled (almost linear)
to 50,000 emulated
slaves,
200 frameworks,
100K tasks (based on
normal distribution
with mean = 30 sec len)




3. Fault Tolerance

- **Mesos master has only *soft state*:** list of currently running frameworks and tasks
- **Rebuild** a new master: frameworks and slaves re-register with the new master after a failure
- **Result: fault detection and recovery in ~10 sec**



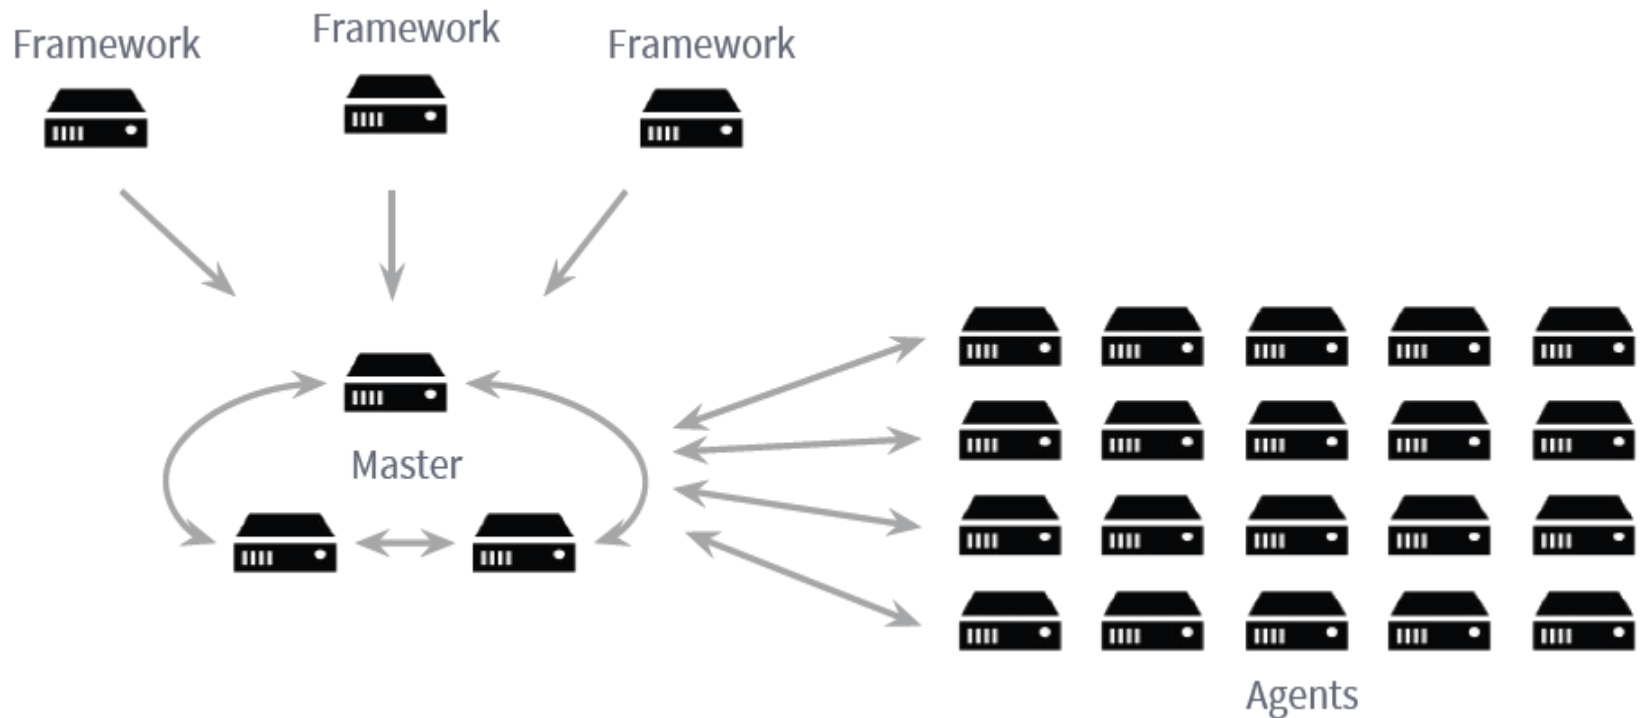
Twitter DC/OS based on Mesos - Aurora

The Datacenter Computer

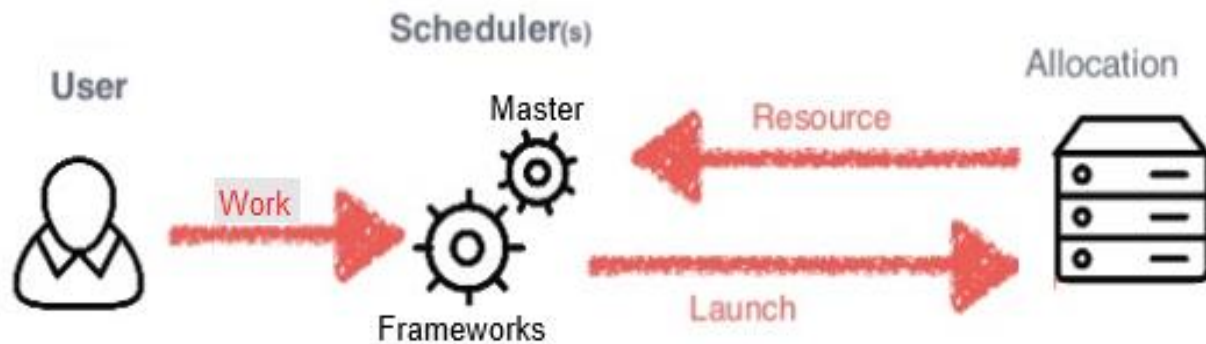


- * Do not confuse Twitter's Aurora with Amazon's Aurora. Amazon Aurora is an RDBMS that is compatible with MySQL and PostgreSQL and is built for the cloud.

Two-levels Scheduling

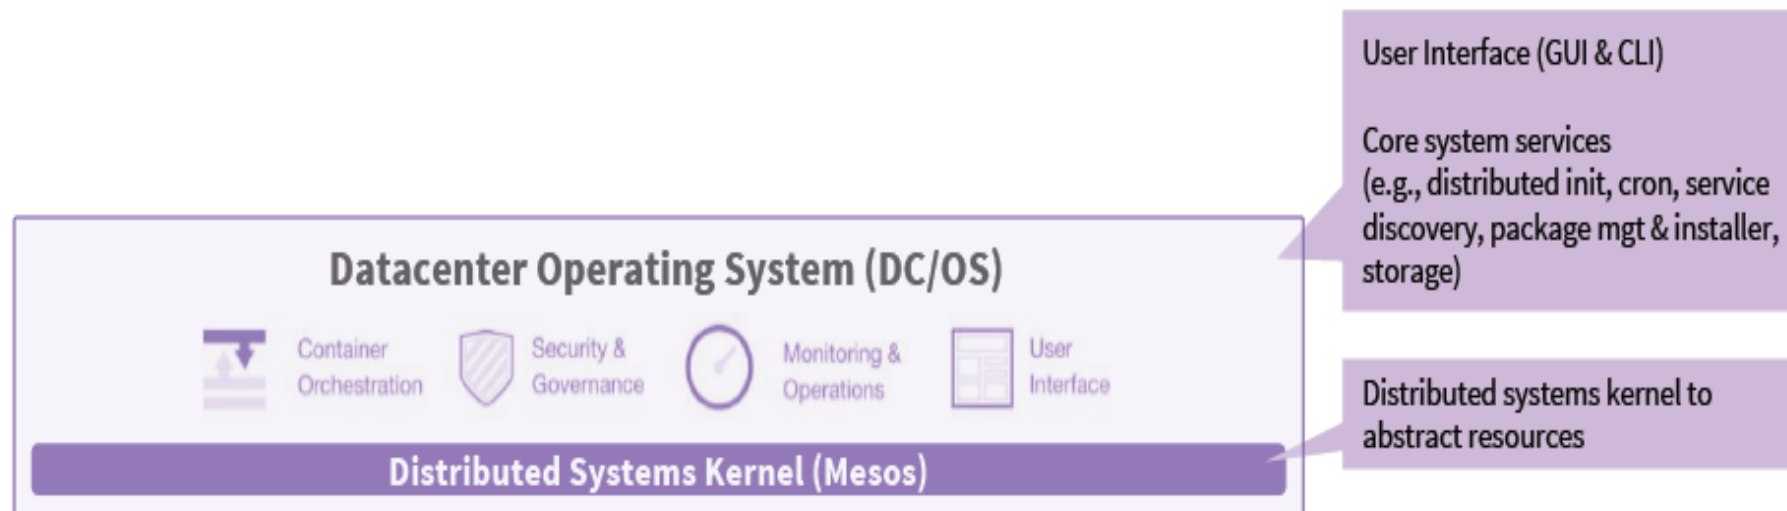


Two-levels Scheduling



1. Master (Mesos+) schedule global resources to a given framework
2. Framework schedule jobs to agents (Mesos agent+)

Datacenter Operating System (DC/OS)



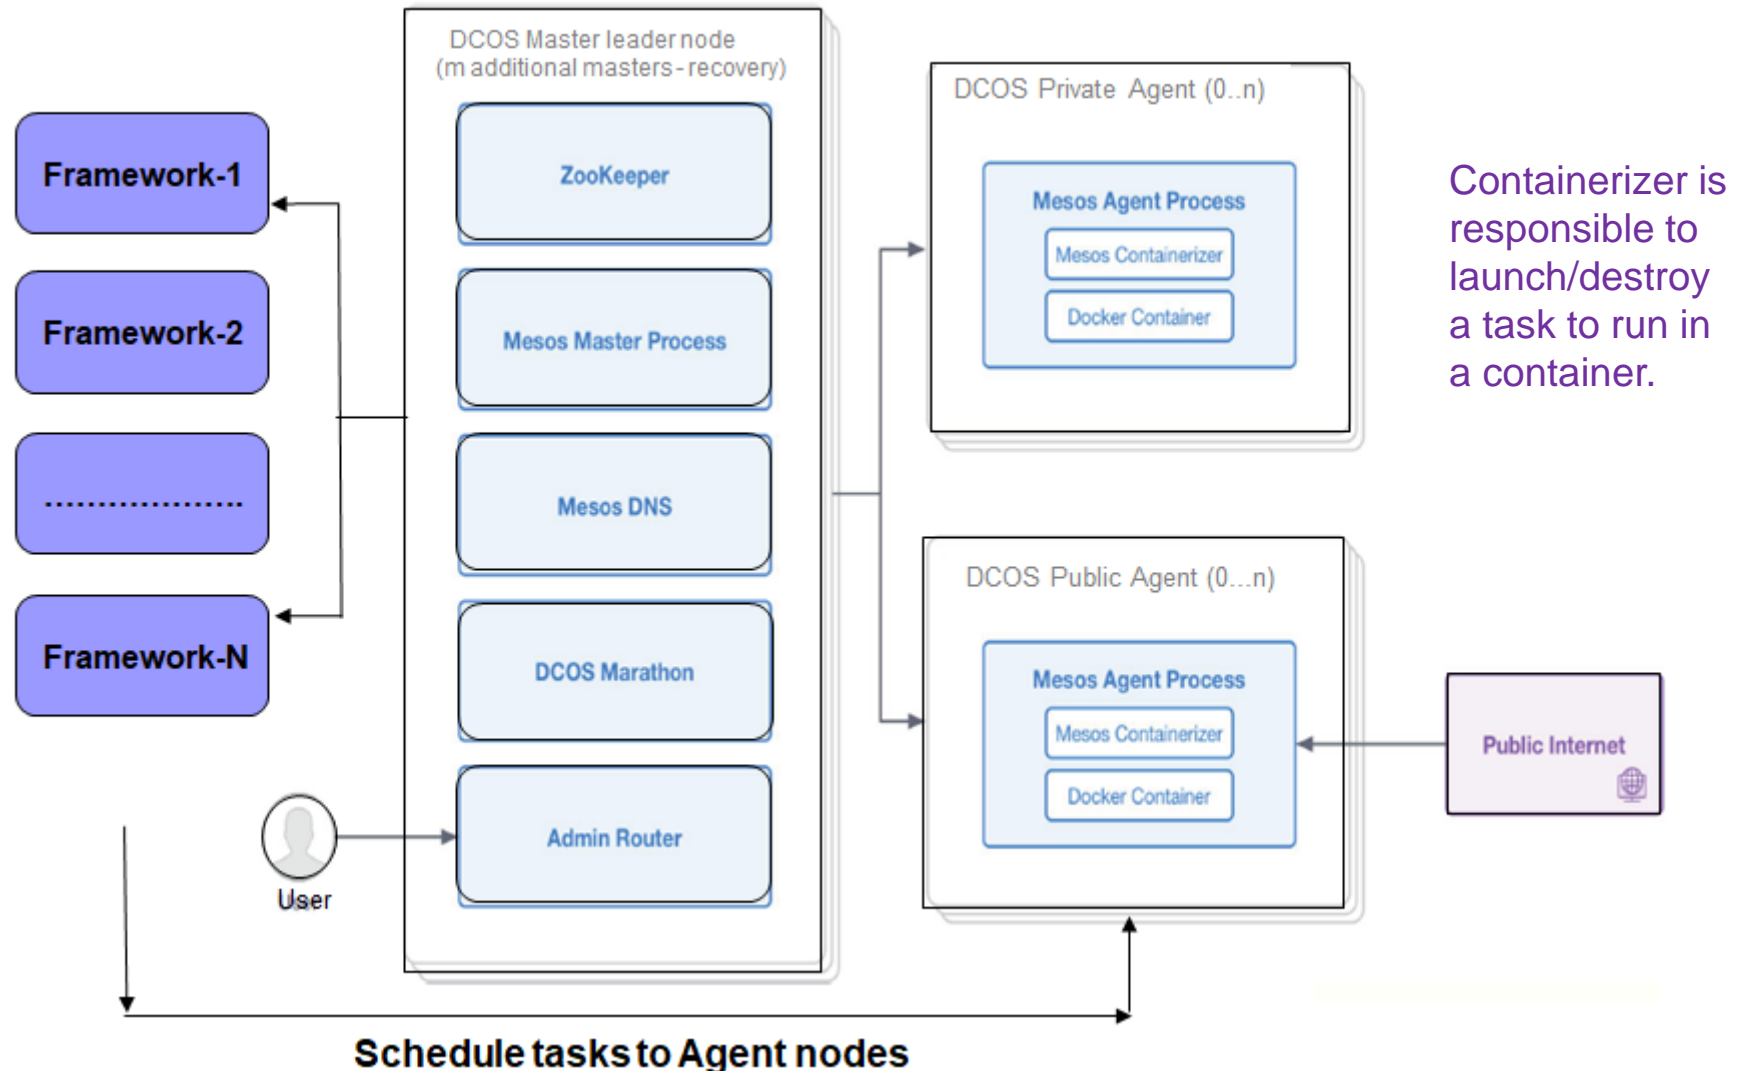
❑ What is included:

- ❖ **Mesos!**
- ❖ API, CLI and GUI
- ❖ Service discovery & Load balancing
- ❖ Storage Volumes
- ❖ Package Manager
- ❖ Installer on Premise and on Cloud

Datacenter OS Architecture

Offers to Frameworks

Resources available (from Agent to Master)



Datacenter OS Features

- ❑ **Kernel == Apache Mesos**, scaling up to 10,000's nodes
- ❑ **Fault tolerance in all components, rolling upgrades throughout**
- ❑ **Containers first class citizens** (LXC, Docker)
- ❑ **Local OS per node** (container enabled)
- ❑ **Scheduling** (long-lived, batch)
- ❑ **Service discovery, monitoring, logging, debugging**

Datacenter Components

1. **Master Node:** one master is assigned as the leader, manage the metadata for the whole system and propagate the metadata to other masters using quorum replication, i.e., not a single point of failure like Hadoop.
2. **Agent Node:** can support multiple tasks executing concurrently belonging to different frameworks with complete isolation, using **cgroups (Linux control groups)**, and managing resources consumption by each task to be limited to what is assigned to that task
3. **Frameworks:** representing different application paradigms, e.g., time sharing, batch, etc. An individual Framework will receive resource offers from the master and schedule tasks to the assigned agent nodes accordingly. **Marathon** is a framework for running long jobs + can be used to launch other frameworks + container orchestration..

Datacenter Components

1. Master Node:

- ❑ Resource allocation
- ❑ Resource de-allocation
- ❑ Resource reservation
- ❑ Resource isolation
- ❑ Resource monitoring
- ❑ Failure detection
- ❑ Package distribution
- ❑ Task starting, killing, cleanup
- ❑ Volume management
- ❑

2. Agent Node:

- ❑ Notify Master with available/free resources
- ❑ Execute tasks assigned to them by different Frameworks
- ❑ Use **cgroups** to isolate tasks from each other and to ensure that no task consume resources beyond what was assigned to them

3. Framework:

- ❑ Receive resource offers from the master and assign tasks to be executed on the assigned agent nodes
- ❑ Handle task failure or loss of agent node

Datacenter Components: Master


- **Master** is the cluster coordinator (cluster configuration Mgr).
- **The cluster will have multiple master nodes** (for fault tolerance). One master functions as the **leader** and propagates the metadata updates to the other masters using quorum replication. If the leader dies, one of the remaining masters is elected as the new leader.
- **The leader monitors** the state of the cluster.
- **Leader receives** resources available from the agent nodes and generates “resource offers” to the available frameworks based on some **policy**; framework scheduler may accept (part or all of the offered resources to use) or reject the offer!
- **The most significant factor** of a leader to manage large cluster is available RAM in the leader node.
- **Twitter** with cluster of **30,000 nodes** they use 5 master nodes.

Datacenter Components: Agent

- **Agent node** notifies the leader node with available/free resources (CPU, RAM, Storage, Ports)
- **Agent node** execute tasks assigned to them by any of the frameworks supported in the cluster
- **Two tasks belong to different frameworks** are completely isolated from each other as well the agent node guarantee that a task can use only resources assigned to them, i.e., no task can abuse the resources on that node.
- **You can assign resource to a specific role**, i.e., restricting this resource to a specific framework; otherwise, resources can be used by any application.

Datacenter Components: Framework

- **Tasks** are assigned to **Mesos containerizer** on the agent node by framework. In return, the containerizer will launch/destroy a task to run in a container.
- **Framework** is an application that uses Mesos APIs to receive **resource offers** from the Master and replies to instruct agent nodes to execute appropriate tasks.
- **Framework** handles task failure or loss of an agent node.



Datacenter (DC-OS) Components: Others

- Zookeeper
- Admin Router
- Metronome (Cron Service for DC/OS)
- Mesosphere/Marathon (Container Orchestration / Service scheduler)
- Kubernetes (orchestration s/w to deploy/manage containers)
- Messos-DNS: Name service discovery
- Cosmos: package manager; run on all master nodes. Used to install DC/OS Services



END